

CPSC 304 – September 18/19, 2024

Administrative notes

- Reminder: syllabus quiz due 20nd @10pm
- Reminder: “In-class” exercise 1 due 18th @10pm
- Reminder: Project groups due September 20
 - Please look on Canvas (milestone 0 description)
 - Submit the survey to have a group created (see Canvas for link)
- Tutorial next week: the relational model

CPSC 304

Introduction to Database Systems

The Relational Model

Textbook Reference

Database Management Systems: 3.1 - 3.5

Databases – the continuing saga



- So far we've learned that databases are handy for many reasons
- Before we can use them, we must design them
- In our last very exciting episode, we showed how to use ER diagrams to design the *conceptual schema*
- But the conceptual schema can only get us so far; we need to store data!
- Now we'll learn to use a *logical schema* to actually store the data. We'll be using the *relational model*.

Learning Goals



- Compare and contrast *logical* and *physical data independence*.
- Define the components (and synonyms) of the relational model: tables, rows, columns, keys, associations, etc.
- Create tables, including the attributes, keys, and field lengths, using Data Definition Language (DDL)
- Explain and differentiate the kinds of integrity constraints in a database
- Explain the purpose of referential integrity.
- Enforce referential integrity in a database using DML. Determine which delete, insert, or update policy to use when coding rules/defaults for referential integrity. Analyze the impact that a poor choice has.
- Map ER diagrams to the relational model (i.e., DDL), including constraints, weak entity sets, etc.

What do we want out of our logical schema representation?

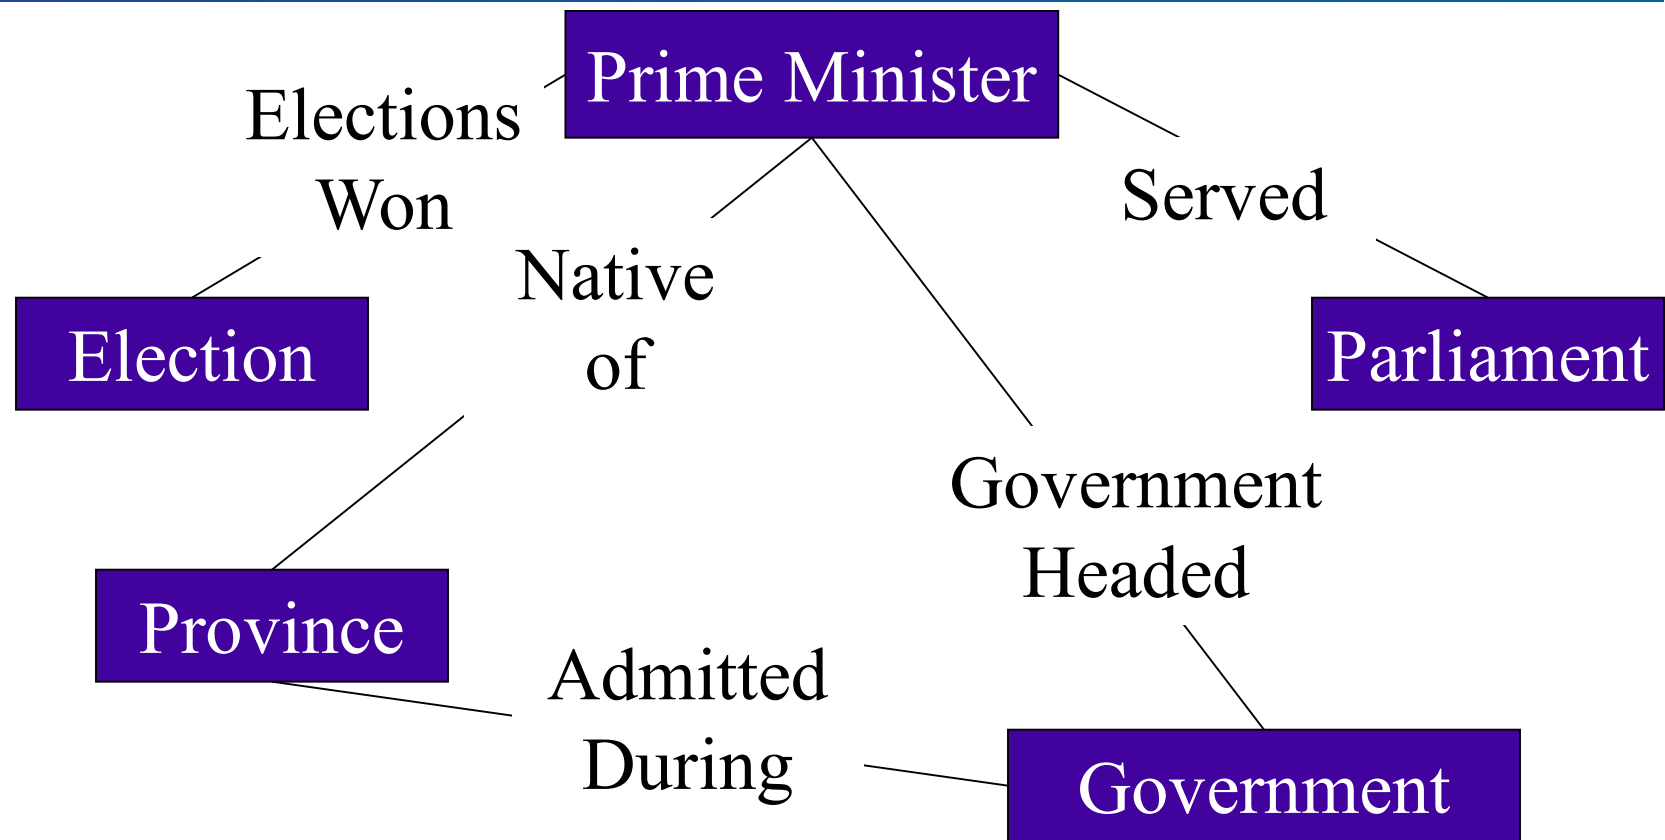
- Ability to store data – w/o worrying about blocks on disk
- Ability to query data easily
- A representation that is easy to understand
- A representation that we can easily adapt from conceptual schema
- Separate from application programming language



How did we get the relational model?

- Prior to the relational model, there were two main contenders
 - Network databases
 - Hierarchical databases
- Network databases had a complex data model
- Hierarchical databases integrated the application in the data model

Example Hierarchical Model



Example IMS (Hierarchical) query: Print the names of all the provinces admitted during a Liberal Government

```
DLITPLI:PROCEDURE (QUERY_PCB) OPTIONS (MAIN);
```

```
DECLARE QUERY_PCB POINTER;
```

```
/*Communication Buffer*/
```

```
DECLARE 1 PCB BASED(QUERY_PCB),
```

```
2 DATA_BASE_NAME CHAR(8),
```

```
2 SEGMENT_LEVEL CHAR(2),
```

```
2 STATUS_CODE CHAR(2),
```

```
2 PROCESSING_OPTIONS CHAR(4),
```

```
2 RESERVED_FOR_DLI FIXED BINARY(31,0),
```

```
2 SEGMENT_NAME_FEEDBACK CHAR(8)
```

```
2 LENGTH_OF_KEY_FEEDBACK_AREA FIXED BINARY(31,0),
```

```
2 NUMBER_OF_SENSITIVE_SEGMENTS FIXED BINARY(31,0),
```

```
2 KEY_FEEDBACK_AREA CHAR(28);
```

```
/* I/O Buffers*/
```

```
DECLARE PM_IO_AREA CHAR(65),
```

```
1 PRIME MINISTER DEFINED PM_IO_AREA,
```

```
2 PM_NUMBER CHAR(4),
```

```
2 PM_NAME CHAR(20),
```

```
2 BIRTHDATE CHAR(8)
```

```
2 DEATH_DATE CHAR(8),
```

```
2 PARTY CHAR(10),
```

```
2 SPOUSE CHAR(15);
```

```
DECLARE SADMIT_IO_AREA CHAR(20),
```

```
1 PROVINCE_ADMITTED DEFINED SADMIT_IO_AREA,
```

```
2 PROVINCE_NAME CHAR(20);
```

```
/* Segment Search Arguments */
```

```
DECLARE 1 PM_SSA STATIC UNALIGNED,
```

```
2 SEGMENT_NAME CHAR(8) INIT('PM '),
```

```
2 LEFT_PARENTHESIS CHAR(1) INIT('('),
```

```
2 FIELD_NAME CHAR(8) INIT('PARTY '),
```

```
2 CONDITIONAL_OPERATOR CHAR(2) INIT('='),
```

```
2 SEARCH_VALUE CHAR(10) INIT('Liberal '),
```

```
2 RIGHT_PARENTHESIS CHAR(1) INIT(')');
```

```
DECLARE 1 PROVINCE_ADMITTED_SSA STATIC UNALIGNED,
```

```
2 SEGMENT_NAME CHAR(8) INIT('SADMIT ');
```

```
/* Some necessary variables */
```

```
DECLARE GU CHAR(4) INIT('GU '),
```

```
GN CHAR(4) INIT('GN '),
```

```
GNP CHAR(4) INIT('GNP '),
```

```
FOUR FIXED BINARY(31) INIT(4),
```

```
SUCCESSFUL CHAR(2) INIT(' '),
```

```
RECORD_NOT_FOUND CHAR(2) INIT('GE');
```

```
/*This procedure handles IMS error conditions */
```

```
ERROR;PROCEDURE(ERROR_CODE);
```

```
*
```

```
*
```

```
*
```

```
END ERROR;
```

```
/*Main Procedure */
```

```
CALL PLITDLI(FOUR,GU,QUERY_PCB,PM_IO_AREA,PRESIDENT_SSA);
```

```
DO WHILE(PCB.STATUS_CODE=SUCCESSFUL);
```

```
CALL PLITDLI(FOUR,GNP,QUERY_PCB,SADMIT_IO_AREA,PROVINCE_ADMITTED_SSA);
```

```
DO WHILE(PCB.STATUS_CODE=SUCCESSFUL);
```

```
PUT EDIT(province_NAME)(A);
```

```
CALL PLITDLI(FOUR,GNP,QUERY_PCB,SADMIT_IO_AREA,PROVINCE_ADMITTED_SSA);
```

```
END;
```

```
IF PCB.STATUS_CODE NOT = RECORD_NOT_FOUND
```

```
THEN DO;
```

```
CALL ERROR(PCB.STATUS_CODE);
```

```
RETURN;
```

```
END;
```

```
CALL PLITDLI(FOUR,GN,QUERY_PCB,PM_IO_AREA,PRIME_MINISTER_SSA);
```

```
END;
```

```
IF PCB.STATUS_CODE NOT = RECORD_NOT_FOUND
```

```
THEN DO;
```

```
CALL ERROR(PCB.STATUS_CODE);
```

```
RETURN;
```

```
END;
```


```
END DLITPLI;
```


Relational model to the rescue!



- Introduced by Edgar Codd (IBM) in 1970
- Most widely used model today.
 - Vendors: IBM, Informix, Microsoft, Oracle, Sybase, etc.
- Competitor: object-oriented model
 - ObjectStore, Versant, Ontos
 - A synthesis emerging: *object-relational model*
 - Informix Universal Server, UniSQL, O2, Oracle, DB2
- Recent competitors (triggered by the needs of the web):
 - XML
 - NoSQL/Key-value stores/etc.

Popularity comparison of database management systems (DBMSs) worldwide as of June 2024, by category




Source
DB-Engines
© Statista 2024

Additional Information:
Worldwide; June 2024

Key points of the relational model

- Exceedingly simple to understand – main abstraction is represented as a table
- **Physical Data Independence** – ability to modify physical schema w/o changing logical schema
- **Logical Data Independence** – done with views
 - Ability to change the conceptual schema without changing applications



Structure of Relational Databases

- **Relational database**: a set of **relations**
- **Relation**: made up of 2 parts:
 - **Schema** : specifies name of relation, plus name and **domain** (type) of each **attribute**.
 - e.g., Student (*sid*: char[20], *name*: char[20], *address*: char[20], *phone*: char[20], *major*: char[20]).
 - **Instance** : a **table**, with rows and columns.
#Rows = cardinality
#Columns = arity / degree
- **Relational Database Schema**: collection of schemas in the database
- **Database Instance**: a collection of instances of its relations

Example of a Relation Instance

attribute,
column name

relation name → Student

| sid | name | address | phone | major |
|----------|----------|-------------------------------------|----------|-------|
| 99111120 | G. Jones | 1234 W. 12 th Ave., Van. | 889-4444 | CPSC |
| 92001200 | G. Smith | 2020 E. 18 th St., Van | 409-2222 | MATH |
| 94001020 | A. Smith | 2020 E. 18 th St., Van | 222-2222 | CPSC |
| 94001150 | S. Wang | null | null | null |

tuple, row, record →

domain value →

- degree/arity = 5; Cardinality = 4,
- Order of rows isn't important
- Order of attributes isn't important (except in some query languages)

Formal Structure

- Formally, a relation r is a set (a_1, a_2, \dots, a_n) where a_i is in D_i , the domain (set of allowed values) of the i -th attribute.
- Attribute values are atomic, i.e., integers, floats, strings
- A domain contains a special value ***null*** indicating that the value is not known.
- If A_1, \dots, A_n are attributes with domains D_1, \dots, D_n , then
 $(A_1:D_1, \dots, A_n:D_n)$
is a ***relation schema*** that defines a relation type –
sometimes we leave off the domains

Student (*sid*: char[20], *name*: char[20], *address*: char[20],
phone: char[20], *major*: char[20])

Example of a formal definition

Student

| sid | name | address | phone | major |
|----------|----------|--|----------|-------|
| 99111120 | G. Jones | 1234 W. 12 th Ave., Van. | 889-4444 | CPSC |
| 92001200 | G. Smith | 2020 E. 18 th St., Van | 409-2222 | MATH |
| 94001020 | A. Smith | 2020 E. 18 th St., Van | 222-2222 | CPSC |
| 94001150 | S. Wang | null | null | null |

Student (sid: integer, name: char[20], address: char[20],
phone: char[20], major: char[20])

Or, without the domains:

Student (sid, name, address, phone, major)

Relational Query Languages

- A major strength of the relational model: simple, powerful *querying* of data.
- Queries can be written intuitively; DBMS is responsible for efficient evaluation.
 - Precise semantics for relational queries.
 - Allows optimizer to extensively re-order operations, while ensuring that the answer does not change.

Clicker Question

- Here is a table representing a relation named R. Identify the attributes, schema, and tuples of R. Which of the following is **NOT** a true statement about R?

| A | B | C |
|---|----|----|
| 0 | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |
| 9 | 10 | 11 |

- A. R has four tuples.
- B. B is an attribute of R.
- C. (6,7,8) is a tuple of R.
- D. The schema of R is R(A,B,C).
- E. None of the above



The SQL Query Language



- SQL was **NOT** the first relational query language
- Developed by IBM (System R) in the 1970s
- Standards:
 - SQL-86
 - SQL-89 (minor revision)
 - SQL-92 (major revision, current standard)
 - SQL-99 (major extensions)

A peek at the SQL Query Language

Students

| sid | name | address | phone | major |
|----------|----------|-------------------------------------|----------|-------|
| 99111120 | G. Jones | 1234 W. 12 th Ave., Van. | 889-4444 | CPSC |
| 92001200 | G. Smith | 2020 E. 18 th St., Van | 409-2222 | MATH |
| 94001020 | A. Smith | 2020 E. 18 th St., Van | 222-2222 | CPSC |

- Find the id's, names and phones of all CPSC students:

```
SELECT sid, name, phone
FROM Students
WHERE major="CPSC"
```

| sid | name | phone |
|----------|----------|----------|
| 99111120 | G. Jones | 889-4444 |
| 94001020 | A. Smith | 222-2222 |

- To select whole rows , replace "SELECT sid, name, phone " with "SELECT * "

Simple, eh?

- We'll see more about how to query (**data manipulation language**) in Chapter 5.
- But you can't query without having a place to store your data, so back to how to create relations (**data definition language**)

Creating Relations in SQL/DDL

- The statement on the right creates the Student relation
 - the type (**domain**) of each attribute is specified and enforced when tuples are added or modified

```
CREATE TABLE Student
(sid      INTEGER,
 name     CHAR(20),
 address  CHAR(30),
 phone    CHAR(13),
 major    CHAR(4))
```

- The statement on right creates Grade information about courses that a student takes

```
CREATE TABLE Grade
(sid      INTEGER,
 dept     CHAR(4),
 course#  CHAR(3),
 mark     INTEGER)
```

Destroying and Altering Relations

`DROP TABLE Student`

- Destroys the relation Student. Schema information *and* tuples are deleted.

`ALTER TABLE Student`

`ADD COLUMN gpa REAL;`

- The schema of Students is altered by adding a new attribute; every tuple in current instance is extended with a *null* value in the new attribute.

Adding and Deleting Tuples

- Can insert a single tuple using:

```
INSERT  
INTO    Student (sid, name, address, phone, major)  
VALUES ('52033688', 'G. Chan', '1235 W. 33, Van',  
        '882-4444', 'PHYS')
```

- Can delete all tuples satisfying some condition (e.g., name = 'Smith'):

```
DELETE  
FROM    Student  
WHERE   name = 'Smith'
```

Powerful variants of these commands exist; more later

Integrity Constraints (ICs)

“Integrity is doing the right thing, even when no one is watching” - CS Lewis

- **IC**: condition that must be true for *any* instance of the database; e.g., domain constraints
 - ICs are specified when schema is defined
 - ICs are checked when relations are modified
- A *legal* instance of a relation is one that satisfies all specified ICs
 - DBMS should not allow illegal instances
 - Avoids data entry errors, too!
- The types of IC's depend on the data model.
 - What did we have for ER diagrams?
 - Next up: constraints for relational databases

Keys Constraints (for Relations)



- Similar to those for entity sets in the ER model
- One or more attributes in a relation form a **key** (or **candidate key**) for a relation, where S is the set of all attributes in the key, if :
 1. No distinct tuples can have the same values for all attributes in the key, and
 2. No subset of S is itself a key.
(it has to be minimal)
- One of the possible keys is chosen (by the DBA) to be the **primary key** (PK). `CREATE TABLE Student`
 - e.g.

| | |
|----------------------|-----------------------------------|
| <code>(sid</code> | <code>INTEGER PRIMARY KEY,</code> |
| <code>name</code> | <code>CHAR(20),</code> |
| <code>address</code> | <code>CHAR(30),</code> |
| <code>phone</code> | <code>CHAR(13),</code> |
| <code>major</code> | <code>CHAR(4))</code> |
 - ***sid*** is the primary key for Students

Quick Detour: Keys

Student

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |

Candidate keys:

- sid
- CWL
- SIN

Quick Detour: Keys

Clicker Question

Student

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |

Candidate keys:

- sid
- CWL
- SIN

In this database, can a student have multiple majors? A. Yes. B. No

Quick Detour: Keys

Clicker Question

Student

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |

Candidate keys:

- sid
- CWL
- SIN

In this database, can a student have multiple majors? A. Yes. B. No

Quick Detour: Keys

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |
| 4 | bPPuff | 222 | Blossom | Education | 18 | |

In this strange school, every student within the same major must have a unique name.

E.g., There is a student called Blossom in music so the music major can never admit another student named Blossom.

Quick Detour: Keys

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |
| 4 | bPPuff | 222 | Blossom | Education | 18 | |

Given only these four tuples, could {major, name} possibly be a candidate key?

A) Yes

B) No

1. No distinct tuples can have the same values for all attributes in the key, and
2. No proper subset of the potential key is itself a key.

Can you uniquely identify a tuple with just major? What about just name?

Quick Detour: Keys

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |
| 4 | bPPuff | 222 | Blossom | Education | 18 | |

Given only these four tuples, can we tell if {major, name} is candidate key?

A) Yes

B) No

No. There might exist other instances that do not adhere to this constraint. A key has to be a key in ALL instances.

But for the next few slides, let's assume that the database designer has declared that it is a candidate key.

Quick Detour: Keys (primary keys)

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |
| 4 | bPPuff | 222 | Blossom | Education | 18 | |

Candidate keys:

- sid
- CWL
- SIN
- {name, major}

Pick a candidate key to be your primary key

Quick Detour: Keys (Superkeys)

| sid | CWL | SIN | name | major | age | ... |
|-----|--------|-----|-----------|-----------|-----|-----|
| 1 | bpuff1 | 123 | Blossom | Music | 18 | ... |
| 2 | bpuff2 | 234 | Buttercup | Physics | 18 | ... |
| 3 | bpuff3 | 456 | Bubbles | Education | 18 | ... |
| 4 | bPPuff | 222 | Blossom | Education | 18 | |

Candidate keys:

- sid
- CWL
- SIN
- {name, major}

What is a superkey? A key plus zero or more additional attributes. Examples (not exhaustive) All the candidate keys plus:

- {sid, name}
- {CWL, major}
- {name, major, age}

Just a preview. I promise we'll do more superkeys later.

Keys Constraints in SQL

- A **PRIMARY KEY** constraint specifies a table's primary key
 - values for primary key must be unique
 - a primary key attributes cannot be *null*
- Other keys are specified using the **UNIQUE** constraint
 - values for a group of attributes must be unique (if they are not null)
 - these attributes can be *null*
- Key constraints are checked when
 - new values are inserted
 - values are modified

Clicker question

Does the constraint PRIMARY KEY(Dept, Course#) hold for this instance?

| Course# | Dept | Title |
|---------|------|--------------------------------------|
| 304 | CPSC | Introduction to Relational Databases |
| 304 | PSYC | Brain and Behaviour |

A. Yes

B. No

C. It depends

Clicker question

Does the constraint PRIMARY KEY(Dept, Course#) hold for this instance?

| Dept | Course# | Term | Section |
|------|---------|--------|---------|
| CPSC | 304 | 2017W2 | 201 |
| CPSC | 304 | 2017W2 | 202 |

A. Yes

B. No

C. It depends

Keys Constraints in SQL (cont')

(Ex.1- Normal) “For a given student and course, there is a single grade.”

```
CREATE TABLE Grade
(sid      INTEGER,
dept     CHAR(4),
course#  CHAR(3),
mark     INTEGER,
PRIMARY KEY (sid,dept,course#) )
```

VS.

(Ex.2 - Silly) “Students can take a course once, and receive a single grade for that course; further, no two students in a course receive the same grade.”

```
CREATE TABLE Grade2
(sid      INTEGER,
dept     CHAR(4),
course#  CHAR(3),
mark     CHAR(2),
PRIMARY KEY (sid,dept,course#),
UNIQUE (dept,course#,mark) )
```

Keys Constraints in SQL (cont')

For single attribute keys, can also be declared on the same line as the attribute.

```
CREATE TABLE Student
    (sid      INTEGER PRIMARY KEY,
     name     CHAR(20),
     address  CHAR(30),
     phone    CHAR(13),
     major    CHAR(4))
```

Foreign Keys Constraints



- **Foreign key** : Set of attributes in one relation used to 'reference' a tuple in another relation.
 - **Must correspond to the primary key of the other relation.**
 - Like a 'logical pointer'.
- E.g.:
Grade(*sid*, *dept*, *course#*, *grade*)
 - *sid* is a foreign key referring to **Student**:
 - (*dept*, *course#*) is a foreign key referring to **Course**
- **Referential integrity**: All foreign keys reference existing entities.
 - i.e. there are no dangling references
 - all foreign key constraints are enforced

Let's look at students and grades again...

Student

| sid | name | ... |
|-----|-----------|-----|
| 1 | Blossom | ... |
| 2 | Buttercup | ... |
| 3 | Bubbles | ... |
| 4 | Blossom | |

Grade

| sid | dept | cnum | grade |
|-----|------|------|-------|
| 2 | CPSC | 304 | 90 |
| 2 | MATH | 221 | 90 |
| 2 | EPSE | 223 | 90 |
| 1 | MUSC | 103 | 90 |

Do we want tuples in Grade to ever refer to a sid that does not exist?

NO!

Idea: Check the Student table each time we insert a tuple into the Grade table to see if sid exists.

Enforcing Referential Integrity

Student

| sid | name | ... |
|-----|-----------|-----|
| 1 | Blossom | ... |
| 2 | Buttercup | ... |
| 3 | Bubbles | ... |
| 4 | Blossom | |

Grade

| sid | dept | cnum | grade |
|-----|------|------|-------|
| 2 | CPSC | 304 | 90 |
| 2 | MATH | 221 | 90 |
| 2 | EPSE | 223 | 90 |
| 1 | MUSC | 103 | 90 |

A foreign key is a set of attributes in one relation (e.g., Grades.sid) used to 'reference' a tuple in another relation (e.g., Students.sid).

Enforcing Referential Integrity

```
CREATE TABLE Grade (  
    sid INTEGER,  
    dept CHAR(4),  
    cnum CHAR(3),  
    mark INTEGER,  
    PRIMARY KEY (sid,dept,cnum),  
    FOREIGN KEY (sid) REFERENCES Student,  
    FOREIGN KEY (dept, cnum)  
        REFERENCES Course(dept, cnum)  
    )
```

Grade

| sid | dept | cnum | grade |
|-----|------|------|-------|
| 2 | CPSC | 304 | 90 |
| 2 | MATH | 221 | 90 |
| 2 | EPSE | 223 | 90 |
| 1 | MUSC | 103 | 90 |

Foreign Keys in SQL

Only students listed in the Student relation should be allowed to have grades for courses that are listed in the Course relation.

CREATE TABLE Grade

(sid INTEGER, dept CHAR(4), course# CHAR(3), mark INTEGER,
PRIMARY KEY (sid,dept,course#),

FOREIGN KEY (sid) REFERENCES Student,

FOREIGN KEY (dept, course#) REFERENCES Course(dept, cnum))

Sometimes you can not specify which attributes are referenced, but in this case they are needed. Never hurts to include them!

Grade

| sid | dept | course# | mark |
|-------|------|---------|------|
| 53666 | CPSC | 101 | 80 |
| 53666 | RELG | 100 | 45 |
| 53650 | MATH | 200 | null |
| 53666 | HIST | 201 | 60 |

Student

| sid | name | address | Phone | major |
|-------|----------|---------|-------|-------|
| 53666 | G. Jones | | ... | ... |
| 53688 | J. Smith | | ... | ... |
| 53650 | G. Smith | | ... | ... |

Enforcing Referential Integrity

- *sid* in Grade is a foreign key that references Student.
- What should be done if a Grade tuple with a non-existent student id is inserted? (*Reject it!*)
- What should be done if a **Student tuple** is deleted?
 - Also delete all Grade tuples that refer to it?
 - Disallow deletion of this particular Student tuple?
 - Set *sid* in Grade tuples that refer to it, to *null*, (the special value denoting '*unknown*' or '*inapplicable*'.)
 - problem if *sid* is part of the primary key
 - Set *sid* in Grade tuples that refer to it, to a *default sid*.
- Similar if primary key of a Student tuple is updated

Referential Integrity in SQL/92

- SQL/92 supports all 4 options on deletes and updates.
 - Default is **NO ACTION** (*delete/update is rejected*)
 - **CASCADE** (also updates/deletes all tuples that refer to the updated/deleted tuple)
 - **SET NULL / SET DEFAULT** (referencing tuple value is set to the default foreign key value)

```
CREATE TABLE Grade
(sid CHAR(8), dept CHAR(4),
 course# CHAR(3), mark INTEGER,
PRIMARY KEY (sid,dept,course#),
FOREIGN KEY (sid)
REFERENCES Student(sid)
ON DELETE CASCADE
ON UPDATE CASCADE
FOREIGN KEY (dept, course#)
REFERENCES
Course(dept,course#)
ON DELETE SET DEFAULT
ON UPDATE CASCADE );
```

Clicker question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...  
                    PRIMARY KEY (bid),  
                    FOREIGN KEY (sid) REFERENCES STUDENTS  
                    ON DELETE CASCADE);
```

If bid = 1000 and sid = 5678 for a row in Table BMW, choose the best answer

- A. If the row for sid value 5678 in STUDENTS is deleted, then the row with bid = 1000 in BMW is automatically deleted.
- B. If a row with sid value 5678 in BMW is deleted, then the row with sid=5678 in STUDENTS is automatically deleted.
- C. Both of the above.

Clicker question

Consider the following table definition.

```
CREATE TABLE BMW ( bid INTEGER, sid INTEGER, ...  
                    PRIMARY KEY (bid),  
                    FOREIGN KEY (sid) REFERENCES STUDENTS  
                    ON DELETE CASCADE);
```

If bid = 1000 and sid = 5678 for a row in Table BMW, choose the best answer

- A. If the row for sid value 5678 in STUDENTS is deleted, then the row with bid = 1000 in BMW is automatically deleted.
- B. If a row with sid value 5678 in BMW is deleted, then the row with sid=5678 in STUDENTS is automatically deleted.
- C. Both of the above.

BMW

| bid | Sid |
|------|------|
| 1000 | 5678 |

Student

| sid | name | Address |
|------|-------|---------|
| 5678 | James | Null |




Where do ICs Come From?

- ICs are based upon the real-world semantics being described (in the database relations).
- We *can* check a database instance to verify an IC, but we *cannot* tell the ICs by looking at the instance.
 - For example, even if all student names differ, we cannot assume that name is a key.
 - An IC is a statement about *all possible* instances.
- All constraints must be identified during the conceptual design.
- Some constraints can be explicitly specified in the conceptual model
 - Key and foreign key ICs are shown on ER diagrams.
- Others are written in a more general language.

Logical DB Design: ER to Relational


Entity sets to tables.

- Each entity set is mapped to a table.
 - entity attributes become table attributes
 - entity keys become table keys



```
CREATE TABLE MoviePeople  
(ID CHAR(11),  
Name CHAR(20),  
BirthDate DATE,  
PRIMARY KEY (ID))
```


Relationship Sets to Tables



- A relationship set id is mapped to a single relation (table).
- Simple case: relationship has no constraints (i.e. many-to-many)
- In this case, attributes of the table must include:
 - Keys for each participating entity set as foreign keys.
 - This is a *key* for the relation.
 - All descriptive attributes.

```
CREATE TABLE WorkOn(  
  MPID    CHAR(11),  
  MID     INTEGER,  
  Role    CHAR(20),  
  PRIMARY KEY (MPID, MID),  
  FOREIGN KEY (MPID)  
    REFERENCES MoviePeople,  
  FOREIGN KEY (MID)  
    REFERENCES Movies)
```

Example: Many to Many Relationships




Scrooge McDuck
has worked on
two movies




Pua has worked
on one movie
(Moana)



Example: Many to Many Relationships



Can we reduce redundancy by combining these tables in any way?



MoviePeople

| <u>ID</u> | Name | Birthdate |
|-----------|---------|-----------|
| 1 | Scrooge | ... |
| 2 | Pua | ... |

WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1 | 1 | ... |
| 1 | 2 | ... |
| 2 | 3 | ... |

Movies

| <u>ID</u> | Title |
|-----------|-------------------|
| 1 | Ducktales |
| 2 | A Christmas Carol |
| 3 | Moana |

Example: Many to Many Relationships

MoviePeople

| <u>ID</u> | Name | Birthdate |
|-----------|---------|-----------|
| 1 | Scrooge | ... |
| 2 | Pua | ... |

Movies

| <u>ID</u> | Title |
|-----------|-------------------|
| 1 | Ducktales |
| 2 | A Christmas Carol |
| 3 | Moana |

WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1 | 1 | ... |
| 1 | 2 | ... |
| 2 | 3 | ... |

Can we integrate the information in WorkOn into MoviePeople?

Example: Many to Many Relationships

MoviePeople

| <u>ID</u> | Name | Birthdate | MID |
|-----------|---------|-----------|-----|
| 1 | Scrooge | ... | |
| 2 | Pua | ... | |

Movies

| <u>ID</u> | Title |
|-----------|-------------------|
| 1 | Ducktales |
| 2 | A Christmas Carol |
| 3 | Moana |

WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1 | 1 | ... |
| 1 | 2 | ... |
| 2 | 3 | ... |

Do we put MID 1 or MID 2? We can't have both in the column (i.e., we can't store a list there)

Not much we can do in terms of reducing tables

Example: Many to Many Relationships

MoviePeople

| <u>ID</u> | Name | Birthdate |
|-----------|---------|-----------|
| 1 | Scrooge | ... |
| 2 | Pua | ... |

Movies

| <u>ID</u> | Title | <u>MPID</u> |
|-----------|-------------------|-------------|
| 1 | Ducktales | |
| 2 | A Christmas Carol | |
| 3 | Moana | |




WorkOn

| <u>MPID</u> | <u>MID</u> | Role |
|-------------|------------|------|
| 1 | 1 | ... |
| 1 | 2 | ... |
| 2 | 3 | ... |

What if we have more than one person work on this movie?
Same problem as before!

Not much we can do in terms of reducing tables

Relationship Sets to Tables



MoviePeople

| ID | Name | Birth Date |
|----|------|------------|
| 1 | | |
| 2 | | |

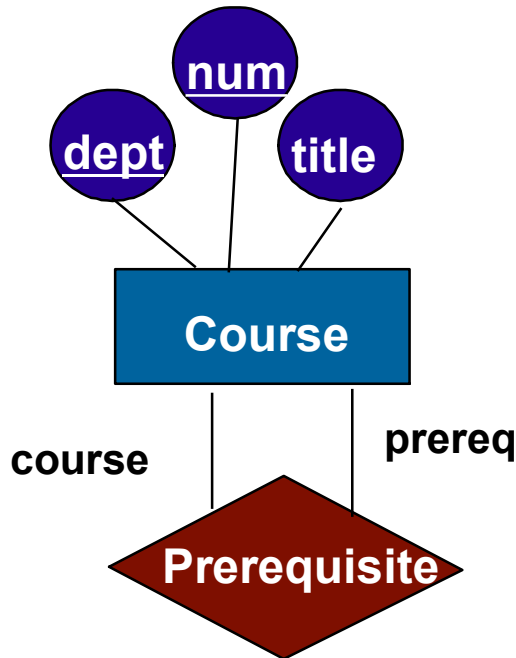
WorkOn

| MoviePeople ID | Movie ID |
|----------------|----------|
| 1 | 1 |
| 2 | 1 |

Movies

| ID | Title |
|----|-------|
| 1 | |
| 2 | |

Relationship Sets to Tables (cont')



- In some cases, we need to use the roles:

```
CREATE TABLE Prerequisite(  
  course_dept  CHAR(4),  
  course_num   CHAR(3),  
  prereq_dept  CHAR(4),  
  prereq_num   CHAR(3),  
  PRIMARY KEY (course_dept, course_num,  
               prereq_dept, prereq_num),  
  FOREIGN KEY (course_dept, course_num)  
    REFERENCES Course(dept, num),  
  FOREIGN KEY (prereq_dept, prereq_num)  
    REFERENCES Course(dept, num))
```

Clicker question: Relationship sets to table

Given the table below, does every course have to have a prereq?

```
CREATE TABLE Prerequisite(  
  course_dept CHAR(4),  
  course_num CHAR(3),  
  prereq_dept CHAR(4),  
  prereq_num CHAR(3),  
  PRIMARY KEY (course_dept, course_num, prereq_dept, prereq_num),  
  FOREIGN KEY (course_dept, course_num) REFERENCES Course(dept, num),  
  FOREIGN KEY (prereq_dept, prereq_num) REFERENCES Course(dept, num))
```

B. Yes

C. No

Clicker question: Relationship sets to table

Given the table below, does every course have to have a prereq?

```
CREATE TABLE Prerequisite(  
  course_dept CHAR(4),  
  course_num CHAR(3),  
  prereq_dept CHAR(4),  
  prereq_num CHAR(3),  
  PRIMARY KEY (course_dept, course_num, prereq_dept, prereq_num),  
  FOREIGN KEY (course_dept, course_num) REFERENCES Course(dept, num),  
  FOREIGN KEY (prereq_dept, prereq_num) REFERENCES Course(dept, num))
```

B. Yes

C. No

Relationship sets to table clicker question explained

Course

| DEPT | NUM | Title |
|------|-----|-------|
| CPSC | 100 | ... |
| CPSC | 103 | ... |
| CPSC | 107 | ... |
| CPSC | 210 | ... |

Prerequisite

| COURSE_DEPT | COURSE_NUM | PREREQ_DEPT | PREREQ_NUM |
|-------------|------------|-------------|------------|
| CPSC | 107 | CPSC | 103 |
| CPSC | 210 | CPSC | 107 |

There is no requirement that every course appears in the prerequisite table!

```
CREATE TABLE Prerequisite(  
  course_dept CHAR(4),  
  course_num CHAR(3),  
  prereq_dept CHAR(4),  
  prereq_num CHAR(3),  
  PRIMARY KEY (course_dept,  
               course_num, prereq_dept,  
               prereq_num),  
  FOREIGN KEY (course_dept,  
               course_num) REFERENCES  
    Course(dept, num),  
  FOREIGN KEY (prereq_dept,  
               prereq_num) REFERENCES  
    Course(dept, num))
```