

# CPSC 304 – Administrative notes

## October 16 & October 17, 2024

---

- Final exam: December 16 @ noon
- Project:
  - Milestone 3: Project check in – due October 25
    - Sign up coming next week
  - Milestone 4: Project implementation – due November 29
  - Milestone 5: Group demo – week of December 2
  - Milestone 6: Individual Assessment – Due November 29
- October 22: Midterm @ 6PM
  - See Piazza for important midterm information, including how to find where you will write the midterm
    - If you can't find your location on PrairieTest, make sure you're on the US site:  
<https://us.prairietest.com/>
  - **All blank answers will be marked as incorrect!**
- Lecture before the midterm is open office hours (this Friday/next Tuesday)
- Tutorials: pivoting to project work/open office hours – next week is nominally SQL Plus
- Note that your SQL accounts and repositories go away with term end!

# CPSC 304

## **Introduction to Database Systems**

---

Structured Query Language (SQL)

Textbook Reference  
Database Management Systems: Chapter 5

# Databases: the continuing saga

---



When last we left databases...

- We had decided they were great things
- We knew how to conceptually model them in ER diagrams
- We knew how to logically model them in the relational model
- We knew how to normalize our database relations
- We could formally specify queries in Relational Algebra

Now: how do most people write queries? SQL!

# Learning Goals

---



- Given the schemas of a relation, create SQL queries using: SELECT, FROM, WHERE, EXISTS, NOT EXISTS, UNIQUE, NOT UNIQUE, ANY, ALL, DISTINCT, GROUP BY and HAVING.
- Show that there are alternative ways of coding SQL queries to yield the same result. Determine whether or not two SQL queries are equivalent.
- Given a SQL query and table schemas and instances, compute the query result.
- Translate a query between SQL and RA.
- Comment on the relative expressive power of SQL and RA.
- Explain the purpose of NULL values and justify their use. Also describe the difficulties added by having nulls.
- Create and modify table schemas and views in SQL.
- Explain the role and advantages of embedding SQL in application programs.
- Write SQL for a small-to-medium sized programming application that requires database access.
- Identify the pros and cons of using general table constraints (e.g., CONSTRAINT, CHECK) and triggers in databases.

# Coming up in SQL...

---

- Data Definition Language (reminder)
- Basic Structure
- Set Operations
- Aggregate Functions
- Null Values
- Nested Subqueries
- Modification of the Database
- Views
- Integrity Constraints
- Putting SQL to work in an application

# The SQL Query Language

---

- Need for a standard since relational queries are used by many vendors
- Consists of several parts:
  - Data Definition Language (DDL)  
(a blast from the past (Chapter 3))
  - Data Manipulation Language (DML)
    - Data Query
    - Data Modification

# This is going to be more interesting if you can write your own queries

---

- The sqlplus tutorial will walk you through the steps (do it now, or in tutorial next week)
- <https://www.students.cs.ubc.ca/~cs-304/resources.html>

# Creating Tables in SQL(DDL) Revisited

- A SQL relation is defined using the **create table** command:

**create table**  $r$  ( $A_1 D_1, A_2 D_2, \dots, A_n D_n,$   
                  (integrity-constraint<sub>1</sub>),  
                  ...,  
                  (integrity-constraint<sub>k</sub>))

- *Integrity constraints can be:*

- *primary and candidate keys*
- *foreign keys*

- Example:

```
CREATE TABLE Student
(sid    CHAR(20),
 name  CHAR(20),
 address CHAR(20),
 phone CHAR(8),
 major  CHAR(4),
PRIMARY KEY (sid))
```



# Domain Types in SQL

## Reference Sheet

---

- **char(*n*)**. Fixed length character string with length *n*.
- **varchar(*n*)**. Variable length character strings, with maximum length *n*.
- **int**. Integer (machine-dependent).
- **smallint**. Small integer (machine-dependent).
- **numeric(*p*,*d*)**. Fixed point number, with user-specified precision of *p* digits, with *d* digits to the right of decimal point.
- **real, double precision**. Floating point and double-precision floating point numbers, with machine-dependent precision.
- **float(*n*)**. Floating point number, with user-specified precision of at least *n* digits.
- Null values are allowed in all the domain types.  
To prohibit null values declare attribute to be **not null**
- **create domain** in SQL-92 and 99 creates user-defined domain types  
**create domain *person-name* char(20) not null**

# Date/Time Types in SQL

## Reference Sheet

---

- **date.** Dates, containing a (4 digit) year, month and date
  - E.g. **date** '2001-7-27'
- **time.** Time of day, in hours, minutes and seconds.
  - E.g. **time** '09:00:30'      **time** '09:00:30.75'
- **timestamp:** date plus time of day
  - E.g. **timestamp** '2001-7-27 09:00:30.75'
- **Interval:** period of time
  - E.g. Interval '1' day
  - Subtracting a date/time/timestamp value from another gives an interval value
  - Interval values can be added to date/time/timestamp values
- Relational DBMS offer a variety of functions to
  - extract values of individual fields from date/time/timestamp
  - convert strings to dates and vice versa
  - For instance in Oracle (date is a timestamp):
    - TO\_CHAR( date, format)
    - TO\_DATE( string, format)
    - format looks like: 'DD-Mon-YY HH:MI.SS'

# Running Example (should look familiar)

---

Movie(MovieID, Title, Year)

StarsIn(MovieID, StarID, Character)

MovieStar(StarID, Name, Gender)

# Basic SQL Query

- SQL is based on set and relational operations
- A typical SQL query has the form:

**SELECT**  $A_1, A_2, \dots, A_n$   
**FROM**  $r_1, r_2, \dots, r_m$   
**WHERE**  $P$

SELECT	<i>target-list</i>
FROM	<i>relation-list</i>
WHERE	<i>qualification</i>

- $A_i$ s represent attributes
  - $r_i$ s represent relations
  - $P$  is a predicate.
- $\pi \rightarrow$  SELECT clause  
 $\sigma \rightarrow$  WHERE clause  
 $\bowtie \rightarrow$  FROM and WHERE clause
- The result of a SQL query is a table (relation)
  - By default, duplicates are not eliminated in SQL relations, which are **bags** or **multisets** and not sets
  - Let's compare to relational algebra...

# Basic SQL/RA Comparison example 1

---

- Find the titles of movies

Movie( <u>MovieID</u> , Title, Year)
StarsIn( <u>MovieID</u> , <u>StarID</u> , Character)
MovieStar( <u>StarID</u> , Name, Gender)

# Basic SQL/RA Comparison example 1

---

- Find the titles of movies

$\pi_{\text{Title}}(\text{Movie})$

- In SQL,  $\pi$  is in the SELECT clause
- Select only a subset of the attributes

```
SELECT Title  
FROM   Movie
```

- Note duplication can happen!
  - You can get the same value multiple times

# Basic SQL/RA Comparison example 1

- You can also refer to an attribute by (relation name).(attribute name)

```
SELECT  Movie.Title  
FROM    Movie
```

- Since we are only working with one relation, you don't need to specify where the attribute comes from
  - This is useful when you have multiple relations that share the same attribute name

# Clicker Question: SQL projection

- Given the table scores:

what is result of  
`SELECT Score1,  
          Score2  
FROM Scores`

Team1	Team2	Score1	Score2
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

- Which of the following rows is in the answer?

- A. (1,2)
- B. (5,3)
- C. (8,6)
- D. All are in the answer
- E. None are in the answer



clickerprojection.sql

## Clicker Question: SQL projection

- Given the table scores:

what is result of  
`SELECT Score1,  
          Score2  
FROM Scores`

Team1	Team2	Score1	Score2
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12

- Which of the following rows is in the answer?

A. (1,2)

B. (5,3) **Correct**  $\pi_{score1,score2}(Scores)$

C. (8,6)

D. All are in the answer

E. None are in the answer

# In SQL, $\sigma$ is in *WHERE* clause

---

```
SELECT *  
FROM   Movie  
WHERE  Year > 1939
```

Movie( <u>MovieID</u> , Title, Year) StarsIn( <u>MovieID</u> , <u>StarID</u> , Character) MovieStar( <u>StarID</u> , Name, Gender)
--

You can use:

- attribute names of the relation(s) used in the FROM.

- comparison operators: =, <>, <, >, <=, >=

- apply arithmetic operations: rating\*2

- operations on strings (e.g., "||" for concatenation).

- Lexicographic order on strings.

- Pattern matching: s LIKE p

- Special stuff for comparing dates and times.

## Basic SQL/RA Comparison example 2

---

Find female movie stars

Movie( <u>MovieID</u> , Title, Year)
StarsIn( <u>MovieID</u> , <u>StarID</u> , Character)
MovieStar( <u>StarID</u> , Name, Gender)

## Basic SQL/RA Comparison example 2

---

Find female movie stars

$\sigma_{\text{Gender} = \text{'female'}} \text{MovieStar}$

```
SELECT *  
FROM   MovieStar  
WHERE  Gender='female'
```

Movie( <u>MovieID</u> , Title, Year) StarsIn( <u>MovieID</u> , <u>StarID</u> , Character) MovieStar( <u>StarID</u> , Name, Gender)
--

# Clicker Question: Selection

- Consider Scores(Team, Opponent, RunsFor, RunsAgainst) and query

```
SELECT *  
FROM Scores  
WHERE  
    RunsFor > 5
```

- Which tuple is in the result?

- A. (Swallows, Carp, 6, 4)
- B. (Swallows, Carp, 4)
- C. (12)
- D. (\*)

Team	Opponent	RunsFor	RunsAgainst
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

clickerselection.sql

# Clicker Question: Selection

- Consider Scores(Team, Opponent, RunsFor, RunsAgainst) and query

```
SELECT *  
FROM Scores  
WHERE  
    RunsFor > 5
```

- Which tuple is in the result?

- A. (Swallows, Carp, 6, 4)
- B. (Swallows, Carp, 4)
- C. (12)
- D. (\*)

Team	Opponent	RunsFor	RunsAgainst
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

answer A

# Selection & Projection – together forever in SQL

---



We can put these together:

- What are the names of female movie stars?
- What are the titles of movies from prior to 1939?

Movie( <u>MovieID</u> , Title, Year) StarsIn( <u>MovieID</u> , <u>StarID</u> , Character) MovieStar( <u>StarID</u> , Name, Gender)
--

# Selection & Projection – together forever in SQL



We can put these together:

- What are the names of female movie stars?

```
SELECT name  
FROM MovieStar  
WHERE Gender = 'female'
```

- What are the titles of movies from prior to 1939?

```
SELECT title  
FROM Movie  
WHERE year < 1939
```

Movie( <u>MovieID</u> , Title, Year) StarsIn( <u>MovieID</u> , <u>StarID</u> , Character) MovieStar( <u>StarID</u> , Name, Gender)
--



# Selection example (dates)

---

reserves

SID	BID	Day
22	101	2010-10-10
22	102	2010-10-10
22	103	2010-10-08
22	104	2010-07-10
31	102	2010-11-10
31	103	2010-11-06
31	104	2010-11-12
58	102	2010-11-08
58	103	2010-11-12

```
SELECT *  
FROM reserves  
WHERE day < DATE'2010-11-01'
```

SID	BID	Day
22	101	2010-10-10
22	102	2010-10-10
22	103	2010-10-08
22	104	2010-07-10

## Basic SQL/RA comparison example 3

---

- Find the person names and character names of those who have been in movies

$\pi_{\text{Character, Name}}(\text{StarsIn} \bowtie_{\text{StarsIn.StarID} = \text{MovieStar.StarID}} \text{MovieStar})$

- In order to do this we need to use joins.  
How can we do joins in SQL?

$\pi \rightarrow$  SELECT clause

$\sigma \rightarrow$  WHERE clause

$\bowtie \rightarrow$  FROM and WHERE clause

# Joins in SQL

---

```
SELECT Character, Name  
FROM StarsIn s, MovieStar m  
WHERE s.StarID = m.StarID
```

- Cross product specified by FROM clause
- Can alias relations (e.g., “StarsIn s”)
- Conditions specified in WHERE clause

# Clicker Question: Simple Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.a, S.b  
FROM R, S  
WHERE R.b = S.a

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) appears once.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

# Simple Joins Results

R.a	R.b	S.a	S.B
0	0	0	0
1	0	0	0
0	0	0	1
1	0	0	1
0	1	1	0
1	1	1	0
0	1	1	1
1	1	1	1

# Clicker Question: Simple Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.a, S.b  
FROM R, S  
WHERE R.b = S.a

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) appears once.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

False. Only R(0,1) and S(1, 0)

False. R.b <> S.a

True. R(1,1) and S(1, 0)

# Clicker Question: Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.b, T.b  
FROM R, S, T  
WHERE R.b = S.a AND S.b <> T.b (note: <> == 'not equals')

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) does not appear.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

Example of  
the cross  
product  
obtained  
from the  
**first tuple**  
in  $R \times S \times T$ .

R.a	R.b	S.a	S.b	T.a	T.b
0	0	0	0	0	0
0	0	0	0	0	1
0	0	0	0	1	0
0	0	0	0	1	1
0	0	0	1	0	0
0	0	0	1	0	1
0	0	0	1	1	0
0	0	0	1	1	1
0	0	1	0	0	0
0	0	1	0	0	1
0	0	1	0	1	0
0	0	1	0	1	1
0	0	1	1	0	0
0	0	1	1	0	1
0	0	1	1	1	0
0	0	1	1	1	1



Dropping all  
the tuples  
that don't  
fulfill the  
WHERE  
clause.

Repeat this  
process for  
the other  
tuples in R.

R.a	R.b	S.a	S.b	T.a	T.b
0	0	0	0	0	0
0	0	0	0	0	1
0	0	0	0	1	0
0	0	0	0	1	1
0	0	0	1	0	0
0	0	0	1	0	1
0	0	0	1	1	0
0	0	0	1	1	1
0	0	1	0	0	0
0	0	1	0	0	1
0	0	1	0	1	0
0	0	1	0	1	1
0	0	1	1	0	0
0	0	1	1	0	1
0	0	1	1	1	0
0	0	1	1	1	1

What we  
get after  
processing  
the FROM  
and WHERE  
clauses.

R.a	R.b	S.a	S.b	T.a	T.b
0	0	0	0	0	1
0	0	0	0	1	1
0	0	0	1	0	0
0	0	0	1	1	0
0	1	1	0	0	1
0	1	1	0	1	1
0	1	1	1	0	0
0	1	1	1	1	0
1	0	0	0	0	1
1	0	0	0	1	1
1	0	0	1	0	0
1	0	0	1	1	0
1	1	1	0	0	1
1	1	1	0	1	1
1	1	1	1	0	0
1	1	1	1	1	0

R.a	R.b	S.b	T.b
0	0	0	1
0	0	0	1
0	0	1	0
0	0	1	0
0	1	0	1
0	1	0	1
0	1	1	0
0	1	1	0
1	0	0	1
1	0	0	1
1	0	1	0
1	0	1	0
1	1	0	1
1	1	0	1
1	1	1	0
1	1	1	0

# Clicker Question: Joins

Consider R :

a	b
0	0
0	1
1	0
1	1

S:

a	b
0	0
0	1
1	0
1	1

T:

a	b
0	0
0	1
1	0
1	1

SELECT R.a, R.b, S.b, T.b  
FROM R, S, T  
WHERE R.b = S.a AND S.b <> T.b (note: <> == 'not equals')

Compute the results

Which of the following are true:

- A. (0,1,1,0) appears twice.
- B. (1,1,0,1) does not appear.
- C. (1,1,1,0) appears once.
- D. All are true
- E. None are true

True R(0,1) S(1,1), T(0,0)&  
R(0,1), S(1,1), T(1,0),

False: R(1,1), S(1,0), T(0,1)

False: like A but use R(1, 1)

# So how does a typical SQL query relate to relational algebra then?

---

SQL:

**SELECT**  $A_1, A_2, \dots, A_n$   
**FROM**  $r_1, r_2, \dots, r_m$   
**WHERE**  $P$

Is approximately equal to  
Relational algebra

$$\pi_{A_1, A_2, \dots, A_n}(\sigma_P(r_1 \times r_2 \times \dots \times r_m))$$

Difference? Duplicates.  
Remove them? **Distinct**

# Using DISTINCT

---

- Find the names of actors who have been in at least one movie

```
SELECT  DISTINCT Name
FROM    StarsIn S, MovieStar M
WHERE   S.StarID = M.StarID
```

- Would removing DISTINCT from this query make a difference?
- Note: on the exams, if we ask for a general question like “find all the names”, we expect duplicates to be removed. When in doubt, keep the DISTINCT rather than reasoning through if you need it!

# Clicker question: distinction

Consider the relation:  
Scores(Team, Opponent,  
RunsFor, RunsAgainst) and  
the query:

```
SELECT DISTINCT Team,  
                RunsFor  
FROM    Scores
```

Which is true:

- A. 1 appears once
- B. 5 appears twice
- C. 6 appears 4 times
- D. All are true
- E. None are true

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

# Clicker question: distinction

Consider the relation:  
Scores(Team, Opponent,  
RunsFor, RunsAgainst) and  
the query:

```
SELECT DISTINCT Team,  
                RunsFor  
FROM    Scores
```

Which is true:

- A. 1 appears once
- B. 5 appears twice **Correct**
- C. 6 appears four times
- D. All are true
- E. None are true

Team	Opponent	Runs For	Runs Against
Dragons	Tigers	5	3
Carp	Swallows	4	6
Bay Stars	Giants	2	1
Marines	Hawks	5	3
Ham Fighters	Buffaloes	1	6
Lions	Golden Eagles	8	12
Tigers	Dragons	3	5
Swallows	Carp	6	4
Giants	Bay Stars	1	2
Hawks	Marines	3	5
Buffaloes	Ham Fighters	6	1
Golden Eagles	Lions	12	8

clickerdistinction.sql



# Join Example

---

- Find the names of all movie stars who have been in a movie

# Join Example

---

- Find the names of all movie stars who have been in a movie

SELECT Name

FROM StarsIn S, MovieStar M

WHERE S.StarID = M.StarID

Is this totally correct?

StarID	Name	Gender
1	Harrison Ford	Male
2	Vivian Leigh	Female
3	Judy Garland	Female

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Scarlett O'Hara
3	3	Dorothy Gale

Harrison Ford will appear twice

# Join Example

---

- Find the names of all movie stars who have been in a movie

```
SELECT Name  
FROM StarsIn S, MovieStar M  
WHERE S.StarID = M.StarID
```

Is this totally correct?

```
SELECT DISTINCT Name  
FROM StarsIn S, MovieStar M  
WHERE S.StarID = M.StarID
```

What if two movie stars  
had the same name?

- What if I run the following query?

```
SELECT DISTINCT StarID, Name  
FROM StarsIn S, MovieStar M  
WHERE S.StarID = M.StarID
```

Error: Column StarID  
is ambiguous

# Select Project Join example

---

- What are all the titles of movies with female actors?
- Write in Relational Algebra and SQL

# Select Project Join example

---

- What are all the titles of movies with female actors?
- Write in Relational Algebra and SQL

Relational algebra:

$\pi_{\text{title}}(\sigma_{\text{gender} = \text{'female'}}(\text{Movie} \bowtie \text{StarsIn} \bowtie \text{MovieStar}))$

# Select Project Join example

---

- What are all the titles of movies with female actors?
- Write in Relational Algebra and SQL

Relational algebra:

$\pi_{\text{title}}(\sigma_{\text{gender} = \text{'female'}}(\text{Movie} \bowtie \text{StarsIn} \bowtie \text{MovieStar}))$

SQL

```
SELECT DISTINCT Title
FROM Movie m, StarsIn s, MovieStar st
WHERE m.MovieID = s.MovieID and s.StarID =
st.StarID and gender = 'female'
```

# Renaming Attributes in Result

---

- SQL allows renaming relations and attributes using the **as** clause:  
*old-name as new-name*
- Example: Find the title of movies and the IDs of all actors in them, and rename “StarID” to “ID”

```
SELECT  Title, StarID AS ID
FROM    StarsIn S, Movie M
WHERE   M.MovieID = S.MovieID
```

# Congratulations:

## You know select-project-join queries

---

- Very common subset to talk about
  - You saw it in the RA tutorial
- Can do many (but not all) useful things

SQL is *declarative*, not procedural  
how do we know? Lets see what  
procedural would look like...



# Conceptual Procedural Evaluation Strategy

---

1. Compute the cross-product of *relation-list*.
2. Discard resulting tuples if they fail *qualifications*.
3. Delete attributes that are not in *target-list*.
4. If DISTINCT is specified, eliminate duplicate rows.

# Example of Conceptual Procedural Evaluation

```
SELECT Name
FROM MovieStar M, StarsIn S
WHERE S.StarID = M.StarID AND MovieID = 276
```

join

selection

MovieStar X StarsIn

(StarID)	Name	Gender	MovieID	(StarID)	Character
1273	Nathalie Portman	Female	272	1269	Leigh Anne Touhy
1273	Nathalie Portman	Female	273	1270	Mary
1273	Nathalie Portman	Female	274	1271	King George VI
1273	Nathalie Portman	Female	276	1273	Nina Sayers
...	...	...	...	...	...