

CPSC 304 – Administrative notes

October 25 & October 29, 2024

- Midterm is back!
 - See Piazza – in general you all did quite well
- Project:
 - Milestone 3: Project check in – due October 25
 - Sign up now!
 - Milestone 4: Project implementation – due November 29
 - Milestone 5: Group demo – week of December 2
 - Milestone 6: Individual Assessment – Due November 29
- Tutorials: basically project group work time
- Final exam: December 16 at 12pm!
 - Look out for conflict form to come soon!

Now where were we...

- SQL!
- When in doubt, start with
SELECT
FROM
WHERE

Set Operations: EXCEPT/MINUS

- Find the sids of all students who took Operating System Design but did not take Database Systems

```
SELECT snum
FROM enrolled
WHERE cname = 'Operating System Design'
EXCEPT ← Oracle uses MINUS rather than EXCEPT
SELECT snum
FROM enrolled
WHERE cname = 'Database Systems'
```

Can we do it in a different way?
(We'll come back to this)

But what about...

- Select the IDs of all students who have not taken “Operating System Design”
 - One way to do is to find all students that taken “Operating System Design”.
 - Do all students MINUS those who have taken “Operating System Design”

```
SELECT snum
FROM student
EXCEPT ← Oracle uses MINUS rather than EXCEPT
SELECT snum
FROM enrolled
WHERE cname = 'Operating System Design'
```

Motivating Example for Nested Queries

Find ids and names of male stars who have been in movie with ID 248:

Motivating Example for Nested Queries

Find ids and names of male stars who have been in movie with ID 248:

```
SELECT M.StarID, name
FROM   MovieStar M, StarsIn S
WHERE  M.StarID = S.starID AND S.MovieID = 248
AND gender = 'male';
```

- *Find ids and names of male stars who have been in some movie but not been in movie w/ ID 248 w/o using EXCEPT/MINUS:*

Would the following be correct?

```
SELECT DISTINCT M.StarID, name
FROM   MovieStar M, StarsIn S
WHERE  M.StarID = S.starID AND S.MovieID <> 248
and gender = 'male';
```

Nested Queries

- A very powerful feature of SQL:

```
SELECT   $A_1, A_2, \dots, A_n$   
FROM     $R_1, R_2, \dots, R_m$   
WHERE   condition
```

- A nested query is a query that has another query embedded with it.
 - A **SELECT**, **FROM**, **WHERE**, or **HAVING** clause can itself contain an SQL query!
 - Being part of the **WHERE** clause is the most common

Nested Queries (IN/Not IN)

Find ids and names of male stars who have been in movie with ID 248:

Nested Queries (IN/Not IN)

Find ids and names of male stars who have been in movie with ID 248:

```
SELECT DISTINCT M.StarID, M.Name
FROM   MovieStar M
WHERE  M.Gender = 'male' AND
      M.StarID IN (SELECT S.StarID
                  FROM   StarsIn S
                  WHERE  MovieID=248)
```

There's also NOT IN

- To find stars who have *not* been in movie 248, use **NOT IN**.
- To understand nested query semantics, think of a nested loops evaluation:
 - For each MovieStar tuple, check the qualification by computing the subquery.

Nested Queries (IN/Not IN)

Find ids and names of male stars who have been in movie with ID 248:

```
SELECT DISTINCT M.StarID, M.Name
FROM   MovieStar M
WHERE  M.Gender = 'male' AND
       M.StarID IN (SELECT S.StarID
                    FROM   StarsIn S
                    WHERE  MovieID=28)
```

- In this example the inner query does not depend on the outer query so it could be computed just once.
- Think of this as a function that has no parameters.

```
SELECT S.StarID
FROM   StarsIn S
WHERE  MovieID=248
```

StarID
1245
1246

```
SELECT M.StarID, M.Name
FROM   MovieStar M
WHERE  M.Gender = 'male' AND
       M.StarID IN
       (1245,1246)
```

Rewriting EXCEPT Queries Using In

- Using nested queries, find the snums of all students who took Operating System Design but did not take Database Systems

```
Student(snum,sname,major,standing,age)
Class(name,meets_at,room,fid)
Enrolled(snum,cname)
Faculty(fid,fname,deptid)
```

Rewriting EXCEPT Queries Using In

- Using nested queries, find the snums of all students who took Operating System Design but did not take Database Systems

```
SELECT snum
FROM enrolled
WHERE cname = 'Operating System Design' and snum not in
      (SELECT snum
       FROM enrolled
       WHERE cname = 'Database Systems')
```

Rewriting INTERSECT Queries Using IN

Find IDs of stars who have been in movies in 1944 and 1974

Rewriting INTERSECT Queries Using IN

Find IDs of stars who have been in movies in 1944 and 1974

```
SELECT DISTINCT S.StarID
FROM   Movie M, StarsIn S
WHERE  M.MovieID = S.MovieID AND M.year = 1944 AND
       S.StarID IN (SELECT S2.StarID
                    FROM Movie M2, StarsIn S2
                    WHERE  M2.MovieID = S2.MovieID AND M2.year = 1974)
```

The subquery finds stars who have been in movies in 1974

Nested Queries with Correlation

Same idea, subtle difference

Find names of stars who have been in movie w/ ID 248:

```
SELECT DISTINCT M.Name
FROM   MovieStar M
WHERE  EXISTS (SELECT *
                FROM   StarsIn S
                WHERE  MovieID=248 AND S.StarID = M.StarID)
```

- **EXISTS:** *returns true if the set is not empty.*
- **UNIQUE:** *returns true if there are no duplicates.*
- Illustrates why, in general, subquery must be re-computed for each StarsIn tuple.

SQL EXISTS Condition

- The SQL EXISTS condition is used in combination with a subquery and is considered to be met, if the subquery returns at least one row. It can be used in a SELECT, INSERT, UPDATE, or DELETE statement.
- We can also use NOT EXISTS

SQL EXISTS Condition

- Using the EXISTS/ NOT EXISTS operations and correlated queries, find the name and age of the oldest student(s)

```
Student(snum,sname,major,standing,age)
Class(name,meets_at,room,fid)
Enrolled(snum,cname)
Faculty(fid,fname,deptid)
```

SQL EXISTS Condition

- Using the EXISTS/ NOT EXISTS operations and correlated queries, find the name and age of the oldest student(s)

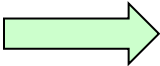
```
SELECT sname, age
FROM student s2
WHERE NOT EXISTS(SELECT *
                  FROM student s1
                  WHERE s1.age > s2.age)
```

SQL EXISTS ...

```
SELECT sname, age
FROM student s2
WHERE NOT EXISTS(SELECT *
                  FROM student s1
                  WHERE s1.age > s2.age)
```

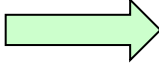


Does there exist a tuple in s1 such that the age of the s1 tuple is greater than the age of the tuple in s2?

Student s2



snum	Name	...
...
...
...

Student s1



snum	Name	...
...
...
...

More on Set-Comparison Operators

- We've already seen **IN** and **EXISTS**. Can also use **NOT IN**, **NOT EXISTS**.
- Also available: **op ANY**, **op ALL**, where **op** is one of: **>**, **<**, **=**, **<=**, **>=**, **<>**
- Find movies made after "Fargo"

More on Set-Comparison Operators

- We've already seen **IN** and **EXISTS**. Can also use **NOT IN**, **NOT EXISTS**.
- Also available: **op ANY**, **op ALL**, where **op** is one of: **>**, **<**, **=**, **<=**, **>=**, **<>**
- Find movies made after "Fargo"

```
SELECT *  
FROM Movie  
WHERE year > ANY (SELECT year  
                  FROM Movie  
                  WHERE Title = 'Fargo')
```

Just returning one column

If we have multiple movies named Fargo then we can use ALL instead of ANY

Clicker nested question

Determine the result of:

```
SELECT Team, Day
```

```
FROM Scores S1
```

```
WHERE Runs <= ALL
```

```
  (SELECT Runs
```

```
    FROM Scores S2
```

```
    WHERE S1.Day = S2.Day )
```

Which of the following is in the result:

- A. (Carp, Sun)
- B. (Bay Stars, Sun)
- C. (Swallows, Mon)
- D. All of the above
- E. None of the above

Scores:			
Team	Day	Opponent	Runs
Dragons	Sun	Swallows	4
Tigers	Sun	Bay Stars	9
Carp	Sun	Giants	2
Swallows	Sun	Dragons	7
Bay Stars	Sun	Tigers	2
Giants	Sun	Carp	4
Dragons	Mon	Carp	6
Tigers	Mon	Bay Stars	5
Carp	Mon	Dragons	3
Swallows	Mon	Giants	0
Bay Stars	Mon	Tigers	7
Giants	Mon	Swallows	5

Clicker nested question

Clickernested.sql

Determine the result of:

```
SELECT Team, Day
FROM Scores S1
WHERE Runs <= ALL
  (SELECT Runs
   FROM Scores S2
   WHERE S1.Day = S2.Day )
```

Which of the following is in the result:

- A. (Carp, Sun)
- B. (Bay Stars, Sun)
- C. (Swallows, Mon)
- D. All of the above
- E. None of the above

Correct

Scores:			
Team	Day	Opponent	Runs
Dragons	Sun	Swallows	4
Tigers	Sun	Bay Stars	9
Carp	Sun	Giants	2
Swallows	Sun	Dragons	7
Bay Stars	Sun	Tigers	2
Giants	Sun	Carp	4
Dragons	Mon	Carp	6
Tigers	Mon	Bay Stars	5
Carp	Mon	Dragons	3
Swallows	Mon	Giants	0
Bay Stars	Mon	Tigers	7
Giants	Mon	Swallows	5

Team/Day pairs such that the team scored the minimum number of runs for that day.

Example

- Using the any or all operations, find the name and age of the oldest student(s)

Example

- Using the any or all operations, find the name and age of the oldest student(s)

```
SELECT sname, age
FROM student s2
WHERE NOT EXISTS(SELECT *
                  FROM student s1
                  WHERE s1.age > s2.age)
```

```
SELECT sname, age
FROM student s2
WHERE s2.age >= all (SELECT age
                    FROM student s1)
```

You can rewrite queries that use **any** or **all** with queries that use **exist** or **not exist**

Clicker Question

- Consider the following SQL query

```
SELECT DISTINCT s1.sname, s1.age  
FROM student s1, student s2  
WHERE s1.age > s2.age
```

- This query returns
 - A: The name and age of one of the oldest student(s)
 - B: The name and age of all of the oldest student(s)
 - C: The name and age of all of the youngest student(s)
 - D: The name and age of all students that are older than the youngest student(s)
 - E: None of the above

Clicker Question

- Consider the following SQL query

```
SELECT DISTINCT s1.sname, s1.age  
FROM student s1, student s2  
WHERE s1.age > s2.age
```

- This query returns
 - A: The name and age of one of the oldest student(s)
 - B: The name and age of all of the oldest student(s)
 - C: The name and age of all of the youngest student(s)
 - D: The name and age of all students who are older than the youngest student(s)
 - E: None of the above

Reminder: Division A/B

A

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

B1

pno
p2

A/B1

sno
s1
s2
s3
s4

B2

pno
p2
p4

A/B2

sno
s1
s4

B3

pno
p1
p2
p4

A/B3

sno
s1

Division in SQL

Find students who've
taken all classes.

(method 1)

```
SELECT sname
FROM Student S
WHERE NOT EXISTS
    ((SELECT C.name
      FROM Class C)
    EXCEPT
    (SELECT E.cname
      FROM Enrolled E
      WHERE E.snum=S.snum))
```

All classes
Classes
taken by S

The hard way (without EXCEPT): (method 2)

```
SELECT sname
FROM Student S
WHERE NOT EXISTS (SELECT C.name
                  FROM Class C
                  WHERE NOT EXISTS (SELECT E.snum
                                    FROM Enrolled E
                                    WHERE C.name=E.cname
                                    AND E.snum=S.snum))
```

Method 2:

*select Student S such that ...
there is no Class C...*

which is not taken by S

Division in SQL

(method 1- use EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE NOT EXISTS
      ((SELECT C.name    All classes
        FROM   Class C)
      EXCEPT
      (SELECT E.cname
        FROM   Enrolled E
        WHERE E.snum=S.snum))
```

Division in SQL

(method 1- use EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE NOT EXISTS
      ((SELECT C.name
        FROM   Class C)
      EXCEPT
      (SELECT E.cname
       FROM   Enrolled E
       WHERE E.snum=S.snum))
```

All classes
taken by S

Division in SQL

(method 1- use EXCEPT)

Find students who've taken all classes.

```
SELECT sname  
FROM Student S  
WHERE NOT EXISTS
```

```
    ((SELECT C.name  
       FROM Class C)  
     EXCEPT  
     (SELECT E.cname  
       FROM Enrolled E  
       WHERE E.snum=S.snum))
```

All classes
that have
not been
taken by S

Division in SQL

(method 1- use EXCEPT)

Find students who've taken all classes.

```
SELECT sname  
FROM Student S  
WHERE NOT EXISTS
```

```
    ((SELECT C.name  
       FROM Class C)  
     EXCEPT  
     (SELECT E.cname  
       FROM Enrolled E  
       WHERE E.snum=S.snum))
```

Only true if
there is no
class that has
not been
taken by S
(i.e., S must
have taken all
the classes)

Division in SQL

(method 2- without using EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE NOT EXISTS (SELECT C.name
                  FROM   Class C
                  WHERE  NOT EXISTS (SELECT E.snum
                                    FROM   Enrolled E
                                    WHERE  C.name=E.cname
                                    AND E.snum=S.snum))
```

Returns a result if student
S is enrolled in class C

Division in SQL

(method 2- without using EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE NOT EXISTS (SELECT C.name
                  FROM   Class C
                  WHERE  NOT EXISTS (SELECT E.snum
                                    FROM   Enrolled E
                                    WHERE  C.name=E.cname
                                    AND E.snum=S.snum))
```

Only true if student S has
never been enrolled in
class C.

Division in SQL

(method 2- without using EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE NOT EXISTS (SELECT C.name
                  FROM   Class C
                  WHERE  NOT EXISTS (SELECT E.snum
                                    FROM   Enrolled E
                                    WHERE  C.name=E.cname
                                    AND E.snum=S.snum))
```

Find the classes that student
S has not enrolled in.

Division in SQL

(method 2- without using EXCEPT)

Find students who've taken all classes.

```
SELECT sname
FROM   Student S
WHERE  NOT EXISTS (SELECT C.name
                   FROM   Class C
                   WHERE  NOT EXISTS (SELECT E.snum
                                     FROM   Enrolled E
                                     WHERE  C.name=E.cname
                                     AND E.snum=S.snum))
```

Only true if there is no class that student S has never been enrolled in (i.e., student S has been enrolled in all the classes).

You're Now Leaving the World of Relational Algebra

- You now have many ways of asking relational algebra queries
 - For this class, you should be able write queries using all of the different concepts that we've discussed & know the terms used
 - In general, use whatever seems easiest, unless the question specifically asks you to use a specific method.
 - Sometimes the query optimizer may do poorly, and you'll need to try a different version, but we'll ignore that for this class.