

CPSC 304

Introduction to Database Systems

Conceptual Database Design
The Entity-Relationship Model

Textbook Reference
Database Management Systems: Chapter 2

Databases: the continuing saga...



- We motivated that databases were great because they:
 - Store large amounts of data
 - Handle transactions
 - Allow efficient querying
 - *And many, many more classic favourites!*
- Before we can do all of these, we must design the database


Learning Goals



- Explain the purpose of an ER diagram, and list the major components.
- Given a problem description, create an ER diagram given a specification. Justify the decisions you make for entities, relationships, keys, key constraints, participation constraints, weak entities, is-a relationships, and aggregations.
- Given a problem description, identify alternative representations of the problem concepts and evaluate the choices
- Compare alternative ER models for the same domain and identify their strengths and weaknesses

Levels of Abstraction

- A major purpose of a DB system is to provide an abstract view of the data.
- Three abstraction levels:
 - **External (or View) level:** describes different part of the database to different users
 - convenience, security, etc.
 - Compare views of student, registrar, and database admin.
 - **Conceptual (or Logical) level:** how data is perceived by the users
 - **Physical level:** how data is actually stored (404) – covers things like indexes, bits on disk



Schema and Instances

- We create the conceptual **schema** – the logical structure of the database (e.g., students take courses)
- Later we'll populate **instances** – the content of the database at a particular point in time
- E.g., Gradebook schema is set, but currently there are no grades for CPSC 304


Conceptual Database Design

- What are the entities and relationships in the enterprise?
 - Entities are usually nouns
 - e.g., Students and Courses
 - Relationships are statements about 2 or more objects. Often, verbs.
 - e.g., an instructor Teaches a course
- What information about these entities and relationships should we store in the database?
- What integrity constraints or other rules hold?
- In relational databases, this is generally encoded in an **Entity-Relationship (ER) Diagram**

ER Model Basics: Entities



- **Entity**: Real-world object distinguishable from other objects.
- An entity is described using a set of **attributes**.




- **Entity Set**: A collection of similar entities.
E.g., all Movie People.
 - All entities in an entity set have the same set of attributes. (Until we consider ISA hierarchies, anyway!)
 - Each attribute has a **domain**. (e.g., *float*, *date*, *int*)
 - Each entity set has a **key**. The key is composed of all underlined attributes




Keys

- Distinguish entities
- A **key** is the **minimal** set of one or more attributes which, taken collectively, identify uniquely an entity in an entity set.
 - In Canada, ~50 addresses share the same postal code
- A **primary key** is the key chosen as the principal means to identify entities in an entity set
- The only keys shown in ER diagrams are primary keys (do not worry about this for now)
- We'll discuss superkeys when we consider normal forms (for now, don't worry about them)




ER Model Basics (Cont.)



- **Relationship**: Association among two or more entities.
 - E.g., Michelle Yeoh worked on Everything Everywhere All at Once.




ER Model Basics (Cont.)



- **Relationship Set:** Collection of similar relationships.
 - Collection of all MoviePeople that have worked in Movies.
- Same entity set could participate in different relationship sets, or in different “roles” in same set. (Kirk Douglas isParentOf Michael Douglas)



ER Model Basics (Cont.)



- A relationship must be uniquely defined by the entities involved (it may not be a key because it may not be minimal, as we'll see shortly)
- A relationship set may have **descriptive attributes** (like Role).
- An n-ary relationship set R relates n entity sets $E_1 \dots E_n$; each relationship in R involves entities $e_1 \in E_1, \dots, e_n \in E_n$
 - **Degree** or **arity**: # of entity sets in relationship (binary, ternary, etc.)

Cardinalities




- A **cardinality ratio** for a relationship set specifies the number of relationships in the set that an entity can participate in.

Let R be a relationship set between sets A and B .
 R can have 1 of 4 cardinalities:

1. **one-to-one** from A to B :

- an entity in A is associated with **at most one** entity in B and vice versa
- e.g. A : student, B : student ID #



Entity sets A and B

Cardinalities




- A **cardinality ratio** for a relationship set specifies the number of relationships in the set that an entity can participate in.

Let R be a relationship set between sets A and B .
 R can have 1 of 4 cardinalities:

1. **one-to-one** from A to B :

- an entity in A is associated with **at most one** entity in B and vice versa
- e.g. A : student, B : student ID #






Cardinalities (cont')

2. **one-to-many** from A to B:


- an entity in A is associated with any number of entities in B
- an entity in B is associated with at most one entity in A
- e.g. A: biological-mother, B: children



1-to Many

3. **many-to-one** from A to B: (switch A and B above)

- an entity in B is associated with any number of entities in A
- an entity in A is associated with at most one entity in B




Many-to-1

Cardinalities (cont')

4. **many-to-many** from A to B:

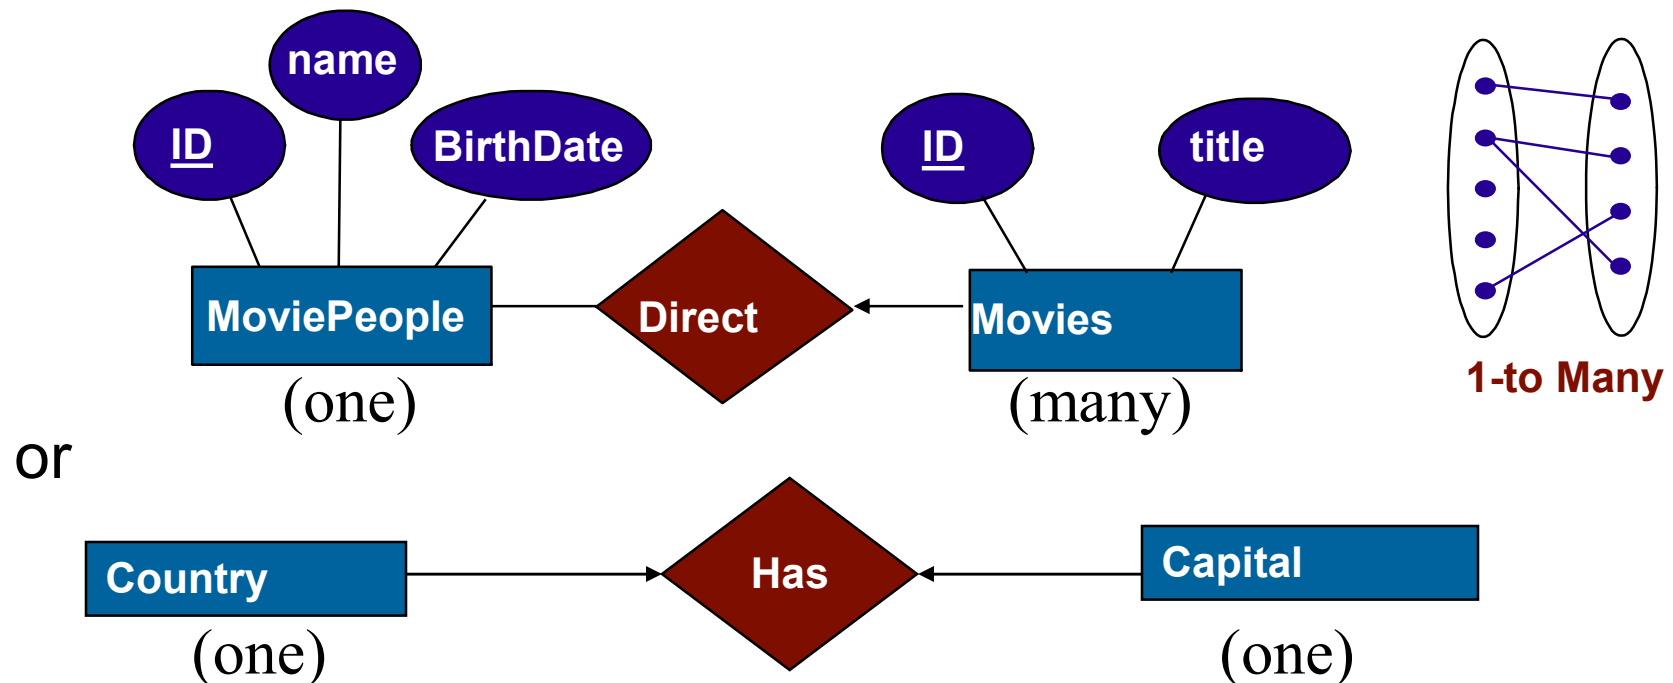
- an entity in A is associated with any number of entities in B and vice versa
- e.g. A: students, B: courses



Key Constraints




- The restriction imposed by a 1-to-1 and 1-to-many ratios are examples of **key constraints**.
- A key constraint is shown with an arrow in the ER diagram.
- Important on insertions



A brief digression on notation


The ER notation we use can be read: “if you know the entity with the arrow, then you know the relationship (and the other entities involved)”
– the arrow points to the thing there is only one of



James
Cameron



How can we uniquely identify a relationship?




How can we identify the role of a specific MoviePerson in a specific movie



Robert Pattinson
as
Bruce Wayne



How can we uniquely identify a relationship?




James
Cameron



The key of a many to one relationship is the key of the entity on the many side.

How can we uniquely identify a relationship?


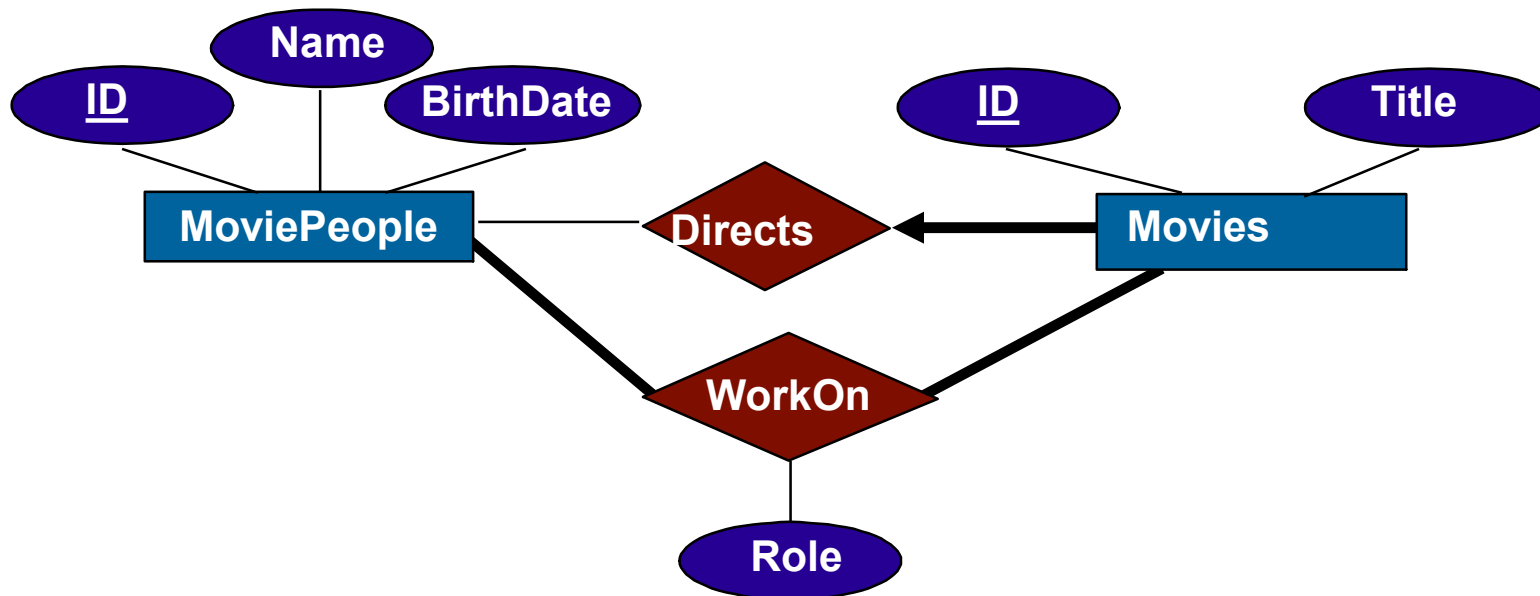


Ottawa

The key of a one-to-one relationship is the key of ONE of the entities.

Participation Constraints


- Participation : Indicates if all entities participate in the relationship.
- An entity's participation can be **total** or **partial**.
- Requiring total participation is a **participation constraint** and it is shown with a thick line
 - Important on deletions
 - i.e., participation of Movie in Directs is total (thick line)
 - Every movie must appear in some relationship in the Directs set




Remember?

Line types summarized

- Plain lines mean many to many:

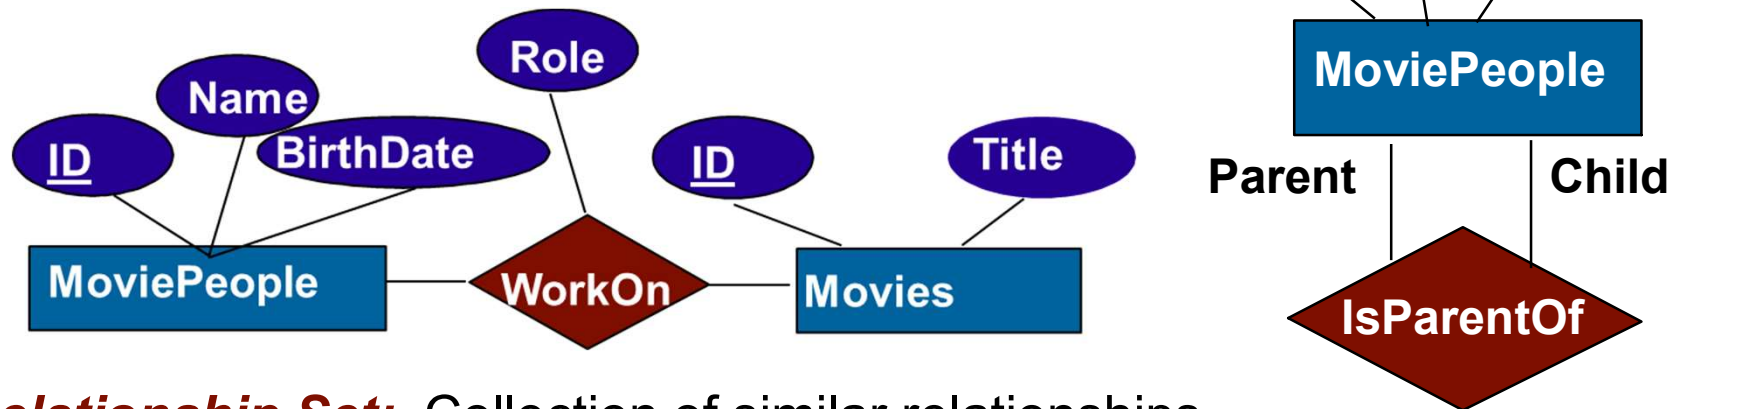


- Arrows mean the other side has a cardinality of one



- A thick line requires total participation and can be added to any line, arrow or not

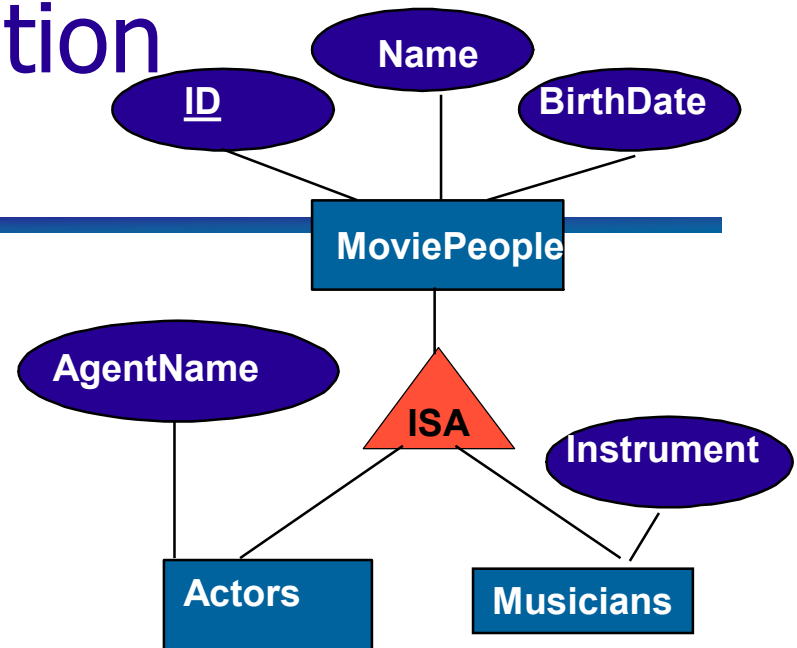
ER Model Basics (Cont.)



- **Relationship Set**: Collection of similar relationships.
 - Collection of all moviePeople that have worked in Movies.
- Same entity set could participate in different relationship sets, or in different “roles” in same set (“Parent”/”Child”)
- **A relationship must be uniquely defined by the entities involved** (it may not be a key because it may not be minimal, as we’ll see shortly)
- An n-ary relationship set R relates n entity sets $E_1 \dots E_n$; each relationship in R involves entities $e_1 \in E_1, \dots, e_n \in E_n$
 - **Degree** or **arity**: # of entity sets in relationship (binary, ternary, etc.)

Generalization/Specialization (ISA relationships)

- As in Java, or other PLs, attributes can be inherited.
- If we declare A **ISA** B, every A entity is a B entity.
- Reasons for using ISA:
 - To add descriptive attributes specific to a subclass.
 - To restrict entities that participate in a relationship.



There are some ISA constraints we can't express in ER diagrams

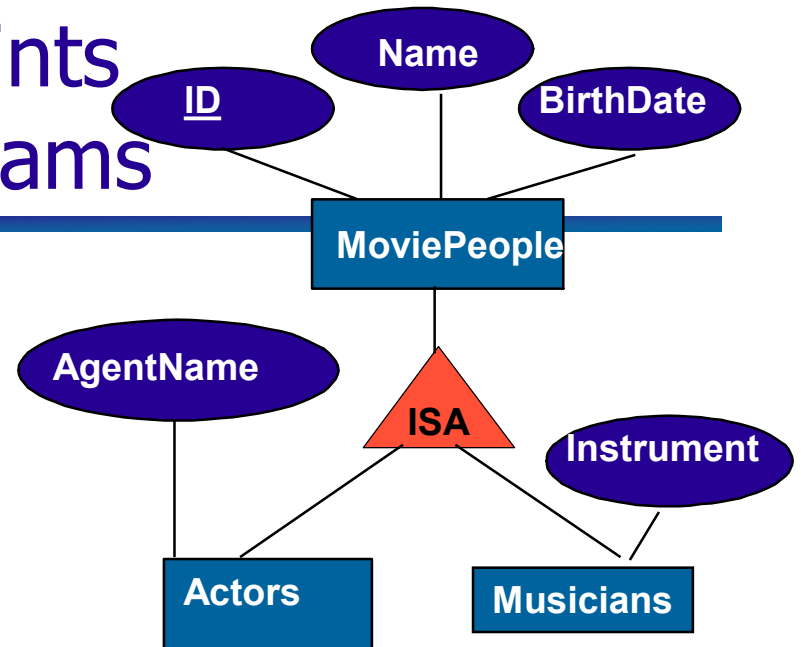
- **Overlap constraints :**

Specializations can be:

- **Disjoint :** a superclass entity belongs to no more than a single subclass
- **Overlapping :** subclasses may overlap

- **Covering constraints :** Specializations can be:


- **Total :** a superclass entity must belong to some subclass
- **Partial :** some superclass entity may not be in any subclass



We can represent these constraints by just writing them in.


ISA constraints Illustrated

Total + Overlap




Actor & Musician

Total + Disjoint

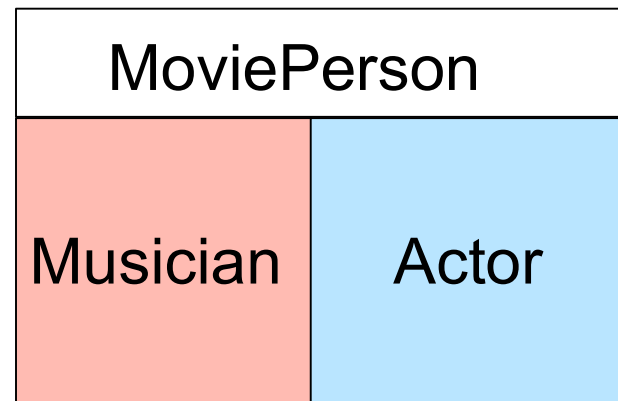


Partial + Overlap

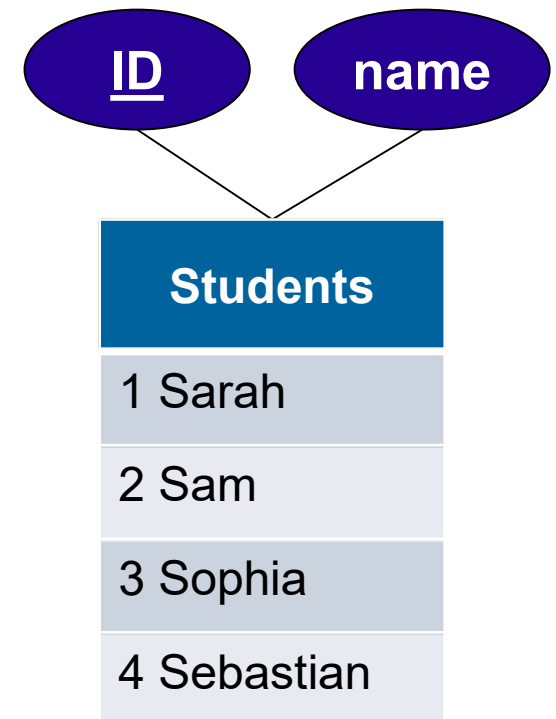
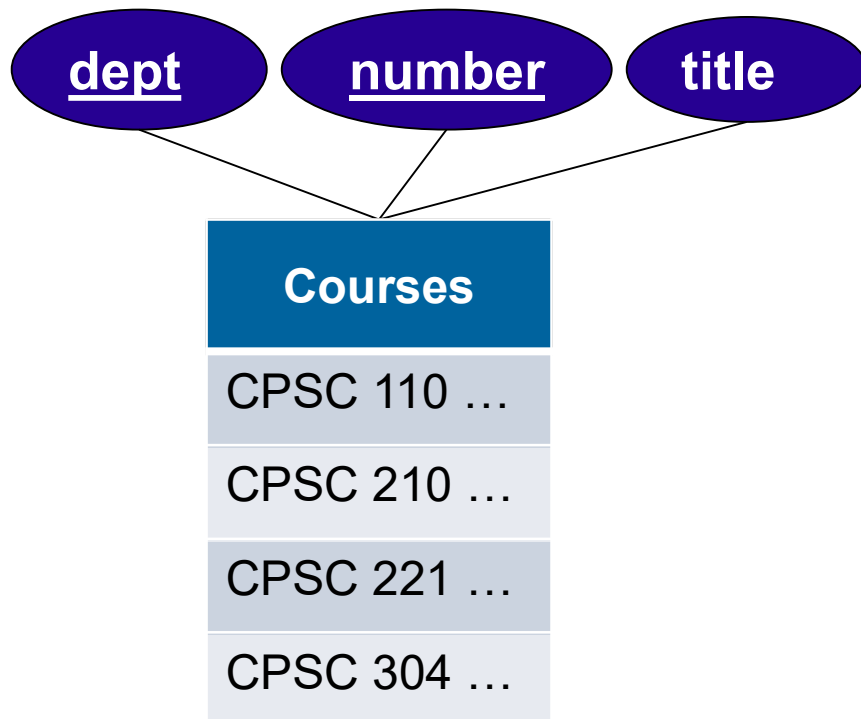


Actor & Musician

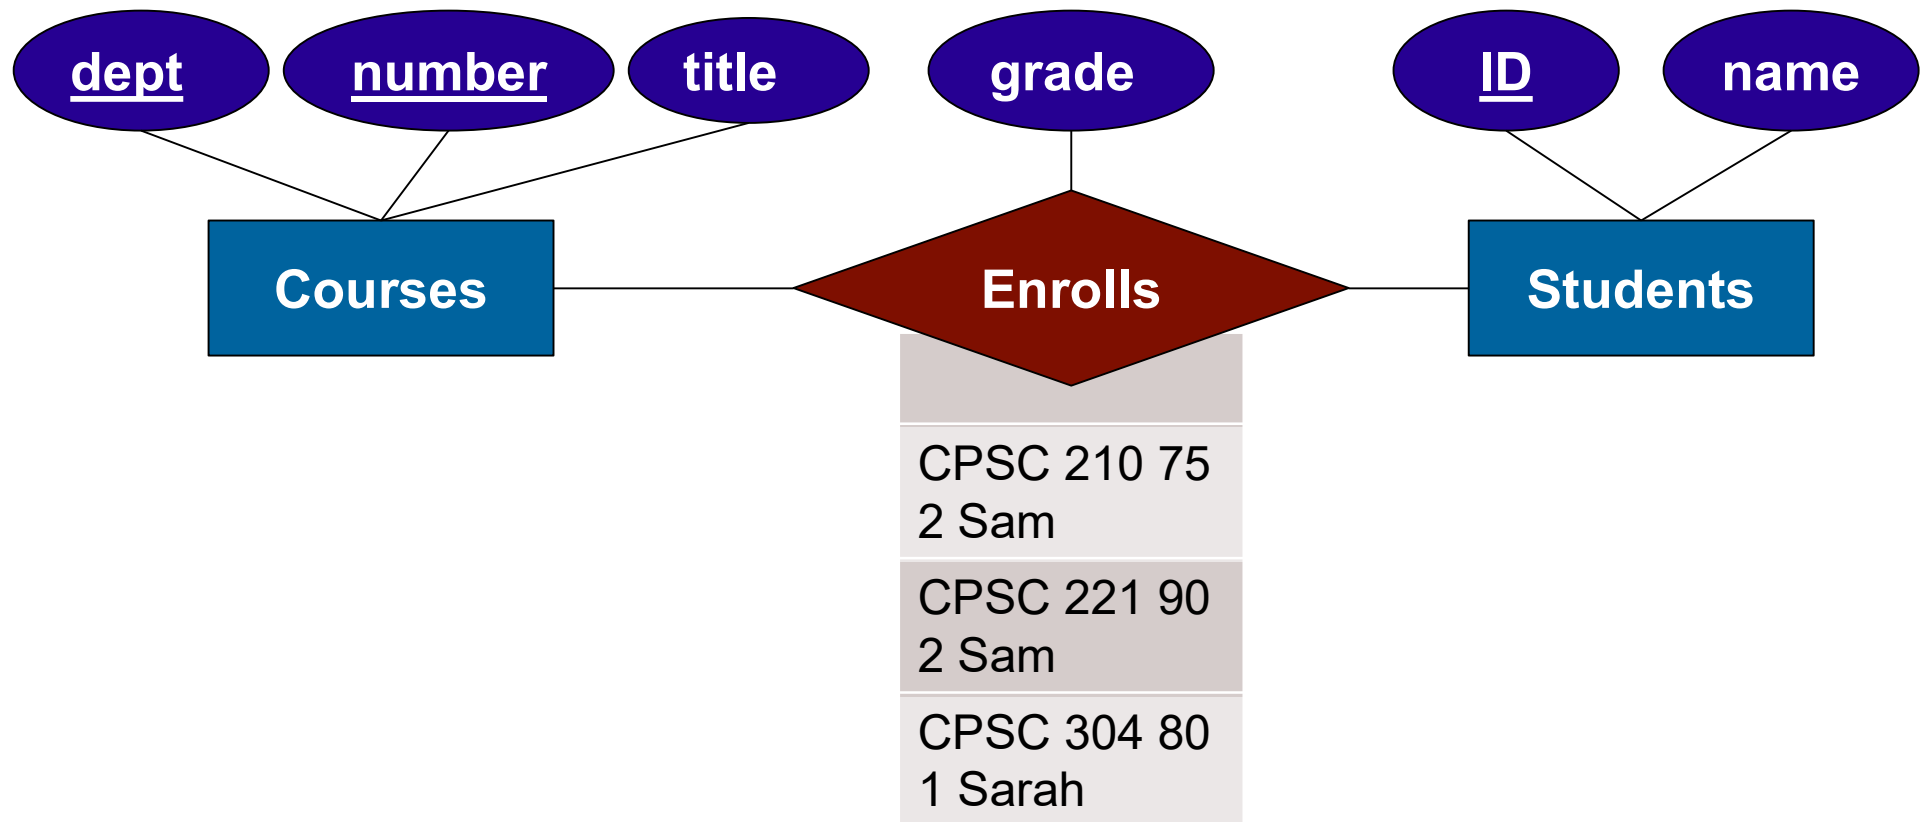
Partial + Disjoint



Review – Entity Sets



Review – Relationship Sets



Now where were we...

- We'd just finished covering the basics of ER diagrams, including:

- Entities

Students

- Attributes

ID

- Relationships

Have

- Key constraints

Students

Have

Vaccination
status

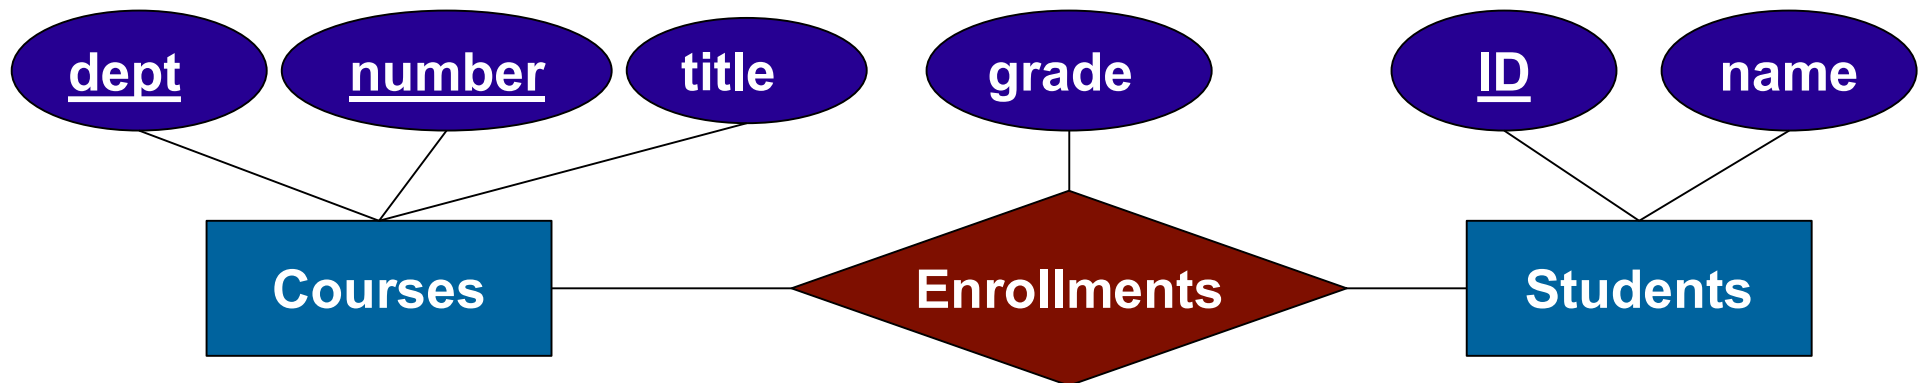
- Total participation

—

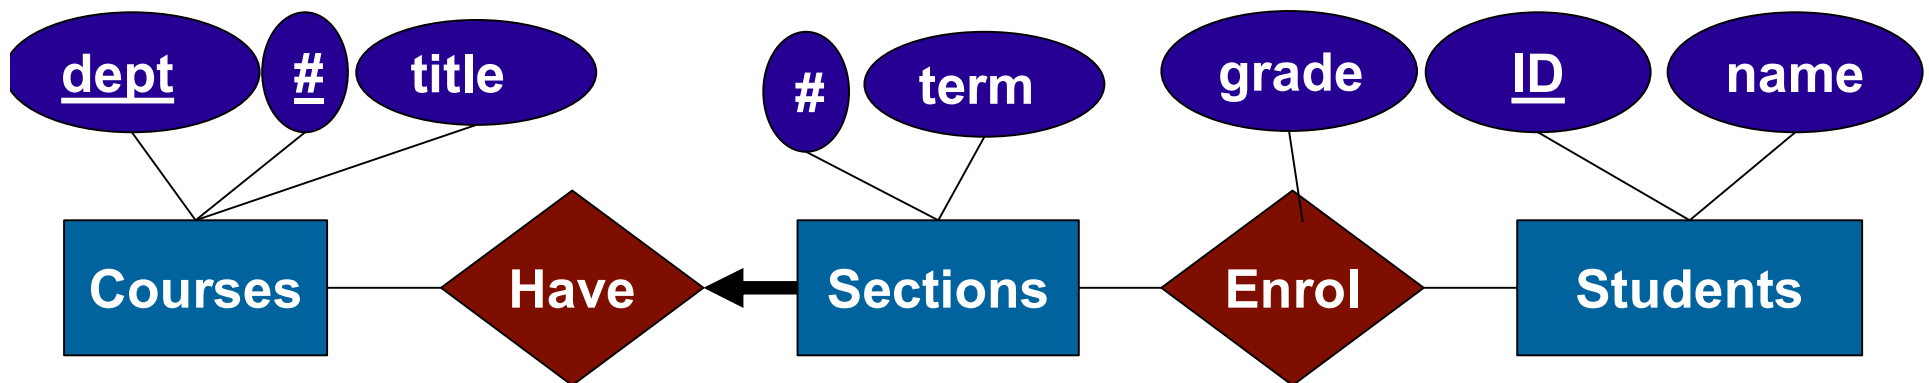
- Generalization/Specialization

ISA

Can we improve the design?

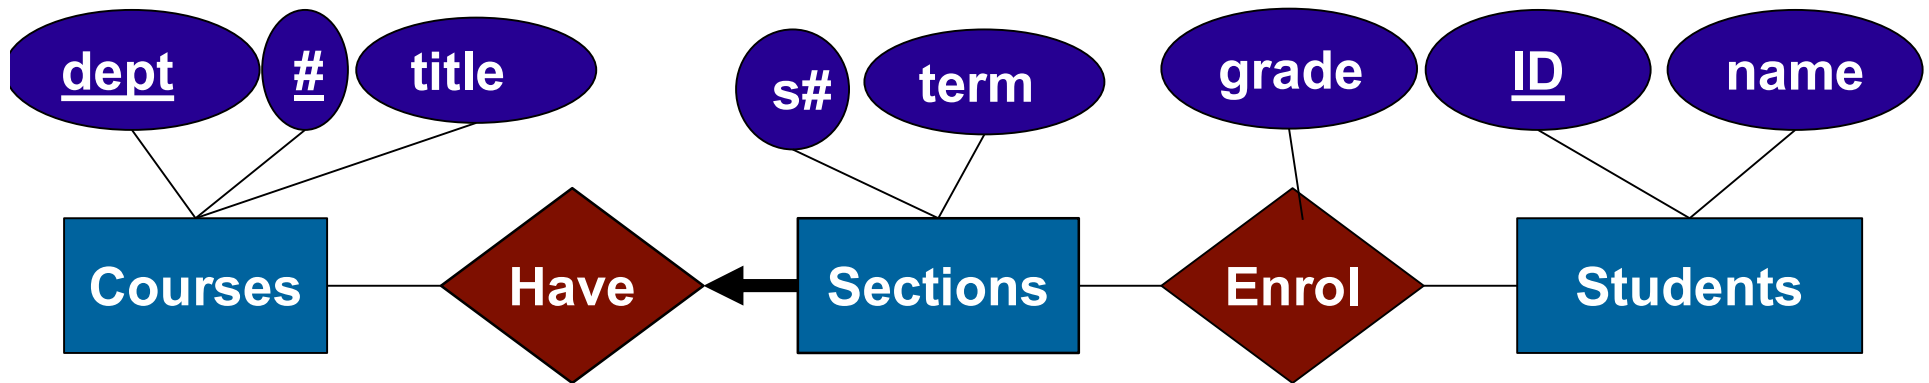


Here a student couldn't take a class more than once. Try:



But what should the key of Sections be?

A better solution would be if Sections could depend on Courses



The key has to be unique, right?


- Looking only at Sections (s# and term):
101 2021W1 is not unique
- Looking at Courses + Sections (dept # s# term):
CPSC 304 101 2021W1 is unique

We can do this with a *weak entity*


Weak Entities



- A **weak entity** can be identified uniquely only by additionally considering the key of another (*owner*) entity.
 - Think of this as a “belongs to” relationship.
- Owner entity set and weak entity set must participate in a **one-to-many** relationship set (one owner, many weak entities).
- Weak entity set must have **total** participation in this identifying relationship set.
- Weak entity sets and their identifying relationship sets are shown with **thick** lines.



A better solution would be if Sections could depend on Courses




{dept, #, s#, term} is the **key** for Have.

{s#, term} is a **partial key**, underlined with a dashed line.

Congratulations, you can now take CPSC 304 twice! (yay? 😊)

Aggregation

- Having a relationship between relationships is forbidden.

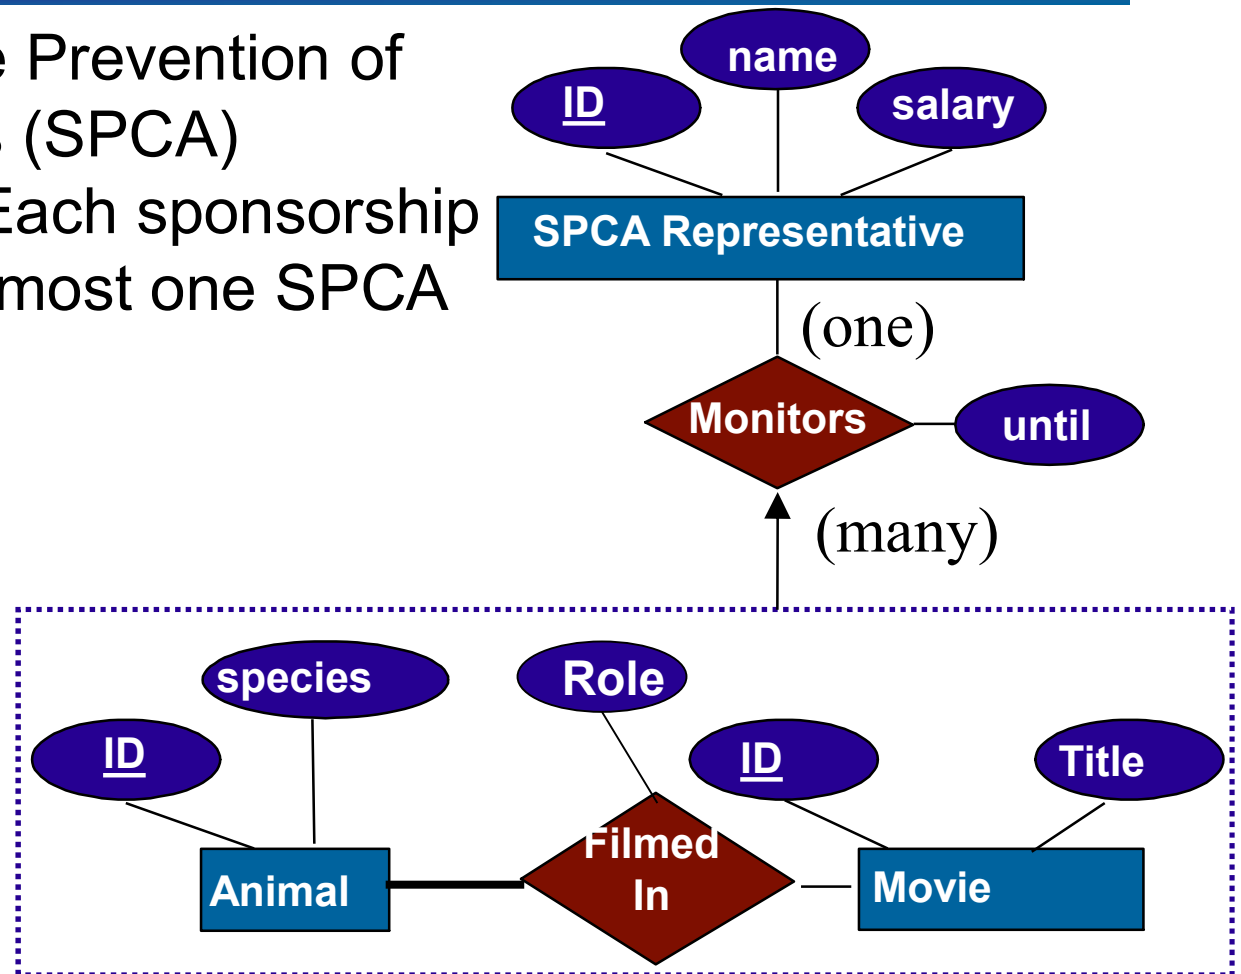


- Aggregation allows us to treat a relationship set as an entity set for purposes of participation in (other) relationships

Aggregation: getting around relationships between relationships




- The Society for the Prevention of Cruelty to Animals (SPCA) monitors movies. Each sponsorship is monitored by at most one SPCA representative



Aggregation: getting around relationships between relationships



- The Society for the Prevention of Cruelty to Animals (SPCA) monitors movies. Each sponsorship is monitored by at most one SPCA representative




- What is the key for FilmedIn?

Animal ID, movie ID





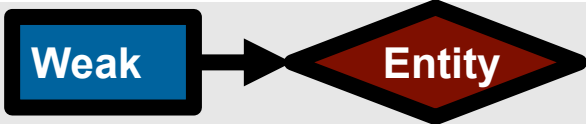
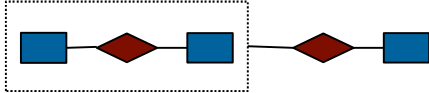

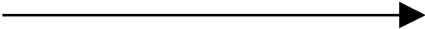
- What is the key for Monitors?

Animal ID, movie ID

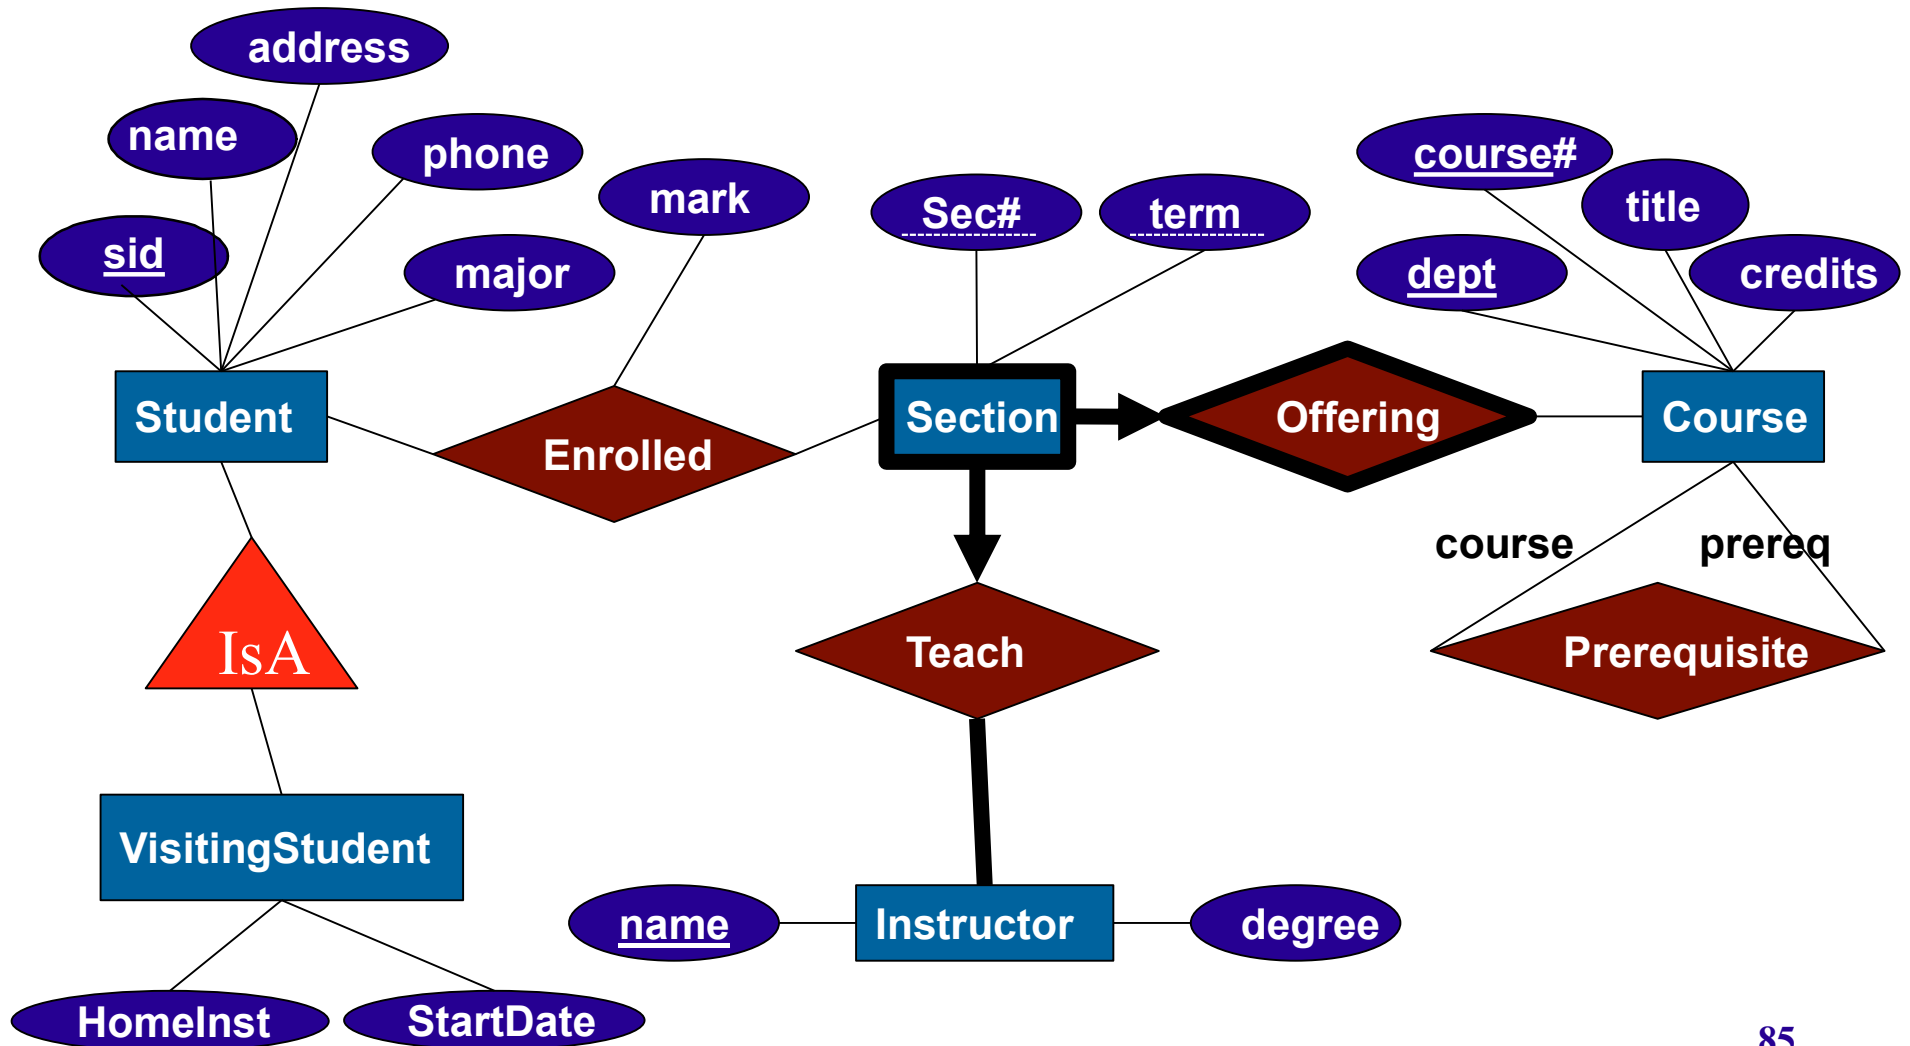


- This differs from a ternary relationship because monitors is its own relationship with a descriptive attribute

Summary


Name	Symbol
Entity	
Attribute	
Relationship	
Generalization/Specialization	
Weak Entity	
Aggregation	
Participation Constraints	
key constraints	

Sample solution



That's all there is to it

- Some ER models differ in expressiveness
- They model *most* concepts people want
- They don't model all of them, e.g.,
 - Functional dependencies – some attributes determine some other attributes, e.g., postal code determines (only) city and province



Conceptual Design Using the ER Model

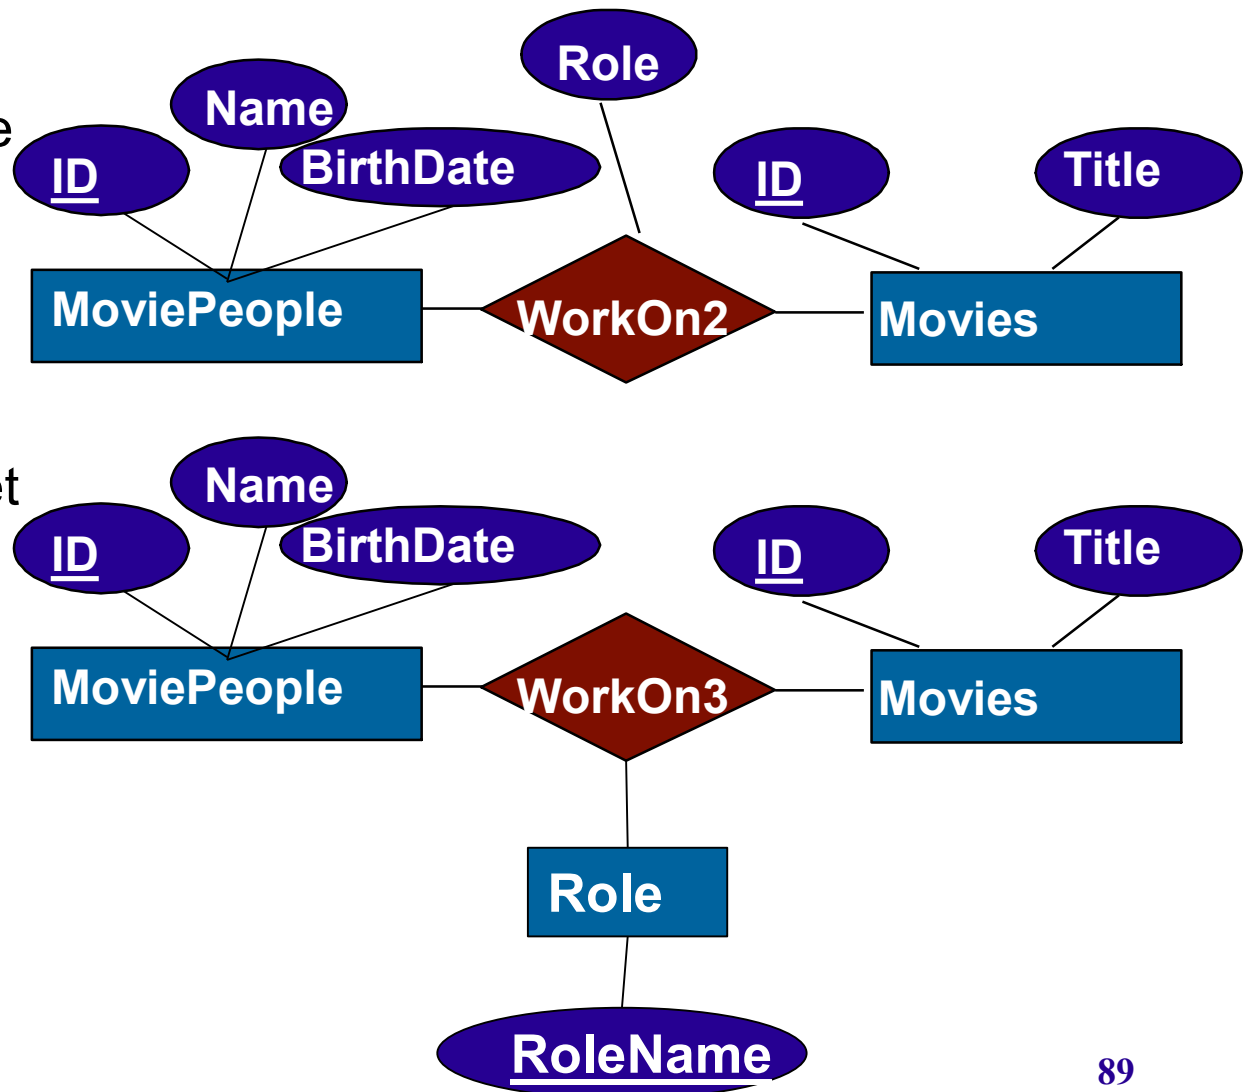
- Design choices:
 - Should a concept be modeled as an entity or an attribute?
 - Should a concept be modeled as an entity or a relationship?
 - Identifying relationships: Binary or ternary? Aggregation?
- Constraints in the ER Model:
 - A lot of data semantics can (and should) be captured.
 - But some constraints cannot be captured in ER diagrams.
 - i.e. domain constraints
 - dependencies

Entity vs. Attribute

- Should an *address* be an attribute of MoviePeople or an entity (connected to MoviePeople by a relationship)?
- Depends upon
 - the use we want to make of address information
 - the semantics of the data:
 - If we have several addresses per person, *address* must be an entity (since attributes cannot be set-valued).
 - If a person has only one street address, one city, one province, one postal code, etc. then these should simply be attributes.

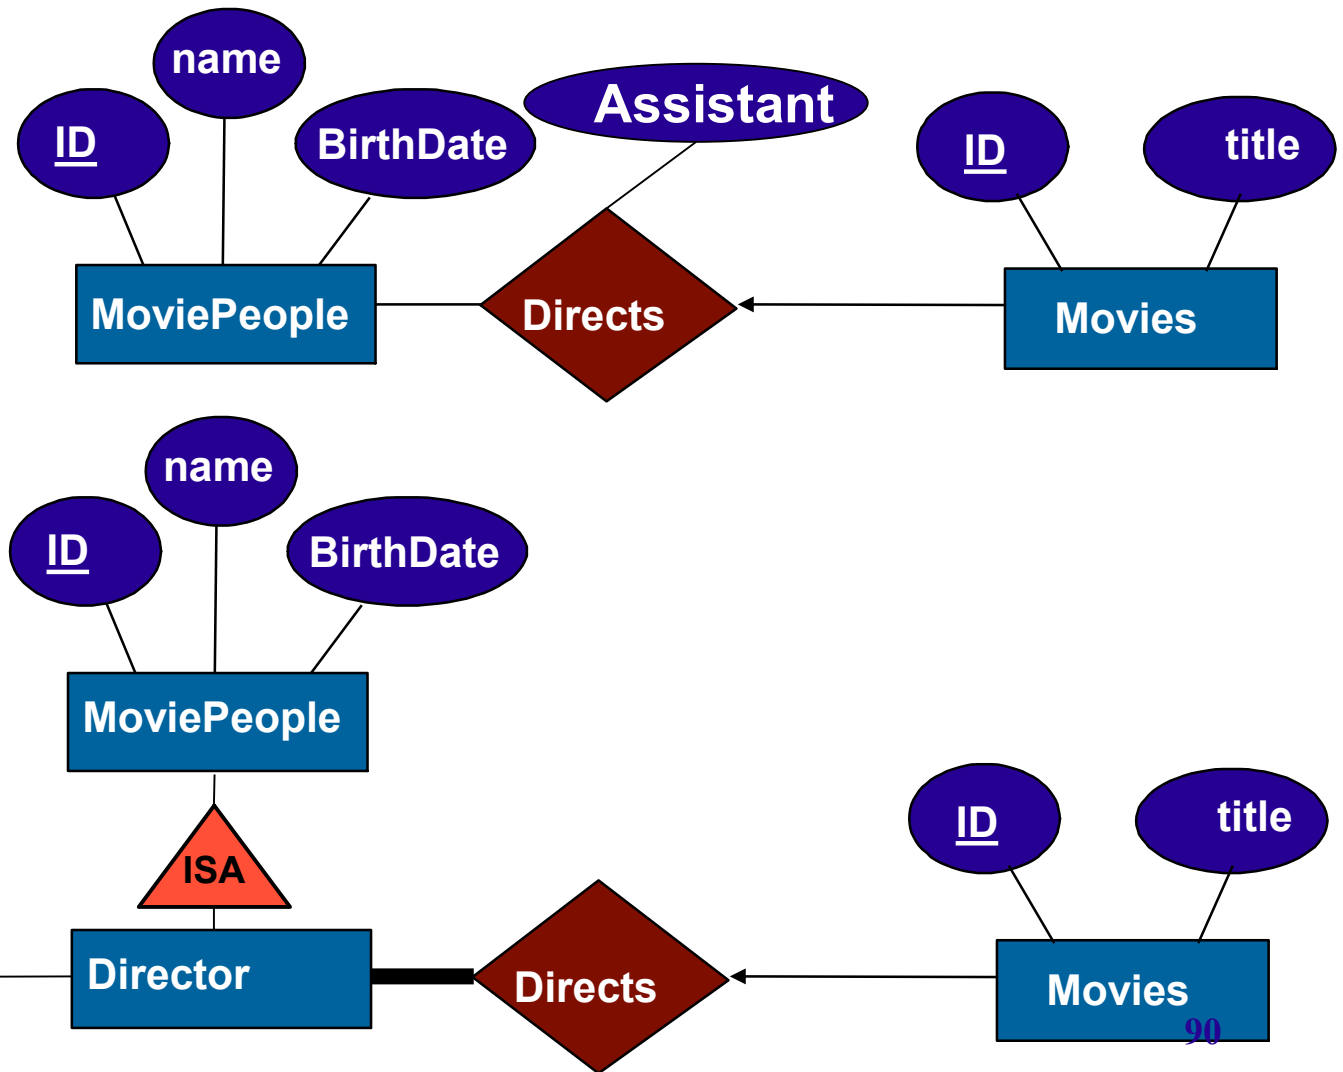
Entity vs. Attribute (Cont.)

- WorkOn2 does not allow a person to have more than one role in the same movie.
- We want to associate the same pair (MoviePerson, Movie) with more than one set of values for the descriptive attributes?
- Solution: change descriptive attributes into entities.



Entity vs. Relationship

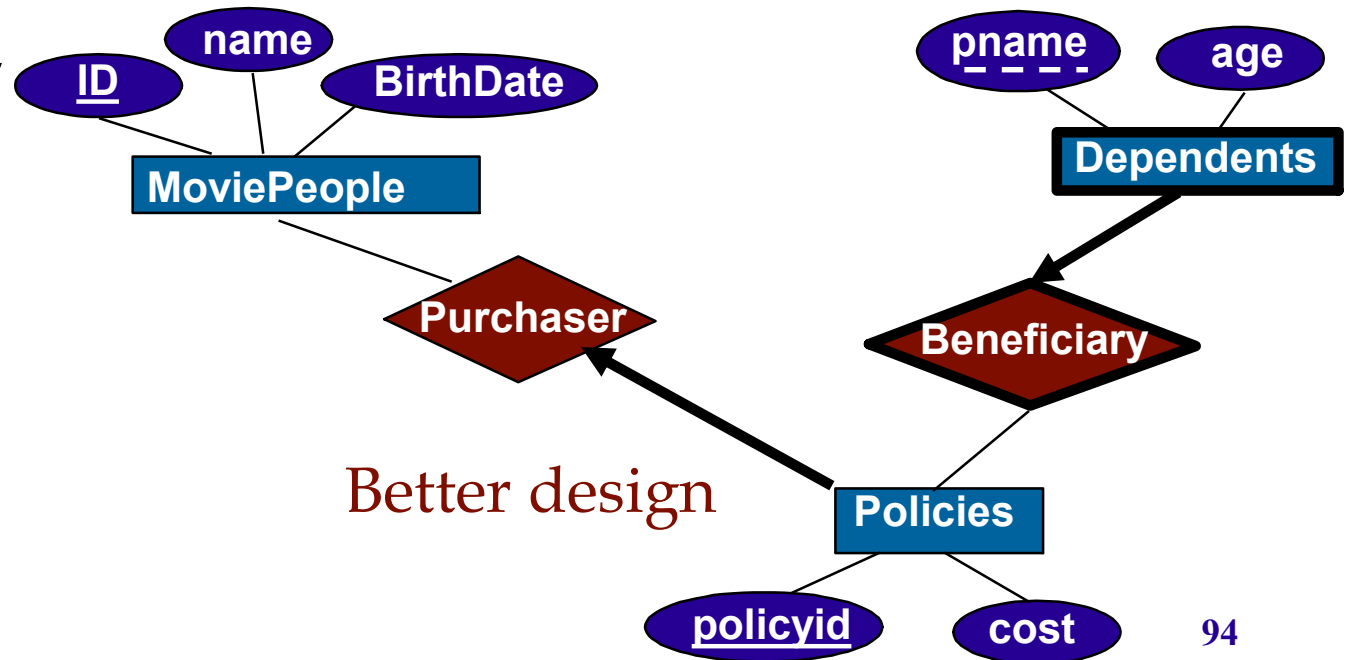
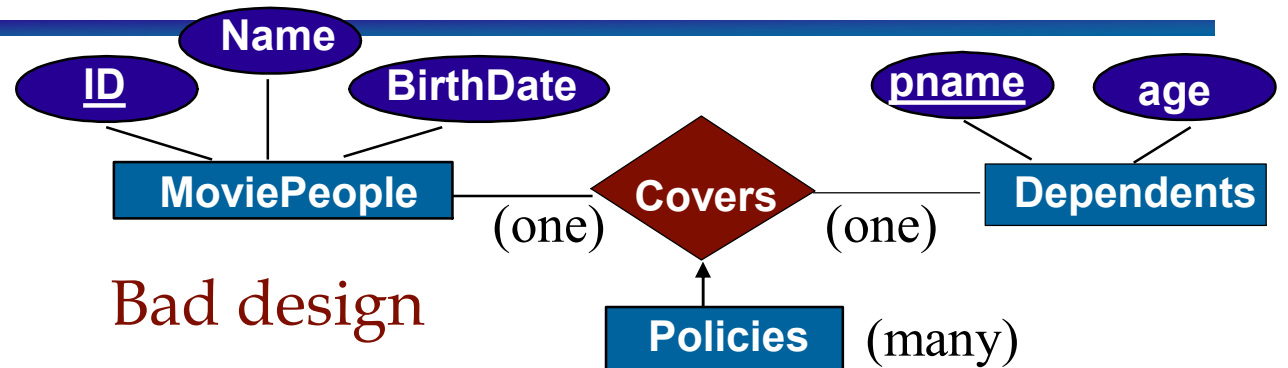
- How are the two ER models different?
- Director can get a separate assistant for each movie.
- All director must direct a movie and have the same assistant for *all* movies



Binary vs. Ternary Relationships

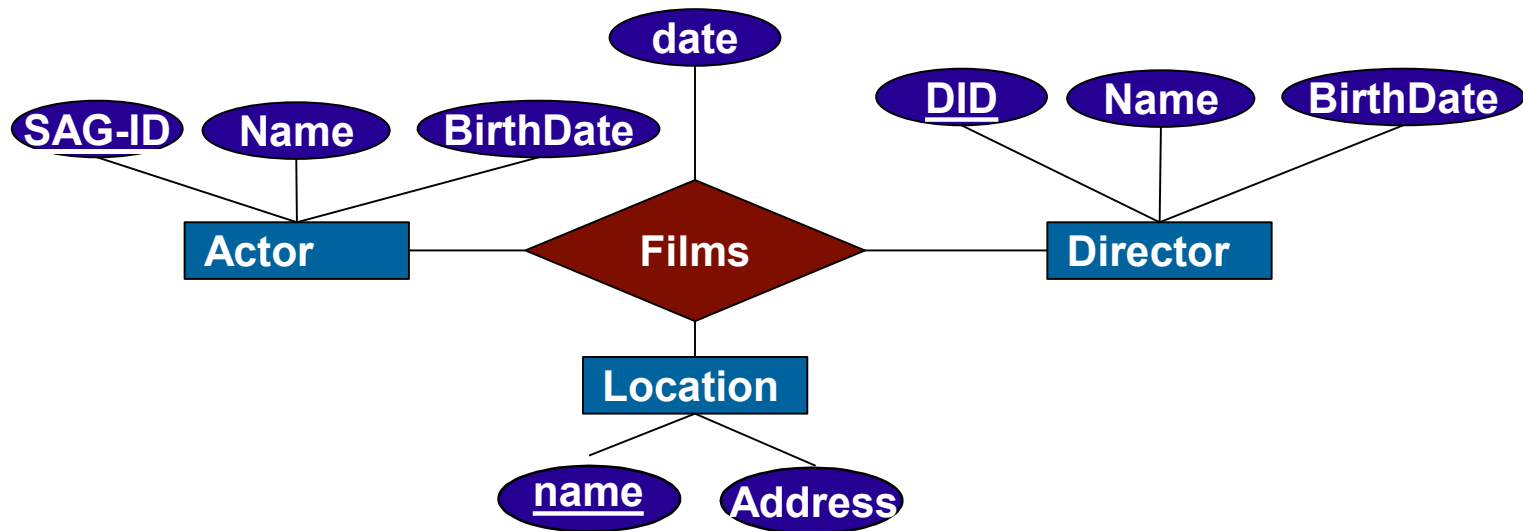
- If each policy is owned by just 1 person:

- Key constraint on Policies would mean policy can only cover 1 dependent!



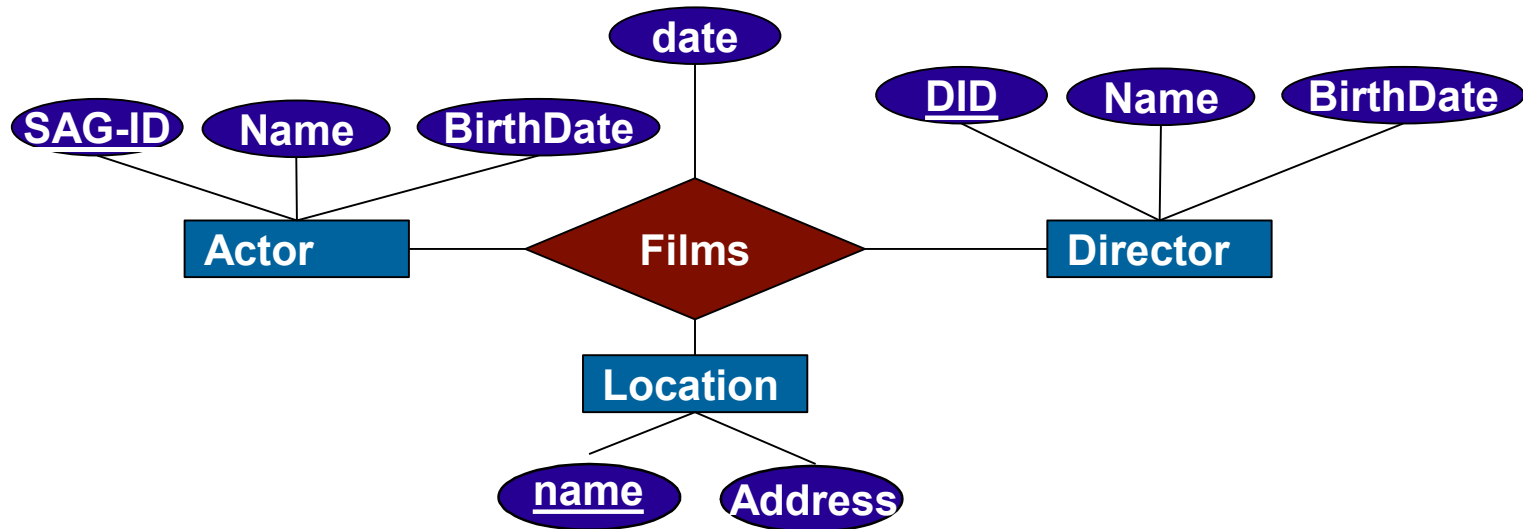
Binary vs. Ternary Relationships

- An example in the other direction: a ternary relationship **Films** relates entity sets **Actor**, **Director** and **Location**, and has descriptive attribute *date*.



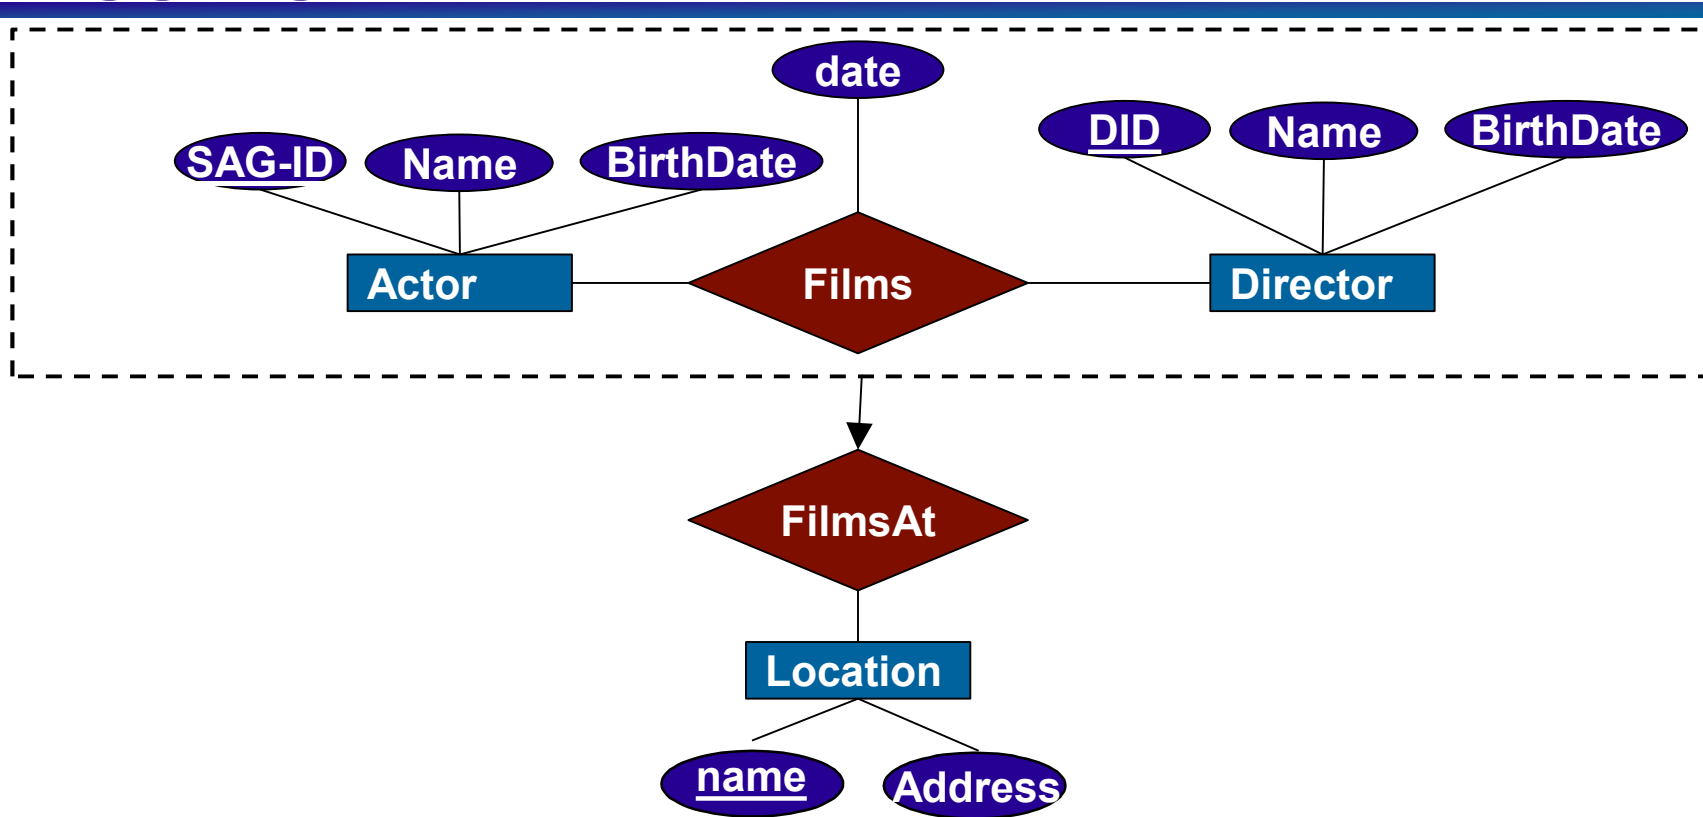
- Can we use two of binary relationships instead?

Binary vs. Ternary Relationships vs. Aggregation




- No combination of binary relationships is an adequate substitute:
 - D “directs” A, D “is filmed at ” L, and D “directs at” L does not imply that D directs A at L.
 - Also, how would we record *date*?

Binary vs. Ternary Relationships vs. Aggregation



- Aggregation can be used instead of a ternary relation if need to impose additional constraints:
 - I.e., If you know the actor and the director, you know the location they filmed at.

Binary vs. Ternary Relationships vs. Aggregation



- Aggregation can be used instead of a ternary relation if need to impose additional constraints:
 - I.e. A movie can only be filmed by one director. An actor can be directed by that director in any location.

Summary of Conceptual Design

- *Conceptual design* follows *requirements analysis*,
 - Yields a high-level description of data to be stored
- ER model popular for conceptual design
 - Constructs are expressive, close to the way people think about their applications.
- Basic constructs: *entities*, *relationships*, and *attributes* (of entities and relationships).
- Some additional constructs: *weak entities*, *ISA relationships*, and *aggregation*.
- Note: There are many variations on ER models.

Summary of ER (Cont.)

- Several kinds of integrity constraints can be expressed in the ER model: *key constraints*, *participation constraints*, and *overlap/covering constraints* for ISA relationships. Some *foreign key constraints* are also implicit in the definition of a relationship set.
 - Some constraints (notably, *functional dependencies*) cannot be expressed in the ER model.
 - Constraints play an important role in determining the best database design for an enterprise.

Summary of ER (Cont.)

- ER design is *subjective*. There are often many ways to model a given scenario! Analyzing alternatives can be tricky, especially for a large enterprise. Common choices include:
 - entity vs. attribute
 - entity vs. relationship
 - binary or n-ary relationship
 - whether or not to use ISA hierarchies
 - whether or not to use aggregation
- Ensuring good database design: resulting relational schema should be analyzed and refined further.

Learning Goals revisited



- Explain the purpose of an ER diagram, and list the major components.
- Given a problem description, create an ER diagram given a specification. Justify the decisions you make for entities, relationships, keys, key constraints, participation constraints, weak entities, is-a relationships, and aggregations.
- given a problem description, identify alternative representations of the problem concepts and evaluate the choices
- compare alternative ER models for the same domain and identify their strengths and weaknesses