

# CPSC 304

## Introduction to Database Systems

---

### Formal Relational Languages

Textbook Reference

Database Management Systems: 4 - 4.2  
(skip the calculii)

# Learning Goals

---



- Identify the basic operators in Relational Algebra (RA).
- Use RA to create queries that include combining RA operators.
- Given an RA query and table schemas and instances, compute the result of the query.

# Databases: the continuing saga

---



When last we left databases...

- We learned that they're excellent things
- We learned how to conceptually model them using ER diagrams
- We learned how to logically model them using relational schemas
- We knew how to normalize our database relations

We're almost ready to use SQL to query it, but first...


# Balance, Daniel-san, is key

---

The mathematical foundations:

- **Relational Algebra**

- Clear way of describing core concepts
- *partially procedural*: describe what you want and how you want it, but the order of operations matters



# Relational Query Languages


---

- Allow data manipulation and retrieval from a DB
- Relational model supports simple, powerful QLs:
  - Strong formal foundation based on logic
  - Allows for much optimization via *query optimizer*
- Query Languages **!=** Programming Languages
  - QLs not intended for complex calculations
  - QLs provide *easy access* to large datasets
  - Users *do not* need to know how to navigate through complicated data structures

# Relational Algebra (RA)

## All in one place

---



- Basic operations:
  - Selection ( $\sigma$ ): Selects a subset of rows from relation.
  - Projection ( $\pi$ ): Deletes unwanted columns from relation.
  - Cross-product ( $\times$ ): Allows us to combine two relations.
  - Set-difference ( $-$ ): Tuples in relation 1, but not in relation 2.
  - Union ( $\cup$ ): Tuples in relation 1 and in relation 2.
  - Rename ( $\rho$ ): Assigns a (another) name to a relation
- Additional, inessential but useful operations:
  - Intersection ( $\cap$ ), join ( $\bowtie$ ), division ( $/$ ), assignment ( $\leftarrow$ )
- All operators take one or two relations as inputs and give a new relation as a result
- For the purposes of relational algebra, relations are sets
- Operations can be **composed**. (Algebra is “closed”)

# Example Movies Database

---

Movie(MovieID, Title, Year)

StarsIn(MovieID, StarID, Character)

MovieStar(StarID, Name, Gender)

# Example Instances

Movie:

MovieID	Title	Year
1	Star Wars	1977
2	Casablanca	1942
3	The Wizard of Oz	1939
4	Indiana Jones and the Raiders of the Lost Ark	1981

StarsIn:

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Ilsa Lund
3	3	Dorothy Gale

MovieStar:

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female



# Selection ( $\sigma$ (sigma))

---

- Notation:  $\sigma_p(r)$
- $p$  is called the **selection predicate**
- ❖ Defined as:

$$\sigma_p(r) = \{t \mid t \in r \text{ and } p(t)\}$$

Where  $p$  is a formula in propositional calculus consisting of:

**connectives** :  $\wedge$  (**and**),  $\vee$  (**or**),  $\neg$  (**not**)

and

**predicates**:

$\langle \text{attribute} \rangle \text{ op } \langle \text{attribute} \rangle$  or

$\langle \text{attribute} \rangle \text{ op } \langle \text{constant} \rangle$

where  $\text{op}$  is one of:  $=, \neq, >, \geq, <, \leq$



Set of  
tuples of  $r$   
satisfying  $p$

# Selection Example

---

Movie:

MovieID	Title	Year
1	Star Wars	1977
2	Casablanca	1942
3	The Wizard of Oz	1939
4	Indiana Jones and the Raiders of the Lost Ark	1981

$\sigma_{\text{year} > 1945}(\text{Movie})$

MovieID	Title	Year
1	Star Wars	1977
4	Indiana Jones and the Raiders of the Lost Ark	1981

# Selection Example

---

Find all male stars from the MovieStar table.

StarID	Name	Gender
1	Harrison Ford	Male

$\sigma_{\text{Gender} = \text{'Male'}}$  MovieStar

# Projection ( $\pi$ (p*i*))

---

- Notation:

$$\pi_{A1, A2, \dots, Ak} (r)$$

where  $A1, \dots, Ak$  are attributes (the projection list) and  $r$  is a relation.

- The result: a relation of the  $k$  attributes  $A1, A2, \dots, Ak$  obtained from  $r$  by erasing the columns that are not listed
- Duplicate rows removed from result (relations are sets)

# Projection Examples

Movie:

$\pi_{\text{Title, Year}}(\text{Movie})$

MovieID	Title	Year	Title	Year
1	Star Wars	1977	Star Wars	1977
2	Casablanca	1942	Casablanca	1942
3	The Wizard of Oz	1939	The Wizard of Oz	1939
4	Indiana Jones and the Raiders of the Lost Ark	1981	Indiana Jones and the Raiders of the Lost Ark	1981

$\pi_{\text{Year}}(\text{Movie})$

Year
1977
1939
1942
1981

What is  $\pi_{\text{Title, Year}}(\sigma_{\text{year} > 1945}(\text{Movie}))$ ?

Title	Year
Star Wars	1977
Indiana Jones and the Raiders of the Lost Ark	1981

# Selection and Projection Example

---

Find the ids of movies made prior to 1950

Movie:

MovieID	Title	Year
1	Star Wars	1977
2	Casablanca	1942
3	The Wizard of Oz	1939
4	Indiana Jones and the Raiders of the Lost Ark	1981

MovieID
2
3

# Union, Intersection, Set-Difference

---

- Notation:  $r \cup s$      $r \cap s$      $r - s$

- Defined as:

$$r \cup s = \{t \mid t \in r \text{ or } t \in s\}$$

$$r \cap s = \{t \mid t \in r \text{ and } t \in s\}$$

$$r - s = \{t \mid t \in r \text{ and } t \notin s\}$$

- For these operations to be well-defined:

1.  $r, s$  must have the *same arity* (same number of attributes)
2. The attribute domains must be *compatible*  
(e.g., 2nd column of  $r$  has same domain of values as the 2nd column of  $s$ )

- What is the schema of the result?

# Union, Intersection, and Set Difference Examples

## MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

## Singer

StarID	SName	Gender
3	Judy Garland	Female
4	Sam Smith	Non-binary

## MovieStar $\cup$ Singer

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female
4	Sam Smith	Non-binary

## MovieStar $\cap$ Singer

StarID	Name	Gender
3	Judy Garland	Female

## MovieStar $-$ Singer

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female



# Set Operator Example

---

## MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

## Singer

StarID	Name	Gender
3	Judy Garland	Female
4	Sam Smith	Non-binary

Find the names of stars that are Singers but not MovieStars

Name
Sam Smith

# Cartesian (or Cross)-Product

---

- Notation:  **$r \times s$**

- Defined as:

$$r \times s = \{ t \ q \mid t \in r \textbf{ and } q \in s \}$$

- It is possible for  $r$  and  $s$  to have attributes with the same name, which creates a naming conflict.

- In this case, the attributes are referred to solely by position.

# Cartesian Product Example

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

StarsIn

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Ilsa Lund
3	3	Dorothy Gale

MovieStar x StarsIn

1	Name	Gender	MovieID	5	Character
1	Harrison Ford	Male	1	1	Han Solo
2	Ingrid Bergman	Female	1	1	Han Solo
3	Judy Garland	Female	1	1	Han Solo
1	Harrison Ford	Male	4	1	Indiana Jones
2	Ingrid Bergman	Female	4	1	Indiana Jones
3	Judy Garland	Female	4	1	Indiana Jones
...	...	...	...	...	...

# Rename ( $\rho$ (rho))

---

- Allows us to name results of relational-algebra expressions.
- Notation

$$\rho (X, E)$$

returns the expression  $E$  under the name  $X$

- We can rename part of an expression, e.g.,  
 $\rho((\mathbf{StarID} \rightarrow \mathbf{ID}), \pi_{StarID, Name}(MovieStar))$
- We can also refer to positions of attributes, e.g.,  
 $\rho((\mathbf{1} \rightarrow \mathbf{ID}), \pi_{StarID, Name}(MovieStar))$   
Is the same as above

# Rename ( $\rho$ (rho))

- We can rename the resulting relation and the attributes in that relation

$\rho(\text{GenderlessStars}(\text{ID}, \text{Nom}), \pi_{\text{StarID}, \text{Name}}(\text{MovieStar}))$

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

$\pi_{\text{StarID}, \text{Name}}(\text{MovieStar})$

StarID	Name
1	Harrison Ford
2	Ingrid Bergman
3	Judy Garland

GenderlessStars

ID	Nom
1	Harrison Ford
2	Ingrid Bergman
3	Judy Garland

# $\rho$ Example

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

StarsIn

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Ilsa Lund
3	3	Dorothy Gale

$\rho((1 \rightarrow \text{StarID1}, 5 \rightarrow \text{StarID2}), \text{MovieStar} \times \text{StarsIn})$

StarID1	Name	Gender	MovieID	StarID2	Character
1	Harrison Ford	Male	1	1	Han Solo
2	Ingrid Bergman	Female	1	1	Han Solo
3	Judy Garland	Female	1	1	Han Solo
1	Harrison Ford	Male	4	1	Indiana Jones
2	Ingrid Bergman	Female	4	1	Indiana Jones
3	Judy Garland	Female	4	1	Indiana Jones
...	...	...	...	...	...

# Additional Operations

---

- They can be defined in terms of the primitive operations
- They are added for convenience
- They are:
  - Join (Condition, Equi-, Natural) ( $\bowtie$ )
  - Division ( $/$ )
  - Assignment ( $\leftarrow$ )

# Joins ( $\bowtie$ )

---

- Condition Join:

$$R \bowtie_c S = \sigma_c(R \times S)$$

- *Result schema* same as cross-product.
- Fewer tuples than cross-product
  - might be able to compute more efficiently
- Sometimes called a *theta-join*.
  - The reference to an attribute of a relation R can be by position (R.i) or by name (R.name)



# Condition Join Example

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

StarsIn

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Ilsa Lund
3	3	Dorothy Gale

MovieStar ⋈<sub>MovieStar.StarID < StarsIn.StarID</sub> StarsIn

1	Name	Gender	MovieID	5	Character
1	Harrison Ford	Male	2	2	Ilsa Lund
1	Harrison Ford	Male	3	3	Dorothy Gale
2	Ingrid Bergman	Female	3	3	Dorothy Gale

MovieStar ⋈<sub>MovieStar.StarID < StarsIn.StarID</sub> StarsIn

MovieStar x StarsIn (first get the cross product)

1	Name	Gender	MovieID	5	Character
1	Harrison Ford	Male	1	1	Han Solo
2	Ingrid Bergman	Female	1	1	Han Solo
3	Judy Garland	Female	1	1	Han Solo
1	Harrison Ford	Male	4	1	Indiana Jones
2	Ingrid Bergman	Female	4	1	Indiana Jones
3	Judy Garland	Female	4	1	Indiana Jones
1	Harrison Ford	Male	2	2	Ilsa Lund
2	Ingrid Bergman	Female	2	2	Ilsa Lund
3	Judy Garland	Female	2	2	Ilsa Lund
1	Harrison Ford	Male	3	3	Dorothy Gale
2	Ingrid Bergman	Female	3	3	Dorothy Gale
3	Judy Garland	Female	3	3	Dorothy Gale

MovieStar ⋈<sub>MovieStar.StarID < StarsIn.StarID</sub> StarsIn

Now remove rows based on the condition stated above.


1	Name	Gender	MovieID	5	Character
1	Harrison Ford	Male	1	1	Han Solo
2	Ingrid Bergman	Female	1	1	Han Solo
3	Judy Garland	Female	1	1	Han Solo
1	Harrison Ford	Male	4	1	Indiana Jones
2	Ingrid Bergman	Female	4	1	Indiana Jones
3	Judy Garland	Female	4	1	Indiana Jones
1	Harrison Ford	Male	2	2	Ilsa Lund
2	Ingrid Bergman	Female	2	2	Ilsa Lund
3	Judy Garland	Female	2	2	Ilsa Lund
1	Harrison Ford	Male	3	3	Dorothy Gale
2	Ingrid Bergman	Female	3	3	Dorothy Gale
3	Judy Garland	Female	3	3	Dorothy Gale

# One more thing you may find helpful: Assignment Operation

---

- Notation:  $t \leftarrow E$   
assigns the result of expression  $E$  to a temporary relation  $t$ .
- Used to break complex queries to small steps.
- Assignment is always made to a temporary relation variable.
- Example: Write  $r \cap s$  in terms of  $\cup$  and/or  $-$

$temp1 \leftarrow r - s$   
 $result \leftarrow r - temp1$



# Equi-Join & Natural Join

- *Equi-Join*: A special case of condition join  $R \bowtie_c S = \sigma_c(R \times S)$ , where  $c$  contains only ***equalities***. Only the first of the repeated columns is retained.
- *Natural Join*: Equijoin on ***all*** common attributes
  - *Result schema*: similar to cross-product, but has only one copy of each common attribute
  - No need to show the condition
  - If the two attributes have no common attributes, this would be the same as cross product.
  - This is what we saw in BCNF & 3NF

# Equi and Natural Join Examples

MovieStar

StarID	Name	Gender
1	Harrison Ford	Male
2	Ingrid Bergman	Female
3	Judy Garland	Female

StarsIn

MovieID	StarID	Character
1	1	Han Solo
4	1	Indiana Jones
2	2	Ilsa Lund
3	3	Dorothy Gale

MovieStar ⋈ StarsIn

StarID	Name	Gender	MovieID	Character
1	Harrison Ford	Male	1	Han Solo
1	Harrison Ford	Male	4	Indiana Jones
3	Judy Garland	Female	3	Dorothy Gale
2	Ingrid Bergman	Female	2	Ilsa Lund

# Join Example

---

- Find the names of all Movie Stars who were in any Movie


Name
Harrison Ford
Ingrid Bergman
Judy Garland

# Assignment Operation

---

- Notation:  $t \leftarrow E$   
assigns the result of expression  $E$  to a temporary relation  $t$ .
- Used to break complex queries to small steps.
- Assignment is always made to a temporary relation variable.
- Example: Write  $r \cap s$  in terms of  $\cup$  and/or  $-$

$temp1 \leftarrow r - s$   
 $result \leftarrow r - temp1$





# Remember this one?

---

- Find the names of all Movie Stars who were in any Movie

$\pi_{\text{name}}(\text{MovieStar} \bowtie \text{StarsIn})$

Name
Harrison Ford
Ingrid Bergman
Judy Garland

- What if we wanted all Movies Stars who were in *all* movies?

# Division

- Notation:  $r / s$  or  $r \div s$
- Useful for expressing queries that include a notion of “**for all**” or “**for every**”, e.g., *Find movie stars who were in all movies.*
- Let  $r$  and  $s$  be relations on schemas  $R$  and  $S$  respectively where
  - $r = (A_1, \dots, A_m, B_1, \dots, B_n)$
  - $s = (B_1, \dots, B_n)$Then  $r / s$  is a relation on schema  
 $r / s = (A_1, \dots, A_m)$   
defined as
$$r / s = \{ t \mid t \in \Pi_{r-s}(r) \wedge \forall u \in s (tu \in r) \}$$
  - i.e.,  **$A/B$  contains all  $x$  tuples (MovieStars) such that for every  $y$  tuple (movies) in  $B$ , there is an  $x,y$  tuple in  $A$ .**

# Examples of Division A/B

*A*

sno	pno
s1	p1
s1	p2
s1	p3
s1	p4
s2	p1
s2	p2
s3	p2
s4	p2
s4	p4

*B1*

pno
p2

*B2*

pno
p2
p4

*B3*

pno
p1
p2
p4

*A/B1*

sno
s1
s2
s3
s4

*A/B2*

sno
s1
s4

*A/B3*

sno
s1

# Find the name of actors who have been in ***all*** movies

---

Be careful in choosing the input relations!

$\text{InAll} \leftarrow \pi_{\text{StarID}, \text{MovieID}} \text{StarsIn} / \pi_{\text{MovieID}}(\text{Movie})$   
 $\pi_{\text{Name}}(\text{InAll} \bowtie \text{MovieStar})$

## Case study of complex relational algebra: build up division of $r/s$ from other operators

---

- Let  $X$  be attributes not in  $R$  and  $Y$  be attributes in  $S$
- Idea: compute all values that are “disqualified” by some value in  $s$ .
  - value  $x$  is *disqualified* if by attaching  $y$  value from  $s$ , we obtain an  $xy$  tuple that is not in  $r$ .
- Take difference from all values

# Expressing r/s Using Basic Operators

- Like a join, can be computed from basic operators
- *Idea:*
  - let  $X$  the set of attributes of  $r$  that are not in  $s$
  - (1) compute the  $X$ -projection of  $r$
  - (2) compute all  $X$ -projection values of  $r$  that are “disqualified” by some value in  $s$ .
    - value  $x$  is *disqualified* if by attaching  $y$  value from  $s$ , we obtain an  $xy$  tuple that is not in  $r$ .
  - result is (1)-(2)
- So,
  - Disqualified  $x$  values:  $\pi_X((\pi_X(r) \times s) - r)$
  - $r/s$  is  $\pi_X(r) - \pi_X((\pi_X(r) \times s) - r)$

# Example: building up division subtract off disqualified answers

$A=R$

Sno	Pno
S1	P1
S1	P2
S1	P3
S1	P4
S2	P1
S2	P2
S3	P2
S4	P2
S4	P4

$B2 = S$

Pno
P2
P4

$\pi_X(R)$

Sno
S1
S2
S3
S4



All possible  
values given R

$\pi_X(R) \times S$

Sno	Pno
S1	P2
S1	P4
S2	P2
S2	P4
S3	P2
S3	P4
S4	P2
S4	P4



$\pi_X(R) \times S - R$

Sno	Pno
S2	P4
S3	P4



These values  
aren't in R

Values needed  
for  $\pi_X(R)$

$$\pi_X(R) - \pi_X(\pi_X(R) \times S - R) = A/B2$$

$A/B2 =$

Sno
S1
S4



Answers not disqualified

# If you want to practice...

---

Try this site:

<https://dbis-uibk.github.io/relax/>

You can play with RA queries and see the results.



# Learning Goals Revisited

---



- Identify the basic operators in RA.
- Use RA to create queries that include combining RA operators.
- Given an RA query and table schemas and instances, compute the result of the query.