

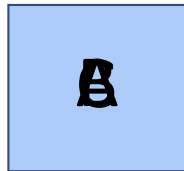
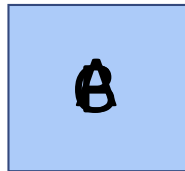
Today

- Learning Objectives
 - Analyze different cache eviction policies
 - When LRU fails: MRU
 - First-in-first-out: FIFO
- How we'll get there
 - Review in-class problem
 - Look at one last replacement algorithm
- Reading
 - 6.4

When LRU is terrible

- We left you in class asking you to derive a sequence of accesses that made LRU perform terribly.
- It's a common and simple access pattern:

A, B, C, A, B, C



Compulsory misses: | | |

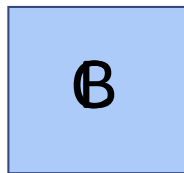
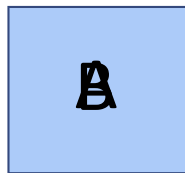
Capacity misses: | | |

Hits:

LRU is always doing the wrong thing!

- Iterating over something is pretty common, and if that something is too big, LRU is awful.
- What would be better?

A, B, C, A, B, C



Compulsory misses: | | |

Capacity misses: |

Hits: | |

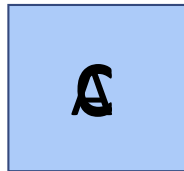
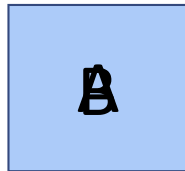
Most-Recently-Used: MRU

- Ideal when iterating over an array that is too large.
 - In other workloads, not necessarily particularly good.

Let's see what FIFO does.

- Evict the item that was loaded into the cache earliest.
 - Assume we have 2 slots in our cache.
 - Assume **fully associative**
 - We access cache blocks in the following order:

A, A, C, A, C, B, B, B, A, A, B, C



Compulsory misses: | | |

Capacity misses: | |

Hits: | | | | | | |

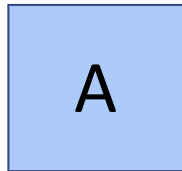
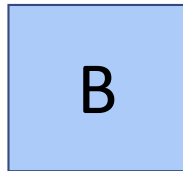
Let's see what FIFO does.

- Evict the item that was loaded into the cache earliest.
 - Assume we have 2 slots in our cache.
 - Assume **fully associative**
 - We access cache blocks in the following order:

A, A, C, A, C, B, B, B, A, A, B, C



Hit Rate: 7/12



Compulsory misses: | | |

Capacity misses: | |

Hits: | | | | | | |

There is no Perfect Policy

- Most commonly used:
 - LRU (frequently just fine)
 - LFU (sometimes better)
 - MRU (when you detect the right pattern)
- For anything but Belady, we can construct an adversarial workload that can be pretty devastating.