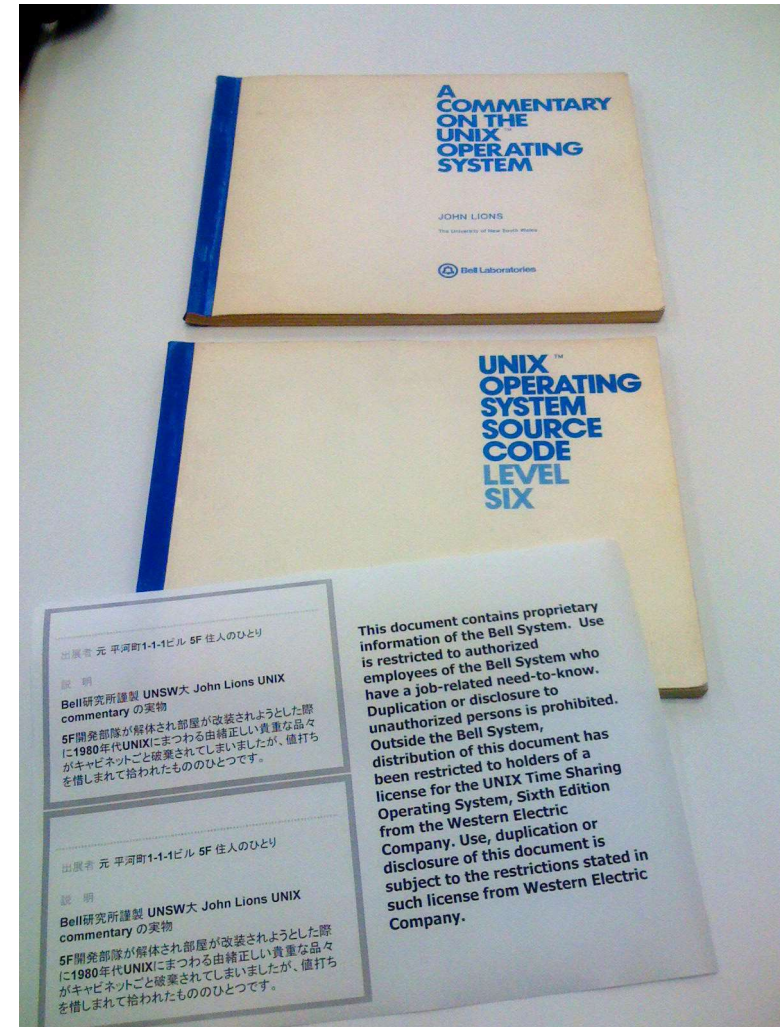
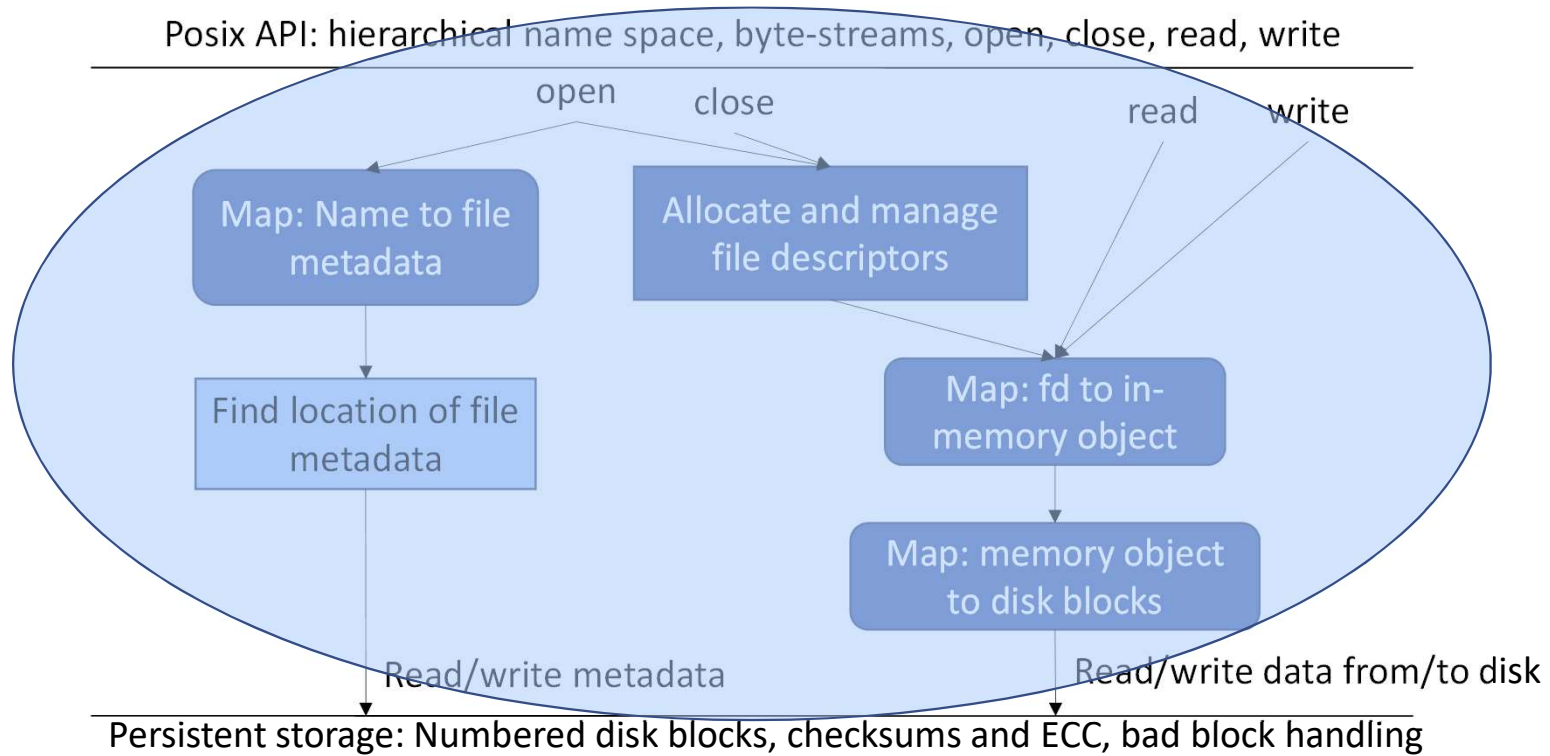


# CPSC313: Computer Hardware and Operating Systems

File Systems  
Case Study: V6



You are here!



# Outline

- How do we lay out a real file system?
  - **Disk Layout:** What goes where, and how big is it?
  - ***File System Metadata*:** How do I find all the parameters of this file system?
  - ***Free Space Management*:** How do I find the next block to allocate?  
How do I deallocate a block?
  - **Logical-Physical Mapping:** How do I find a particular block within a file?
  - **Directory Structures:** How do I make paths/directories work?
- Learning Outcomes
  - Discuss the tradeoffs in file system design choices
  - Describe how the Unix v6 file system represented files

# How Do We Decide?

We're designing all the moving pieces of a file system. How do we pick what it looks like?



Usable! Good for our human end users (and programs!)



This Photo by Unknown Author is licensed under [CC BY-NC](#)

Fast! (When working with big files? Small? Allocating new blocks?)



This Photo by Unknown Author is licensed under [CC BY-NC](#)

Compact! (In terms of internal fragmentation? External? Persistent FS data structures?)



Easy to restore in the face of damage to the media?

# Our Real File System: Unix v6

With enormous thanks to:

Ken Thompson and Dennis Ritchie, John Lions

([https://en.wikipedia.org/wiki/Lions'\\_Commentary\\_on\\_UNIX\\_6th\\_Edition,\\_with\\_Source\\_Code](https://en.wikipedia.org/wiki/Lions'_Commentary_on_UNIX_6th_Edition,_with_Source_Code)), Keith Bostic,

Whoever is behind: <http://v6.czuco.com/>

# The Unix v6 File System

One of the first widely distributed versions of Unix (1975), ran on a “PDP-11” minicomputer from DEC. Critically for us:

- **Released as full source code**
- **Documented** by John Lions of the University of New South Wales.


We start with this *relatively* simple ancestor of most modern FS's.



This Photo by Unknown Author is licensed under [CC BY-NC](#)

In 1975, **disks were tiny**. You'll notice ***compactness*** is a key goal, sometimes at the expense of others!

# Disk Layout



One block is 512 bytes.


# File system level metadata

- `struct filsys`
  - Today, we call this the **superblock**
  - Created when you create the file system
  - Read when you mount the file system

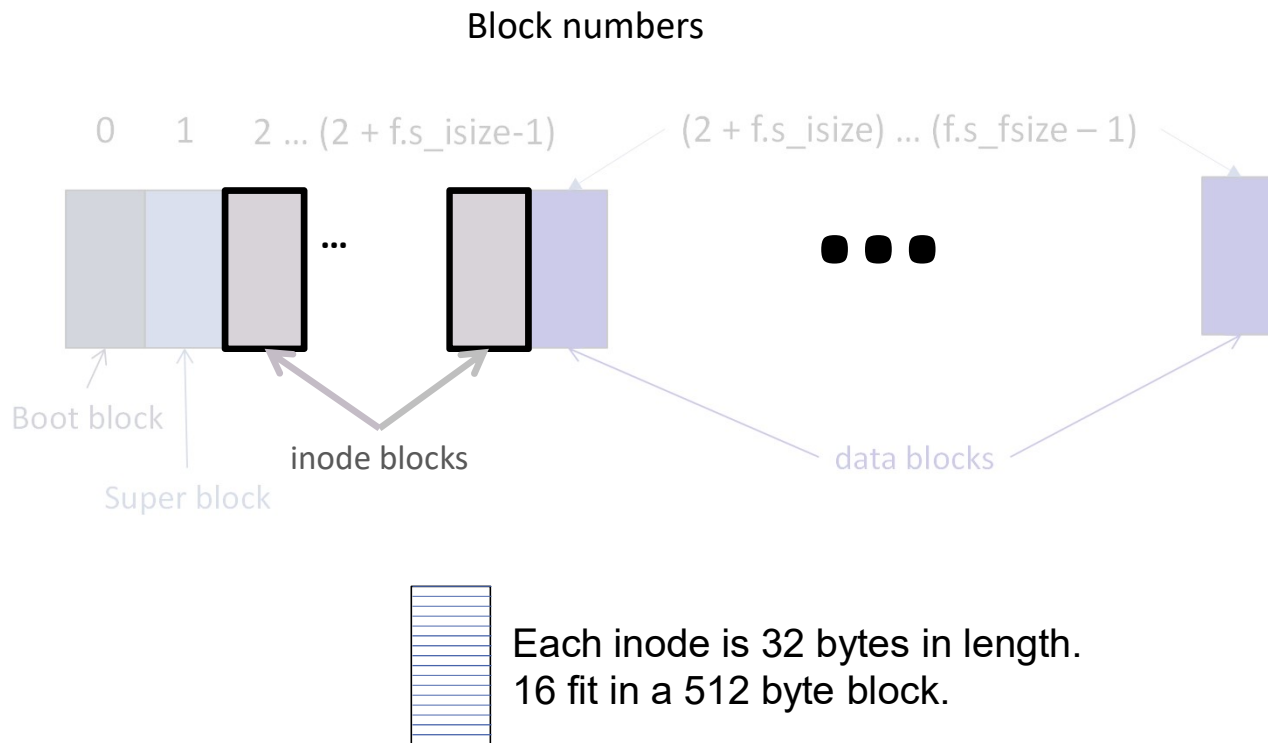
```
struct filsys {  
    int s_isize;      /* size in blocks of the I list */  
    int s_fsize;      /* size in blocks of the entire volume */  
    int s_nfree;      /* number of in core free blocks (between 0 and 100) */  
    int s_free[100];  /* in core free blocks */  
    int s_ninode;     /* number of in core I nodes (0-100) */  
    int s_inode[100]; /* in core free I nodes */  
    char s_flock;     /* lock during free list manipulation */  
    char s_ilock;     /* lock during I list manipulation */  
    char s_fmod;      /* super block modified flag */  
    char s_ronly;     /* super block read-only flag */  
    int s_time[2];    /* current date of last update */  
    int pad[50];  
}
```



# Disk Layout



# Disk Layout: What's in an inode block?



# Per-file Metadata (on-disk) [inode]

Size in bytes..

- On disk inode (note: **int is 2 bytes**):

```
struct ino {  
    2 int    i_mode;          /* File type, size, permissions */  
    1 char   i_nlink;         /* Link count */  
    1 char   i_uid;          /* Owner user id */  
    1 char   i_gid;          /* Group id */  
    1 char   i_size0;         /* most significant bits of size */  
    2 int    i_size1;         /* least sig */  
    16 int   i_addr[8];       /* Disk addresses of blocks */  
    4 int    i_atime[2];      /* Access time */  
    4 int    i_mtime[2];      /* Modified time */  
}  
32!
```

Let's come back to this later when  
we map logical -> physical blocks!




Each inode is 32 bytes in length.  
16 fit in a 512 byte block.

# Outline

- How do we lay out a real file system?
  - **Disk Layout:** What goes where, and how big is it?
  - **File System Metadata:** How do I find all the parameters of this file system?
  - **Free Space Management:** How do I find the next block to allocate?  
How do I deallocate a block?
  - **Logical-Physical Mapping:** How do I find a particular block within a file?
  - **Directory Structures:** How do I make paths/directories work?
- Learning Outcomes
  - Discuss the tradeoffs in file system design choices
  - Describe how the Unix v6 file system represented files

# Lie #1: Just Have a Linked-List of Free Blocks!



Tiny disk! Don't waste a bunch of space tracking free blocks. **Use** the free blocks to **track** the free blocks!

## Problems?

Every allocate or deallocate messes with a random block.

We'd rather touch only the (probably-cached!) super block whenever possible.

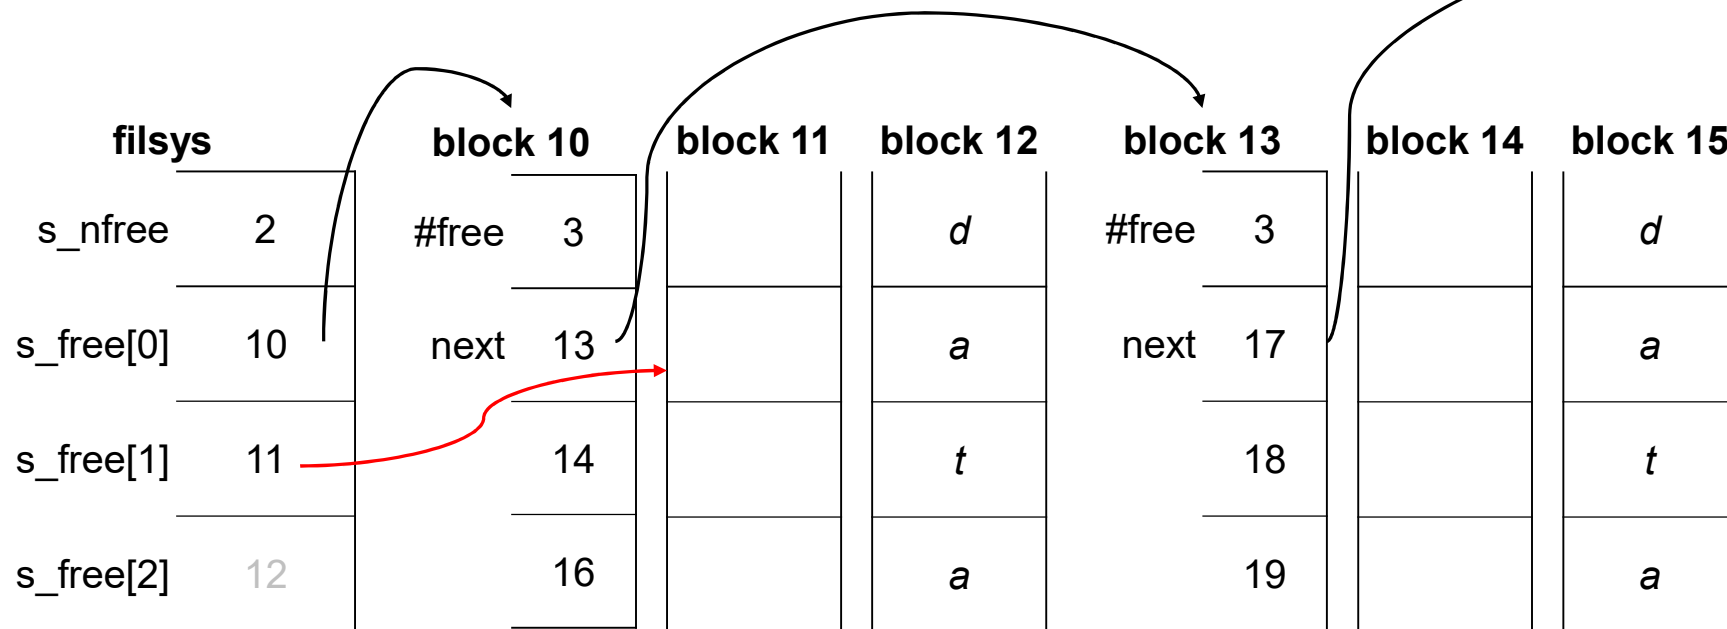
Wait... a block number (pointer) is small! Let's put **three** free blocks in each linked list node instead of one.



This Photo by Unknown Author is licensed under CC BY-NC

Lie #2: Each node has an array of up to 3 block #s.  
 (Our **next** pointer is **s\_free[0]**.)

- Initial state:



# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate a block: *#11*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	2	#free 3		<i>d</i>	#free 3		<i>d</i>
s_free[0]	10	next 13		<i>a</i>	next 17		<i>a</i>
s_free[1]	11	14		<i>t</i>	18		<i>t</i>
s_free[2]	12	16		<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate a block: *#11*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	1	#free 3	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	10	next 13	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	11	14	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	12	16	<i>a</i>	<i>a</i>	19		<i>a</i>



# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate another block: *#10*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	1	#free 3	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	10	next 13	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	11	14	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	12	16	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example


- Assume we store block numbers in groups of 3.
- We allocate another block: *#10 [step 1]*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	1	#free 3	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	10	next 13	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	11	14	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	12	16	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate another block: *#10 [step 1]*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	#free 3	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	13	next 13	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	14	14	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	16	16	<i>a</i>	<i>a</i>	19		<i>a</i>



# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate another block: *#10 [step 2]*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	<i>d</i>	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	16	<i>a</i>	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate another block: *#16*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	<i>d</i>	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	16	<i>a</i>	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We allocate another block: *#16*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	2	<i>d</i>	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	16	<i>a</i>	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We deallocate block **#12**

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	2	<i>d</i>	<i>d</i>	<i>d</i>	#free 3		<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>	<i>a</i>	next 17		<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>	<i>t</i>	18		<i>t</i>
s_free[2]	16	<i>a</i>	<i>a</i>	<i>a</i>	19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We deallocate block **#12**

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	<i>d</i>	<i>d</i>		#free	3	<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>		next	17	<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>			18	<i>t</i>
s_free[2]	12	<i>a</i>	<i>a</i>			19	<i>a</i>



# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We deallocate block *#15*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	<i>d</i>	<i>d</i>		#free 3		<i>d</i>
s_free[0]	13	<i>a</i>	<i>a</i>		next 17		<i>a</i>
s_free[1]	14	<i>t</i>	<i>t</i>		18		<i>t</i>
s_free[2]	12	<i>a</i>	<i>a</i>		19		<i>a</i>

# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We deallocate block *#15 [step 1]*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	3	<i>d</i>	<i>d</i>		#free 3		<i>3</i>
s_free[0]	13	<i>a</i>	<i>a</i>		next 17		<i>13</i>
s_free[1]	14	<i>t</i>	<i>t</i>		18		<i>14</i>
s_free[2]	12	<i>a</i>	<i>a</i>		19		<i>12</i>



# Free Space Management Example

- Assume we store block numbers in groups of 3.
- We deallocate block *#15 [step 2]*

filsys		block 10	block 11	block 12	block 13	block 14	block 15
s_nfree	1	<i>d</i>	<i>d</i>		#free 3		3
s_free[0]	15	<i>a</i>	<i>a</i>		next 17		13
s_free[1]	14	<i>t</i>	<i>t</i>		18		14
s_free[2]	12	<i>a</i>	<i>a</i>		19		12

# Strength? Weaknesses?

We need very little extra space to store our free list!

But the free list works a lot like a linked-list-style malloc/memory manager. Will we end up with a lot of contiguous or nearby blocks in a given file?



This Photo by Unknown Author is licensed under [CC BY-NC](#)

**vs.**



**?**

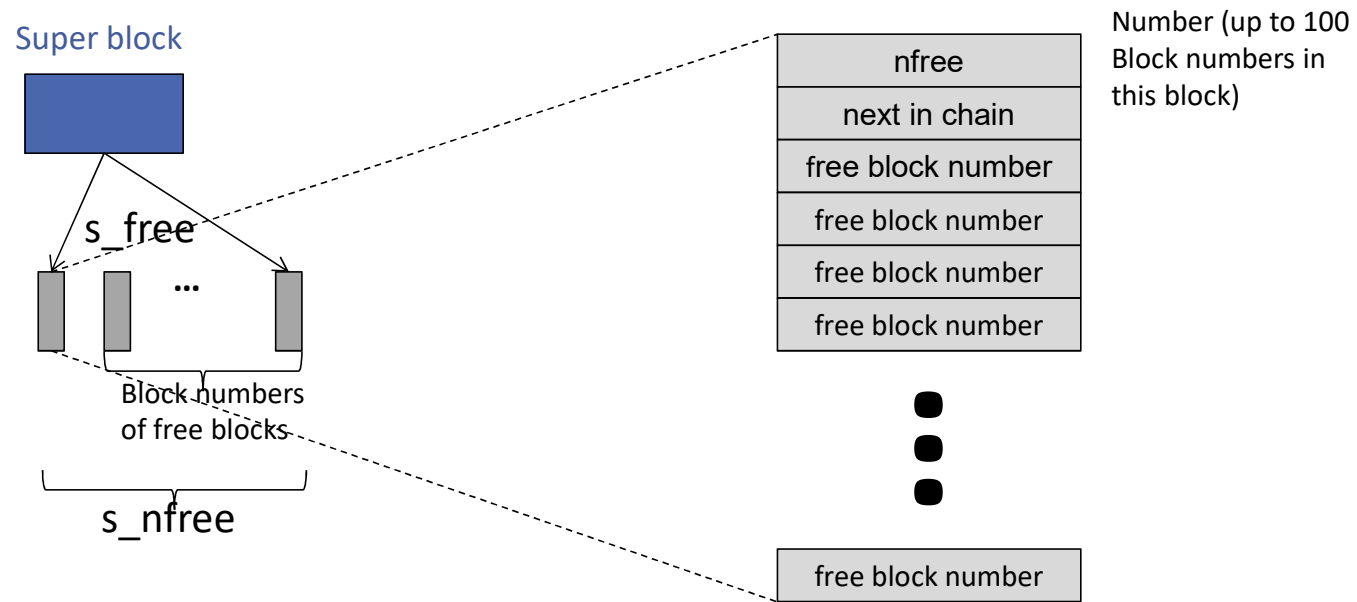
This Photo by Unknown Author is licensed under [CC BY-NC](#)

# Actual V6 Free Space Management: 100 not 3

- struct filsys
  - Today, we call this the **superblock**
  - Created when you create the file system
  - Read when you mount the file system

```
struct filsys {
    int s_ismax;      /* size in blocks of the I list */
    int s_fsize;      /* size in blocks of the entire volume */
    int s_nfree;      /* number of in-core free blocks (between 0 and 100) */
    int s_free[100];  /* in core free blocks */
    int s_ninode;     /* number of in-core I nodes (0-100) */
    int s_inode[100]; /* in core free I nodes */
    char s_flock;      /* lock during free list manipulation */
    char s_ilock;      /* lock during I list manipulation */
    char s_fmod;      /* super block modified flag */
    char s_ronly;      /* super block read-only flag */
    int s_time[2];     /* current date of last update */
    int pad[50];
}
```

# Free Space Management



# Outline

- How do we lay out a real file system?
  - **Disk Layout:** What goes where, and how big is it?
  - **File System Metadata:** How do I find all the parameters of this file system?
  - **Free Space Management:** How do I find the next block to allocate?  
How do I deallocate a block?
  - **Logical-Physical Mapping:** How do I find a particular block within a file?
  - **Directory Structures:** How do I make paths/directories work?
- Learning Outcomes
  - Discuss the tradeoffs in file system design choices
  - Describe how the Unix v6 file system represented files


# Per-file Metadata (on-disk) [inode]

- On disk inode (note: int is 2 bytes):

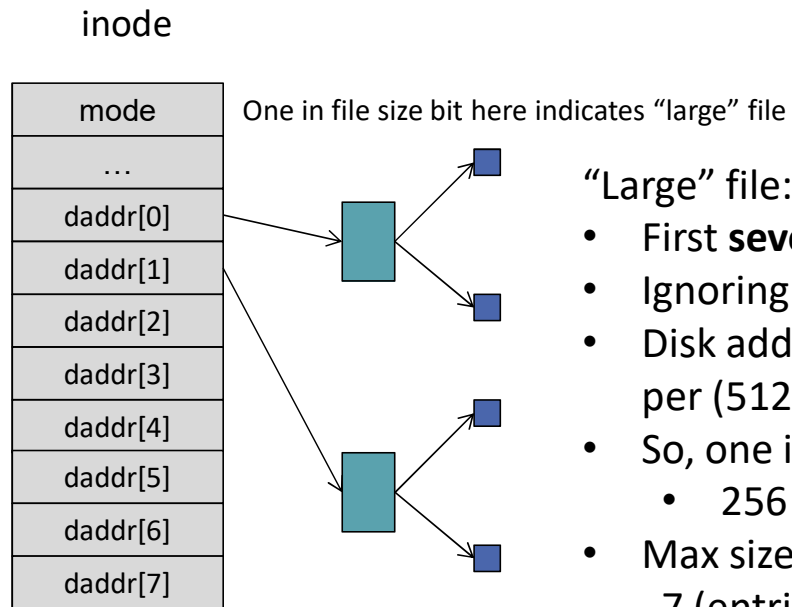
```
struct ino {  
    int    i_mode;           /* File type, size, permissions */  
    char   i_nlink;          /* Link count */  
    char   i_uid;            /* Owner user id */  
    char   i_gid;            /* Group id */  
    char   i_size0;          /* most significant bits of size */  
    int    i_size1;          /* least sig */  
    int    i_addr[8];        /* Disk addresses of blocks */  
    int    i_atime[2];        /* Access time */  
    int    i_mtime[2];        /* Modified time */  
}
```



# Different sized files (1)



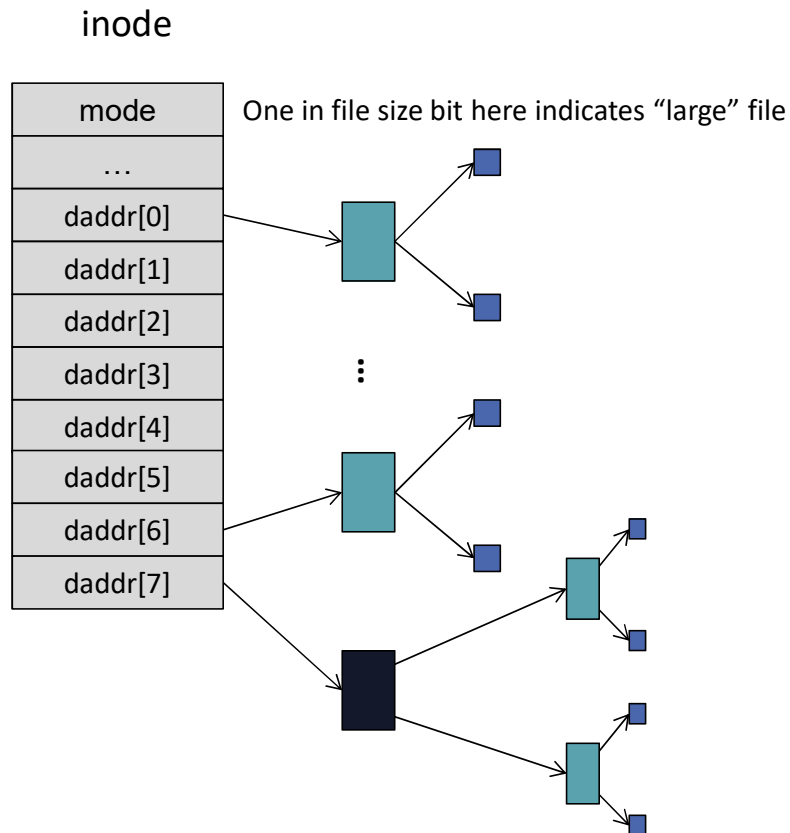
## Different sized files (2)



### "Large" file:

- First **seven** entries contain addresses of **indirect** blocks.
- Ignoring daddr[7] until the next slide ...
- Disk addresses are 2 bytes, so 512/2 or 256 disk addresses per (512-byte) block
- So, one indirect block can access:
  - 256 (daddrs per block) \* 512 bytes per block = **128 KB**
- Max size:
  - 7 (entries in daddr) \* 128 KB/indirect block = **896 KB**

## Different sized files (3)

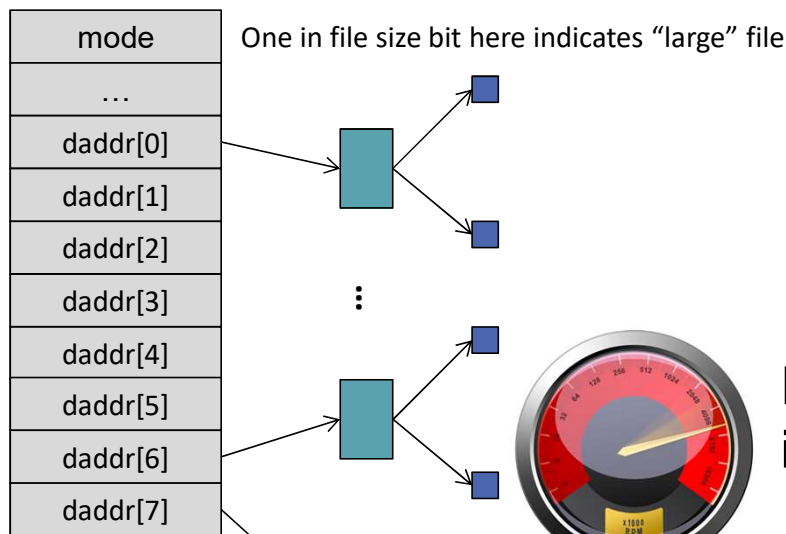


"Huge" file:

- 7 indirect blocks addresses in daddr, AND
- **daddr[7]** contains a **double indirect** block.
- From the previous slide:
  - 1 indirect block reaches 128 KB
- A double indirect contains 256 daddrs
- A double indirect reaches:
  - $256 * 128 \text{ KB} = 32 \text{ MB}$
- Max size:
  - $896 \text{ KB} + 32 \text{ MB}$

## Different sized files (3)

inode



"Huge" file:

- 7 indirect blocks addresses in daddr, AND
- **daddr[7]** contains a double indirect block.
- From the previous slide:
  - 1 indirect block reaches 128 KB

Fast... for small files no "wasted" indirect blocks to pass through.



This Photo by Unknown Author is licensed under [CC BY-NC](#)



Usable? 32MB **was** huge back then!

Compact? Small files don't need "wasted" indirect blocks. Small files were (and are) common.

# Per-file Metadata (in-memory) [vnode]

- In-memory inode (note: int is 2 bytes):

```
struct inode {
    char    i_flag;
    char    i_count;          /* reference count */
    int     i_dev;            /* device where inode resides */
    int     i_number;         /* i number 1:1 w/device addr */
    int     i_mode;
    char    i_nlink;          /* directory entries*/
    char    i_uid;            /* owner */
    char    i_gid;            /* group of owner */
    char    i_size0;          /* most significant of size */
    int     i_size1;          /* least sig */
    int     i_addr[8];        /* device addresses constituting file */
    int     i_lastr;          /* last logical block read */
}
```

# Outline

- How do we lay out a real file system?
  - **Disk Layout:** What goes where, and how big is it?
  - **File System Metadata:** How do I find all the parameters of this file system?
  - **Free Space Management:** How do I find the next block to allocate?  
How do I deallocate a block?
  - **Logical-Physical Mapping:** How do I find a particular block within a file?
  - **Directory Structures:** How do I make paths/directories work?
- Learning Outcomes
  - Discuss the tradeoffs in file system design choices
  - Describe how the Unix v6 file system represented files

# Directory Entries

- Hierarchical directory structure that you know and love, including “.” and “..”.
- Directory entries are 16 bytes (they are fixed size!):
  - 2 bytes of inode number
  - 14 bytes (right padded) of name
- A directory entry with inode = 0 is unused

# Directory Exercise - revisited

```
00000000 00 01 2e 00 00 00 00 00 00 00 00 00 00 00 00 00
00000010 01 01 2e 2e 00 00 00 00 00 00 00 00 00 00 00 00
00000020 64 00 57 6f 6d 62 61 74 00 00 00 00 00 00 00 00
00000030 65 00 4d 61 63 6b 65 72 65 6c 00 00 00 00 00 00
00000040 66 00 44 6f 76 65 00 00 00 00 00 00 00 00 00 00
00000050 67 00 43 61 74 74 6c 65 00 00 00 00 00 00 00 00
00000060 00 00 44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000070 02 01 4b 6f 69 00 00 00 00 00 00 00 00 00 00 00
00000080 00 00 52 61 62 62 69 74 00 00 00 00 00 00 00 00
00000090 68 00 43 6f 6e 64 6f 72 00 00 00 00 00 00 00 00
00000a0 69 00 44 6f 6e 6b 65 79 00 00 00 00 00 00 00 00
00000b0 03 01 41 6c 70 61 63 61 00 00 00 00 00 00 00 00
00000c0 6a 00 44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000d0
```

```
struct dirent {
    int d_ino;    // 16 bits
    char d_name[14];
};
```



# In-class Exercise Revisited

00000000	00	01	2e	00	00	00	00	00	00	00	00	00	00	00	00
00000100	01	01	2e	2e	00	00	00	00	00	00	00	00	00	00	00
00000200	64	00	57	6f	6d	62	61	74	00	00	00	00	00	00	00
00000300	65	00	4d	61	63	6b	65	72	65	6c	00	00	00	00	00
00000400	66	00	44	6f	76	65	00	00	00	00	00	00	00	00	00
00000500	67	00	43	61	74	74	6c	65	00	00	00	00	00	00	00
00000600	00	00	44	6f	6d	65	73	74	69	63	20	72	61	62	62
00000700	02	01	4b	6f	69	00	00	00	00	00	00	00	00	00	00
00000800	00	00	52	61	62	62	69	74	00	00	00	00	00	00	00
00000900	68	00	43	6f	6e	64	6f	72	00	00	00	00	00	00	00
00000a00	69	00	44	6f	6e	6b	65	79	00	00	00	00	00	00	00
00000b00	03	01	41	6c	70	61	63	61	00	00	00	00	00	00	00
00000c00	6a	00	44	6f	6d	65	73	74	69	63	20	72	61	62	62
00000d00															

```
struct dirent {  
    int d_ino;    // 16 bits  
    char d_name[14];  
};
```

# In-class Exercise Revisited

00000000	00 01	2e 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000100	01 01	2e 2e 00 00 00 00 00 00 00 00 00 00 00 00 00
00000200	64 00	57 6f 6d 62 61 74 00 00 00 00 00 00 00 00 00
00000300	65 00	4d 61 63 6b 65 72 65 6c 00 00 00 00 00 00 00
00000400	66 00	44 6f 76 65 00 00 00 00 00 00 00 00 00 00 00
00000500	67 00	43 61 74 74 6c 65 00 00 00 00 00 00 00 00 00
00000600	00 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000700	02 01	4b 6f 69 00 00 00 00 00 00 00 00 00 00 00 00
00000800	00 00	52 61 62 62 69 74 00 00 00 00 00 00 00 00 00
00000900	68 00	43 6f 6e 64 6f 72 00 00 00 00 00 00 00 00 00
00000a00	69 00	44 6f 6e 6b 65 79 00 00 00 00 00 00 00 00 00
00000b00	03 01	41 6c 70 61 63 61 00 00 00 00 00 00 00 00 00
00000c00	6a 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000d00		

```
struct dirent {
    int d_ino;    // 16 bits
    char d_name[14];
};
```

Inode numbers

# In-class Exercise Revisited

00000000	00 01	2e 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000100	01 01	2e 2e 00 00 00 00 00 00 00 00 00 00 00 00
00000200	64 00	57 6f 6d 62 61 74 00 00 00 00 00 00 00 00
00000300	65 00	4d 61 63 6b 65 72 65 6c 00 00 00 00 00 00
00000400	66 00	44 6f 76 65 00 00 00 00 00 00 00 00 00 00
00000500	67 00	43 61 74 74 6c 65 00 00 00 00 00 00 00 00
00000600	00 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000700	02 01	4b 6f 69 00 00 00 00 00 00 00 00 00 00 00
00000800	00 00	52 61 62 62 69 74 00 00 00 00 00 00 00 00
00000900	68 00	43 6f 6e 64 6f 72 00 00 00 00 00 00 00 00
00000a00	69 00	44 6f 6e 6b 65 79 00 00 00 00 00 00 00 00
00000b00	03 01	41 6c 70 61 63 61 00 00 00 00 00 00 00 00
00000c00	6a 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000d00		

```
struct dirent {
    int d_ino;    // 16 bits
    char d_name[14];
};
```

Inode numbers

Names

# In-class Exercise Revisited

00000000	00 01	2e 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000100	01 01	2e 2e 00 00 00 00 00 00 00 00 00 00 00 00
00000200	64 00	57 6f 6d 62 61 74 00 00 00 00 00 00 00 00
00000300	65 00	4d 61 63 6b 65 72 65 6c 00 00 00 00 00 00
00000400	66 00	44 6f 76 65 00 00 00 00 00 00 00 00 00 00
00000500	67 00	43 61 74 74 6c 65 00 00 00 00 00 00 00 00
00000600	00 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000700	02 01	4b 6f 69 00 00 00 00 00 00 00 00 00 00 00
00000800	00 00	52 61 62 62 69 74 00 00 00 00 00 00 00 00
00000900	68 00	43 6f 6e 64 6f 72 00 00 00 00 00 00 00 00
00000a00	69 00	44 6f 6e 6b 65 79 00 00 00 00 00 00 00 00
00000b00	03 01	41 6c 70 61 63 61 00 00 00 00 00 00 00 00
00000c00	6a 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69
00000d00		

```
struct dirent {
    int d_ino;    // 16 bits
    char d_name[14];
};
```

Let's figure out the name and type for inode 0x6a

Inode numbers

Names

Domestic rabbit

# In-class Exercise Revisited

00000000	00 01	2e 00 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000100	01 01	2e 2e 00 00 00 00 00 00 00 00 00 00 00 00 00 00
00000200	64 00	57 6f 6d 62 61 74 00 00 00 00 00 00 00 00 00 00
00000300	65 00	4d 61 63 6b 65 72 65 6c 00 00 00 00 00 00 00 00
00000400	66 00	44 6f 76 65 00 00 00 00 00 00 00 00 00 00 00 00
00000500	67 00	43 61 74 74 6c 65 00 00 00 00 00 00 00 00 00 00
00000600	00 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69 69 69
00000700	02 01	4b 6f 69 00 00 00 00 00 00 00 00 00 00 00 00 00
00000800	00 00	52 61 62 62 69 74 00 00 00 00 00 00 00 00 00 00
00000900	68 00	43 6f 6e 64 6f 72 00 00 00 00 00 00 00 00 00 00
00000a00	69 00	44 6f 6e 6b 65 79 00 00 00 00 00 00 00 00 00 00
00000b00	03 01	41 6c 70 61 63 61 00 00 00 00 00 00 00 00 00 00
00000c00	6a 00	44 6f 6d 65 73 74 69 63 20 72 61 62 62 69 69 69

```
struct dirent {
    int d_ino;    // 16 bits
    char d_name[14];
};
```



Usable? No. Even back then it chafed to have tiny file names.



This Photo by Unknown Author is licensed under CC BY-NC

Compact? Yes.

BTW: Also simple to implement!