

CPSC 313 C Refresher

2023W2

Contents

- Basics (number representation, data types, bit shifting, type casting)
- Memory (structs, alignment, endianness)
- Pointers (segfaults, dangling ptrs, mem leaks, stack vs heap memory)
- C programs (arguments, debugging)

Part 1:

Basics (number representation, data types,
bit shifting, type casting)

Bytes, Bits and Hexits

Remember your conversions: decimal \leftrightarrow binary \leftrightarrow hexadecimal

- ex. $(101)_2 = 2^2 + 0 + 2^0 = (5)_{10}$ (binary \leftrightarrow decimal)
- 1 “hexit” = 1 hex digit = 4 bits, (ex. **0xF** = **0b1111**)
- 1 byte = 8 bits, so 1 byte = 2 hexits



Addressing Bytes

- Each memory address (ex. `0xDEADBEEF`) houses 1 byte of data
 - note: the memory address is not one byte in size, the value housed in each memory address is 1 byte in size

Q: What is the size of the memory address `0xDEADBEEF`?

Addressing Bytes

- Each memory address (ex. 0xDEADBEEF) houses 1 byte of data
 - note: the memory address is not one byte in size, the value housed in each memory address is 1 byte in size

Q: What is the size of the memory address 0xDEADBEEF?

A: 4 bytes

Data Types

Try typing "long long long i;" into gcc

In this course, assume a **64 bit** CPU (since y86) if unspecified

- so assume for now: `char` is 1 byte, `short` is 2 bytes
- TIP: can always use `sizeof()` to double check sizes
- If you `#include <stdint.h>` you may see `uint32_t` explicitly defines an unsigned 32 bit (4 byte) integer
 - similar for `int32_t`, `uint8_t`, etc.

Q: How many bytes is a `uint64_t`?

Data Types

Try typing "long long long i;" into gcc

In this course, assume a **64 bit** CPU (since y86) if unspecified

- so assume for now: `char` is 1 byte, `short` is 2 bytes
- TIP: can always use `sizeof()` to double check sizes
- If you `#include <stdint.h>` you may see `uint32_t` explicitly defines an unsigned 32 bit (4 byte) integer
 - similar for `int32_t`, `uint8_t`, etc.

Q: How many bytes is a `uint64_t`?

A: 8 bytes

Signed vs Unsigned

- **What is the difference between signed and unsigned binary numbers?**
- Recall: use two's complement representation for signed binary numbers (flip all bits, add 1)
- Helpful tip: if we know the number is signed the MSB (most significant bit) is the sign bit

Ranges we can represent with n bits

Given n bits, the range of ints we can represent is:

Unsigned : $[0, 2^n - 1]$

Signed : $[-2^{n-1}, 2^{n-1} - 1]$

Why? (Hint: how many permutations of n bits are there?)

So if we have an 32 bit integer, plug in for respective ranges :

unsigned int: 0 to 4294967295, **int:** -2,147,483,648 to 2,147,483,647

Bit Shifting

<<: left shift operator
>>: right shift operator

```
a = a << 2      a <<= 2  
b = b >> 2      b >>= 2
```

```
char c = 0b00000001;  
c <<= 6; // 0b01000000  
c >>= 7; // 0b00000000
```

Edge cases

```
char c = 0b10000000;  
c >>= 7; // 0b00000001 or 0b11111111  
Specified, Implementation-dependent
```

```
char c = 0b00000001;  
c <<= 8 or above; // ?  
Not well defined!
```

Type Casting (Integer types)

Truncation (casting to smaller type)

```
int16_t i = 0x000A;  
char c = (char) i; // 0x0A
```

```
i = 0xFFFF;  
c = (char) i; // 0xFF
```

Signed extension (casting to larger type)

```
char c = 0x0A;  
int16_t i = (int16_t) c; // 0x000A
```

```
c = 0xFF;  
i = (int16_t) c; // 0xFFFF
```

Bit Masking

Q: How do you get the value of bit number 3 and bit number 5?

```
char c = 0b11011011; // note you can also define in hex using 0x...  
char mask = ???
```

Bit Masking

Q: How do you get the value of bit number 3 and bit number 5?

```
char c = 0b11011011; // note you can also define in hex using 0x...  
char mask = (0b01 << 3) | (0b01 << 5);  
char result = ???
```

Bit Masking

Q: How do you get the value of bit number 3 and bit number 5?

```
char c = 0b11011011; // note you can also define in hex using 0x...  
char mask = (0b01 << 3) | (0b01 << 5);  
char result = c & mask; // 0b00001000
```

Strings

- In C, we represent a string using a null terminated char array
 - This means the end of a string is marked by `'\0'` or `0x0`
- String functions are defined in `<string.h>`
 - `strlen` counts length of null terminated string (excluding null char)
 - `memset` fills memory with the same value (e.g. useful for zeroing data)
 - `strcpy/strncpy/memcpy/memmove` copies data
 - `strcmp/strncmp/memcmp` compares data

Q: What would `printf("%s", str)` output?

```
char str[] = "hello, world!";  
str[5] = '\0';
```


Strings

- In C, we represent a string using a null terminated char array
 - This means the end of a string is marked by `'\0'` or `0x0`
- String functions are defined in `<string.h>`
 - `strlen` counts length of null terminated string (excluding null char)
 - `memset` fills memory with the same value (e.g. useful for zeroing data)
 - `strcpy/strncpy/memcpy/memmove` copies data
 - `strcmp/strncmp/memcmp` compares data

Q: What would `printf("%s", str)` output?

```
char str[] = "hello, world!";  
str[5] = '\0';
```

A: hello

Part 2:

Memory (structs, alignment, endianness)

Structs

```
struct myStruct {  
    char a;      // 1 byte  
    int32_t i;   // 4 bytes  
    char b[3];   // 3 bytes  
};
```

Initialize a struct called ms

```
struct myStruct ms;
```

Set the "a" field to 0x123

```
ms.a = 0x123;
```

Write a function that takes a struct

```
void func(struct myStruct p) { ...
```

Q: What is the size of ms in memory? We will find out...

Alignment Rules

Structs in C may be larger than they appear

1. Whole struct must be aligned by the **alignment of its largest field**

Ex. if the largest thing in the struct is 8 bytes, the size of the entire struct must be a multiple of 8. (8, 16, 24, 32...)

We say "the **alignment of the struct** is 8" or "the struct is **8-aligned**"

2. Arrays are aligned by the size of their elements
e.g. a 100000-long array of `int32_t` (4 bytes) must start at an address that is a multiple of 4. It is 4-aligned.
3. Anything not a struct or an array is aligned by its raw size
e.g. a `short` (2 bytes) must start at an address that is a multiple of 2. It is 2-aligned.

If any rules are not satisfied, add padding (empty space) to the struct until they are.

The answer

```
struct myStruct {  
    char a;    // 1 byte  
    int32_t i; // 4 bytes  
    char b[3]; // 3 bytes  
};
```

The answer

```
struct myStruct {  
    char a;    // 1 byte  
    int32_t i; // 4 bytes  
    char b[3]; // 3 bytes  
};
```

0:	a	
1:	padding	
2:	padding	
3:	padding	(Rule 2: sizeof(i) = 4, so i must be 4-aligned)
4:	i	
5:	i	
6:	i	
7:	i	
8:	b[0]	
9:	b[1]	
10:	b[2]	
11:	padding	(Rule 1: sizeof(s) must be a multiple of 4)

The answer

```
struct myStruct {  
    char a;    // 1 byte  
    int32_t i; // 4 bytes  
    char b[3]; // 3 bytes  
};
```

A: This struct takes 12 bytes of memory! (not 8)

0:	a	
1:	padding	
2:	padding	
3:	padding	(Rule 2: sizeof(i) = 4, so i must be 4-aligned)
4:	i	
5:	i	
6:	i	
7:	i	
8:	b[0]	
9:	b[1]	
10:	b[2]	
11:	padding	(Rule 1: sizeof(s) must be a multiple of 4)

Why do we need alignment?

<https://stackoverflow.com/questions/38875369/what-is-data-alignment-why-and-when-should-i-be-worried-when-typecasting-pointers>

Taken from the above link (note: this is a simplification)

For example I have two different pieces data to store: data1: "ab" data2: "cdef"

So without alignment the memory will look like : |a b c d| |e f 0 0|

How many read cycles to read data1? How about data 2?

Instead, if we use alignment : |a b 0 0| |c d e f|. **How many cycles for data1? And how many for data 2?**

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant byte on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?		
What is the 4-byte integer at location 0x1004?		

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?		
What is the 4-byte integer at location 0x1004?		

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?	0xBE	
What is the 4-byte integer at location 0x1004?		

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 **BE** 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?	0xBE	0xBE
What is the 4-byte integer at location 0x1004?		

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?	0xBE	0xBE
What is the 4-byte integer at location 0x1004?		

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?	0xBE	0xBE
What is the 4-byte integer at location 0x1004?	0x89ABCDEF	

Endianness

The **endianness** of an architecture is the way it represents data.

Big: Most significant digit on the lowest address. Makes more sense to humans.

Little (default in CPSC 313): Most significant digit on the highest address. More commonly seen.

0x1000: 01 23 45 67 89 AB CD EF

0x1008: 03 69 BE 25 8C F1 47 AD

	Big-endian	Little-endian
What is the 1-byte integer at location 0x100A?	0xBE	0xBE
What is the 4-byte integer at location 0x1004?	0x89ABCDEF	0XEFCDAB89

Endianness

0x1000 01 23 45 67 **89 AB** CD EF

0x1008 03 69 BE 25 **8C F1** 47 AD

What is myStruct if read at location 0x1004?

Size of myStruct: 12 bytes

Big-endian:

Little-endian:

```
struct myStruct {  
    char a[2]; // 2 bytes  
    int32_t b; // 4 bytes  
    int16_t c; // 2 bytes  
};
```


Endianness

0x1000 01 23 45 67 **89 AB** CD EF

0x1008 **03 69 BE 25** **8C F1** 47 AD

What is myStruct if read at location 0x1004?

Size of myStruct: 12 bytes

Big-endian:

a: [0x89, 0xAB], b: 0x0369BE25, C: 0x8CF1

Little-endian:

a: [0x89, 0xAB], b: 0x25BE6903, C: 0xF18C

```
struct myStruct {  
    char a[2]; // 2 bytes  
    int32_t b; // 4 bytes  
    int16_t c; // 2 bytes  
};
```

Endianness does not change how items are placed in struct / array like structures.

Part 3:

Pointers (segfaults, dangling ptrs, mem leaks, stack vs heap memory, type casting)

Stack vs Heap Memory Differences

- Global variables are **allocated** (given a spot) at **compile time** (**static** memory allocation)
 - In general, the layout of stack frames are decided at compile time
 - parameters
 - local variables in each function
 - return addresses
 - However, the stack itself is not allocated at compile time. We don't know how many stack frames there will be at any time, nor what the heap will look like, until **runtime** (**dynamic** memory allocation)
 - On the heap, we use malloc/free to manually manage memory
-
- Fun fact: a stack is a stack data structure; a heap is not a heap data structure.

Stack vs Heap Memory Differences (Summary)

	Static or Global	Stack (Local)	Heap (Dynamic)
Layout (Align, Offset, Size)	Compile Time	Compile Time	Run Time
Address	Compile Time	Run Time	Run Time
Data	Compile Time/Run Time	Run Time	Run Time
Per Thread	No	Yes	No
Where?	Executable	OS Provided Memory	OS Provided Memory

Pointers Basics

- `*` : dereference operator, which is also used in pointer declaration
- `&` : (ampersand) "address of" operator
- remember: pointers are addresses! (the address of the thing it is pointing to)
 - this is not the address of the pointer

Initialize a pointer:

```
int x = 10;  
int * p = &x; // pointer p points to x (p contains the address of x in memory)  
*p = 20;      // now if we ask for the value of x, it will be 20
```

Pointer Arithmetic

Q: What is the final value of ptr?

```
int32_t* ptr = 0x7ffffffe0;  
ptr++;
```

Pointer Arithmetic

Q: What is the final value of ptr?

```
int32_t* ptr = 0x7ffffffe0;  
ptr++;
```

A: 0x7ffffffe4

Adding and subtracting an integral value to pointers works like a data index!

- `&ptr[i] == (ptr + i) == (((char*)ptr) + i * sizeof(*ptr))`
- Useful for iterating through data without having to store index separately
- If you need byte arithmetic use something like `char*` and cast back to original type to access

Pass by Value vs Pass by Reference

```
void my_swap1(int* i, int* j) {  
    int* temp = i;  
    i = j;  
    j = temp;  
}
```

```
void my_swap2(int i, int j) {  
    int temp = i;  
    i = j;  
    j = temp;  
}
```

Q: What function will correctly swap a and b?

```
int a = 3;  
int b = 4;  
my_swapN(&a, &b);
```


Pass by Value vs Pass by Reference

```
void my_swap1(int* i, int* j) {  
    int* temp = i;  
    i = j;  
    j = temp;  
}
```

```
void my_swap2(int i, int j) {  
    int temp = i;  
    i = j;  
    j = temp;  
}
```

Q: What function will correctly swap a and b?

```
int a = 3;  
int b = 4;  
my_swapN(&a, &b);
```

A: Neither! We are assigning local variables in either case. We need to dereference the pointers!

Correct Swap Function

```
void my_swap3(int* i, int* j) {  
    int temp = *i;  
    *i = *j;  
    *j = temp;  
}
```

```
int a = 3;  
int b = 4;  
my_swap3(&a, &b);
```



Segmentation Faults

- When you try to access memory that isn't supposed to be accessed
- Common case: you forget to initialize your pointers and dereference NULL

```
char *p = NULL;  
*p = 'a';           // Segmentation fault: core dumped
```

Memory Leaks

Technically memory-safe

When you allocate memory on the heap but do not deallocate it

```
int* ptr = (int *) malloc(sizeof(int));  
  
*ptr = 10;  
  
return 0; // return without freeing ptr  
          // we should call free(ptr); before returning
```

Memory Leaks

Technically memory-safe

When you allocate memory on the heap but do not deallocate it

Q: Why is this a problem?

```
int* ptr = (int *) malloc(sizeof(int));
```

```
*ptr = 10;
```

```
return 0; // return without freeing ptr  
          // we should call free(ptr); before returning
```

Memory Leaks

Technically memory-safe

When you allocate memory on the heap but do not deallocate it

Q: Why is this a problem? A: It's wasteful

```
int* ptr = (int *) malloc(sizeof(int));
```

```
*ptr = 10;
```

```
return 0; // return without freeing ptr  
          // we should call free(ptr); before returning
```

Memory Leaks

Technically memory-safe

When you allocate memory on the heap but do not deallocate it

Q: Why is this a problem? A: It's wasteful

Q: Can we have memory leaks using stack memory?

```
int* ptr = (int *) malloc(sizeof(int));
```

```
*ptr = 10;
```

```
return 0; // return without freeing ptr  
          // we should call free(ptr); before returning
```

Memory Leaks

Technically memory-safe

When you allocate memory on the heap but do not deallocate it

Q: Why is this a problem? A: It's wasteful

Q: Can we have memory leaks using stack memory? A: No.

```
int* ptr = (int *) malloc(sizeof(int));
```

```
*ptr = 10;
```

```
return 0; // return without freeing ptr  
          // we should call free(ptr); before returning
```


Dangling Pointers

Anyone remember lab_debug from 221?

When a pointer points to something that "no longer exists"

Q: Why is this a problem?

```
char *p = NULL;
{
    char c;
    p = &c;
}
// dangling ptr: c does not exist outside the {braces}

// here
```

Dangling Pointers

Anyone remember lab_debug from 221?

When a pointer points to something that "no longer exists"

Q: Why is this a problem? A: Allows undefined behaviour, use-after-free, or (best case) segfault

```
char *p = NULL;
{
    char c;
    p = &c;
}
// dangling ptr: c does not exist outside the {braces}

// here
```

Type Casting (Pointers)

```
int32_t i = 0x12345678;  
char* cp = (char *)&i;
```

```
*cp = ? // 0x78
```



```
char c = 0xAB;  
int32_t* ip=(int32_t *)&c;
```

```
*ip = ? // 0x?????AB  
// UB
```



Part 4:

C Programs (arguments, debugging)

Writing a Main Function

<https://stackoverflow.com/questions/3024197/what-does-int-argc-char-argv-mean>

argv and **argc** are how command line arguments are passed to `main()` in C.

argc will (in practice) be 1 plus the number of actual arguments in **argv**, as virtually all implementations will prepend the name of the program to the array.

Example of a Main Function

```
int main(int argc, char* argv[]) {  
    printf("Have %d arguments:\n", argc);  
    for (int i = 0; i < argc; ++i) {  
        printf("%s\n", argv[i]);  
    }  
}
```

```
gcc test.c -o test
```

Example of a Main Function

```
int main(int argc, char* argv[]) {  
    printf("Have %d arguments:\n", argc);  
    for (int i = 0; i < argc; ++i) {  
        printf("%s\n", argv[i]);  
    }  
}
```

```
gcc test.c -o test
```

Input:

```
./test a1 b2 c3
```

Q: What are in argc/argv?

Example of a Main Function

```
int main(int argc, char* argv[]) {  
    printf("Have %d arguments:\n", argc);  
    for (int i = 0; i < argc; ++i) {  
        printf("%s\n", argv[i]);  
    }  
}
```

```
gcc test.c -o test
```

Input:

```
./test a1 b2 c3
```

Q: What are in argc/argv?

A:

argc = 4

argv[0] = "./test"

argv[1] = "a1"

argv[2] = "b2"

argv[3] = "c3"

Printf Statements

%d : Integer

%f : Float

%c : Character

%s : String

%p : Pointer

%x : Hexadecimal

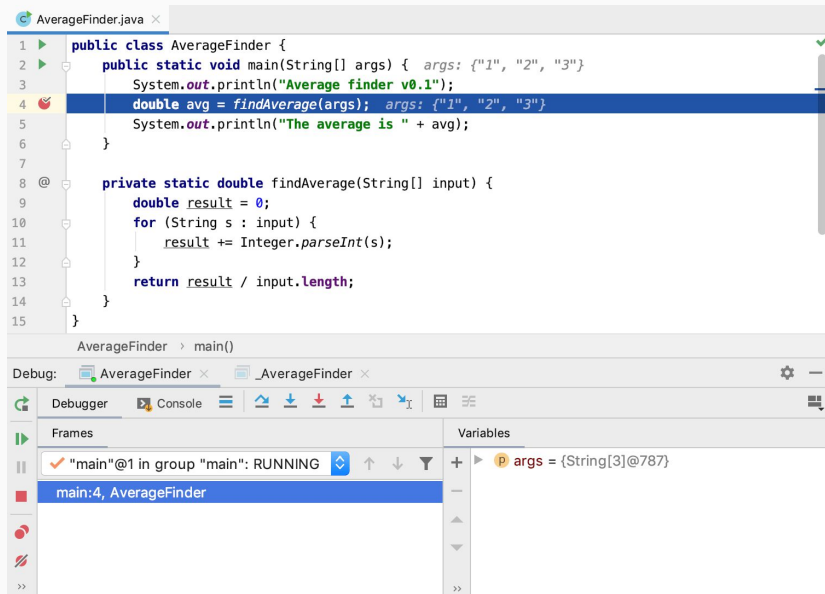
Basic Syntax:

```
char stringVar[] = "Hello World!";  
printf("Value of stringVar: %s\n", stringVar)
```

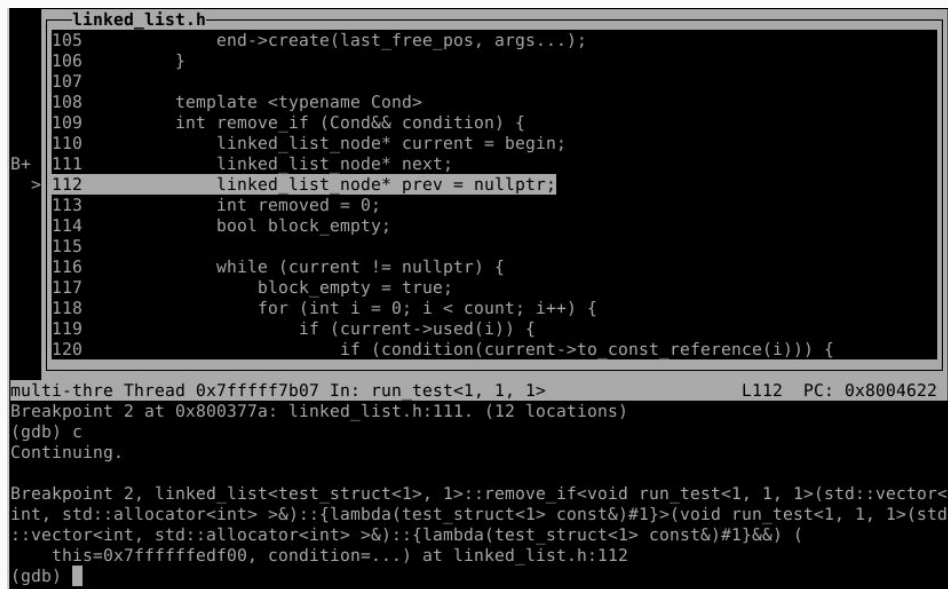
Debugging C programs

1. You should compile your program with `-g` flag for easy debugging
 - a. `gcc -g main.c -o main`
 - b. `make` (if there's already a makefile)
2. One way to run gdb is by passing the your executable's name as its argument:
 - a. `gdb main`
3. Then, you can run your program using `run`
 - a. It'd be helpful to run `layout next` to have code view in your terminal

Debugging C programs



IntelliJ built-in debugger for Java



C program debugger: GDB

GDB Common Commands

`start` / `r(un)`

`p(rint)`

`b(reak)`

`c(ontinue)`

`watch`

`n(ext)`

`bt` backtrace

`s(tep)`

`clear`

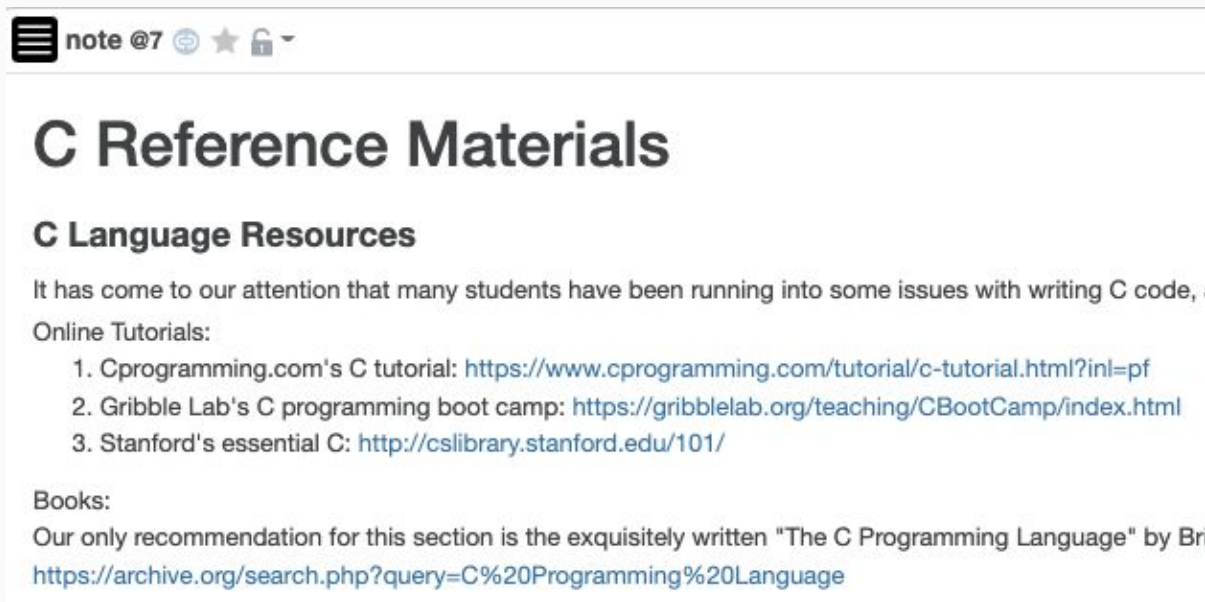
`quit`

Use GDB Reference (Google) for more commands!

You will not learn these commands by looking at the slide. You will learn them by using them.

Demonstration time! (with code snippet)

For Other Resources, See Piazza!



The screenshot shows a Piazza note interface. At the top, there's a header bar with a menu icon, the text 'note @7', and icons for sharing, starring, and locking. Below this is the main content area. The title 'C Reference Materials' is in a large, bold, dark font. Underneath it is a subtitle 'C Language Resources' in a smaller, bold, dark font. The body of the note starts with a paragraph: 'It has come to our attention that many students have been running into some issues with writing C code, ...'. This is followed by a section header 'Online Tutorials:' and a numbered list of three items: 1. 'Cprogramming.com's C tutorial: <https://www.cprogramming.com/tutorial/c-tutorial.html?inl=pf>', 2. 'Gribble Lab's C programming boot camp: <https://gribblelab.org/teaching/CBootCamp/index.html>', and 3. 'Stanford's essential C: <http://cslibrary.stanford.edu/101/>'. Below the list is a section header 'Books:' followed by a paragraph: 'Our only recommendation for this section is the exquisitely written "The C Programming Language" by Bri ...'. At the end of the paragraph is a link: '<https://archive.org/search.php?query=C%20Programming%20Language>'.

note @7

C Reference Materials

C Language Resources

It has come to our attention that many students have been running into some issues with writing C code, ...

Online Tutorials:

1. Cprogramming.com's C tutorial: <https://www.cprogramming.com/tutorial/c-tutorial.html?inl=pf>
2. Gribble Lab's C programming boot camp: <https://gribblelab.org/teaching/CBootCamp/index.html>
3. Stanford's essential C: <http://cslibrary.stanford.edu/101/>

Books:

Our only recommendation for this section is the exquisitely written "The C Programming Language" by Bri ...

<https://archive.org/search.php?query=C%20Programming%20Language>

Also:

- <https://en.cppreference.com/w/c/language> (language reference, standard C)
- <https://www.gnu.org/software/gnu-c-manual/gnu-c-manual.html> (GNU C reference)
- https://www.onlinegdb.com/online_c_compiler (minimal online compiler for C)

Bonus: Right-left rule for reading C declarations

See https://cseweb.ucsd.edu/~gbournou/CSE131/rt_lt.rule.html

```
int *(*mystery())();
```

What is this? A function pointer? A function pointer pointer? What does it return?

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

```
int *p[]
```

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

```
int *p[]
```

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

`int *p[]`



"p" is...

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

`int *p[]`



"p" is...

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

`int *p[]`



"p" is (an)...

... array of...

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything

`int *p[]`



"p" is (an)...

... array of...

Right-left rule

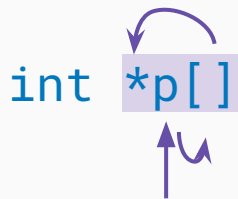
Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything



`int *p[]`

"p" is (an)...

... array of...

... pointer to...

Right-left rule

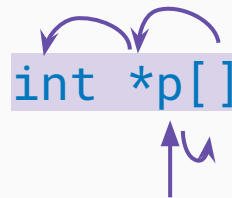
Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read to the **right** until you hit the end, or a closing paren `)`
3. Read to the **left** until you hit the end, or an opening paren `(`

Then repeat steps 2 and 3 until you've read everything



`int *p[]`

"p" is (an)...

... array of...

... pointer to...

int.

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2

`int *(*mystery())();`

mystery is



Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2

`int *(*mystery())();`

mystery is function returning

Right-left rule

Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2

`int *(*mystery())();`

mystery is function returning pointer to

Right-left rule


Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2



mystery is function returning pointer to function returning

Right-left rule


Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2



mystery is function returning pointer to function returning pointer to

Right-left rule


Read `*` as "pointer to"

Read `[]` as "array of"

Read `()` as "function returning"

1. Start at the identifier
2. Read right
3. Read left

// goto 2



mystery is function returning pointer to function returning pointer to int!

Any questions?