

# CPSC313: Computer Hardware and Operating Systems

Unit 4: File Systems  
Fun with File Descriptors



# Admin

- Final examination
  - Reserve your time of PrairieTest if you haven't already done it.
- Tutorial 8 is this week
- Lab 8 has been released and is due Sunday November 17th.
- Code for today is in the course code repository:  
[4.2-fun-with-file-descriptors](#)

# Where we are

- Unit Map:
  - P18: File Systems APIs and How disks work
  - 4.1. Using File Systems APIs
  - P19: File descriptors
  - **4.2. File descriptors management**
  - P20: File Systems implementation overview
  - 4.3. How we represent files

# Today

- Learning Outcomes:
  - Review of file descriptor management data structures.
  - Redirect standard in, standard out, and standard error from within a program.
  - Describe how the shell implements redirection to implement commands that do the following:
    - `./cmd > something`
    - `./cmd >& something`
    - `./cmd < something`
    - `./cmd >> something`
  - Map redirection into operations on file descriptors, the file descriptor table, open file structures and vnodes.

# Questions about how a file system should behave (all answered in 1969)\*

- If you and I are both allowed to read a file in the file system, should we be able to **share the file's data**? **Yes!**
- If **two processes** are reading (writing) the same file, should they be using the same file offset? **No!**
- If a process **opens the same file twice**, should it have one file descriptor or two? In either case, should the two opens share an offset? **Two!**  
**No!**

# Questions about how a file system should behave (all answered in 1969)\*

- If **two threads** are using the same file descriptor, should they be using the same file offset? **Yes!**
- If one process (the parent) creates another process (a child), should the child **inherit** the parent's **file descriptors**? Will they use the **same offset**? **Yes!**  
**Yes!**

# Putting this all together

Per-process


Shared across processes

Shared across processes


File descriptor table

0	
1	
2	
⋮	
N	

Open File Table



Vnode Table



# Recall: Parent and Child have the same FDs


Parent File descriptor table

0	
1	
2	
⋮	
N	


Child File descriptor table

0	
1	
2	
⋮	
N	

Open File Table



Vnode Table



Demo ptest and forktest



# Fork: `pid_t fork(void)`

- **Creates a new process** almost identical to the process that called it.
  - The parent's return value will be the **pid of the child process**.
  - The child process's return value will be **0**.
- Immediately after the fork call, both processes should check their return values and proceed accordingly.
  - The child process often calls a member of the **exec** family.
  - The parent might wait on the child or do something else.

# Exec and friends

- `int execve(const char *path, char *const argv[], char *const envp[])`
  - Typically `path` is the pathname of a command you want to execute, e.g., `./myprog`, `/bin/ls`
  - `argv` is an argument vector -- it is what is passed to `main`, e.g.,  
`int main(int argc, char *argv[])`
  - `envp` is an environment: the environment is a set of name/value pairs that are frequently used to communicate 'environmental' information to processes: where should the process look for commands, what OS are we running, etc.

# Exec and friends


- `int execvp(const char *path, char *const argv[])`
  - If the parameter path does not begin with “/”, or “.”, or “..”, then `execvp` searches for parameter path in each directory listed in the `PATH` environment variable (just like the shell does).

# Recall: Another way to get a refcount of 2


Parent File descriptor table

0	
1	
2	
⋮	
N	

Open File Table



Vnode Table



The **dup** (**dup2**) system call(s) “duplicates a file descriptor”

```
int dup(int fd);  
int dup2(int fd1, int fd2);
```

Demo fdtest and duptest

# When might you use dup?

Modes are traditionally written as OCTAL (base 8) digits to correspond to the 3 bits for each of owner, group, and world. A leading 0 (zero) means the number is octal

- Redirecting standard out:
  - When you want to capture a program's output, you can use:  
`ls > my-stdout-file.txt`
- How does this work?
  - The ugly brute force way (building into a program):  

```
int fd = open("my-stdout-file.txt", O_WRONLY|  
O_CREATE|O_TRUNC, 0644);  
dup2(fd, STDOUT_FILENO); // Closes old standard out
```

# When might you use dup?

- Fun fact about dup:
  - If you dup into an fd that is open (e.g., `STDOUT_FILENO`), this automatically closes the file that had been using that fd.
  - Even though standard IO (e.g., `printf`) is not a system call, when we change where stdout writes, it also changes where `printf` output goes.

# The shell implements redirection for you

- Redirecting standard out:
  - When you want to capture a program's output, you can use:

```
ls > my-stdout-file.txt
```

```
pid_t child_pid = fork();
if (child_pid == 0) {
    /* In the child. */
    int fd = open("my-stdout-file.txt", O_WRONLY|O_CREATE|O_TRUNC, 0644);
    dup2(fd, STDOUT_FILENO);
    close(fd);
    execv("ls", argv); // Changes child to execute 'ls' instead of shell
} else if (child_pid > 0) {
    This is the parent
} else {
    This is an error
}
```


# Redirecting a command in the shell (1)

The red font tells you what happens between this slide and the next one.


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table



Vnode Table



Now, what happens when the parent forks?  
`child_pid = fork();`

Everything is a file!




# Redirecting a command in the shell (2)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table



Vnode Table



Child (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Next, the child opens a file:


```
newfd = open("stdout-out-file.txt",  
             O_WRONLY | O_TRUNC | O_CREAT, 0644);
```

# Redirecting a command in the shell (3)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table





Vnode Table



Child File (shell) descriptor table

0	
1	stdout
2	
3	



Now the child calls dup2  
`dup2(newfd, STDOUT_FILENO);`

# Redirecting a command in the shell (4)

Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table

Position  
Refcnt=2

Vnode Table

Console

Child File (shell) descriptor table

0	
1	stdout
2	
3	

Position  
Refcnt=2


stdout-out-file  
Perms  
Size  
Type  
blocks

# Redirecting a command in the shell (4)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table





Vnode Table



Child File (shell) descriptor table

0	
1	stdout
2	
3	




# Redirecting a command in the shell (4)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table





Vnode Table



Child File (shell) descriptor table

0	
1	stdout
2	
3	




Now it calls close on newfd  
`close(newfd);`

# Redirecting a command in the shell (5)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table





Vnode Table



Child File (shell) descriptor table

0	
1	stdout
2	
3	




Now it calls `execv`  
`execv("ls", argv);`

# Redirecting a command in the shell (6)


Parent (shell) File descriptor table

0	
1	stdout
2	
⋮	
N	

Open File Table





Vnode Table



Child File (ls) descriptor table

0	
1	stdout
2	
3	



# Other fun facts about file descriptors and dup

- Redirect both standard out & standard error:
  - `mycmd >& stdout-and-stderr-file.out`
- Redirect standard in:
  - `mycmd < input-file.txt`
- Redirect standard out, but also let it display on the screen
  - `mycmd | tee stdout-file.out`
- Redirect standard out and append to an existing file
  - `mycmd >> stdout-file.out`