

CPSC313: Computer Hardware and Operating Systems

Unit 5: Virtual Memory

(5.1) Virtual Memory: Achieving Process Isolation



Admin

- Quiz 4 viewings and retake this week
- You should already have a final exam time!
- Labs 9 and 10 are “in progress”. Check PrairieLearn for deadlines!




Today: The start of our last unit, and... MAGIC!

- Learning Outcomes
 - Define:
 - Process
 - Address space
 - Process isolation
 - Virtual memory
 - Explain how virtual memory provides process isolation.
- Reading
 - Section 8.2

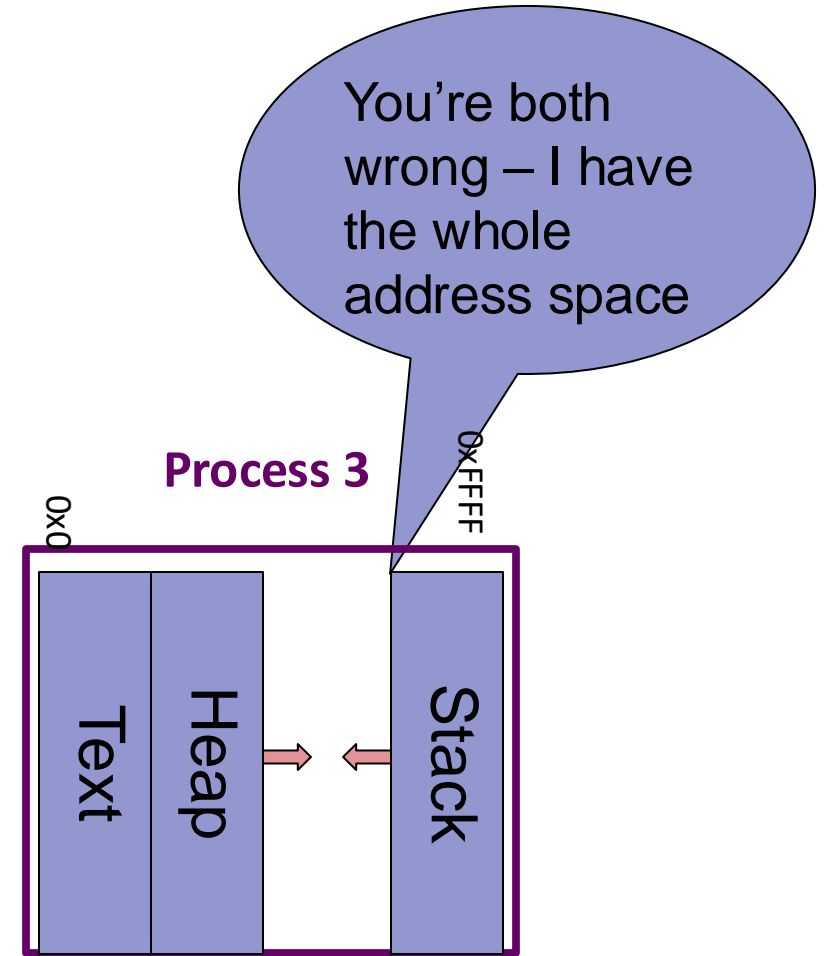

Fundamentals:

- Programs run in different **processes**.
- Each process has its own **address space**.
- Address spaces provide **process isolation**.
- Process isolation means:
 - Anything one process does should not affect what another process does, unless the processes agree (e.g., set up a communication channel).
 - Each process behaves as if it controls the entire machine's resources.

Recall way back from Unit 1



The Illusion



It's all a lie! (illusion)

There are two different kinds of address:

Physical addresses *

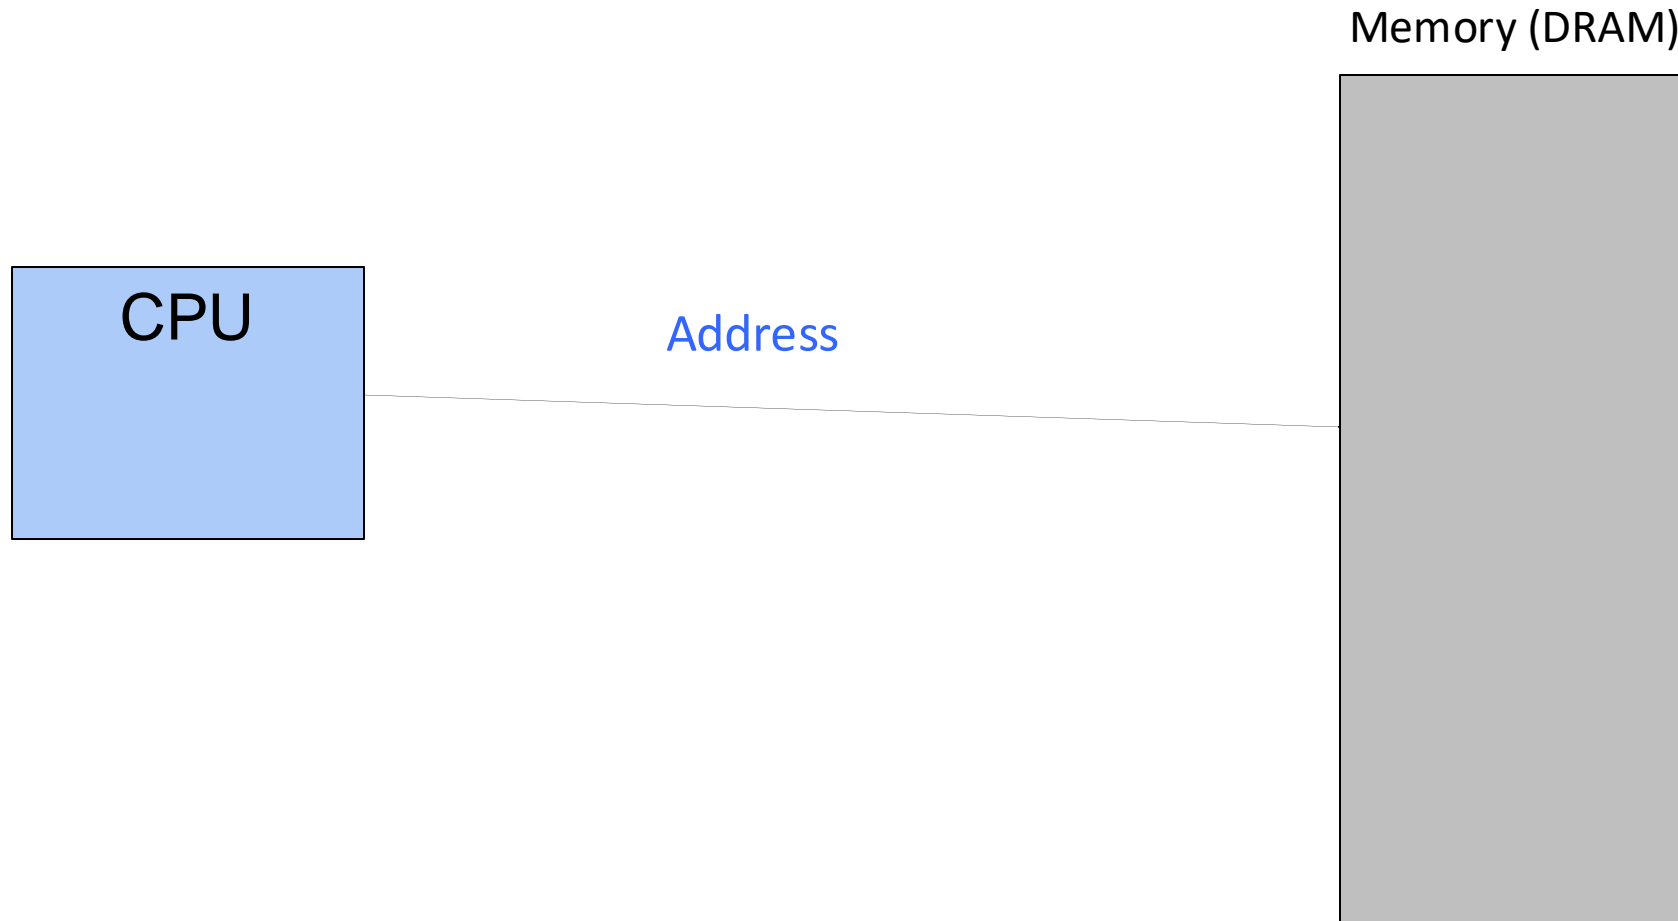
- Refer to specific locations in memory (DRAM).

Virtual addresses


- These are the addresses that programs use.
- The hardware (HW) & software (SW) collaborate to **map from virtual addresses to physical address**.
- A machine can have less physical memory than the size of the virtual address space!
- You can have a machine with a 64-bit architecture supporting program with virtual addresses measuring in Petabytes, without having a machine with that much memory.

* It turns out that this is all a lie too: out of scope for this course, but if you're interested, check out <https://www.usenix.org/system/files/conference/hotos15/hotos15-paper-gerber.pdf>

How?



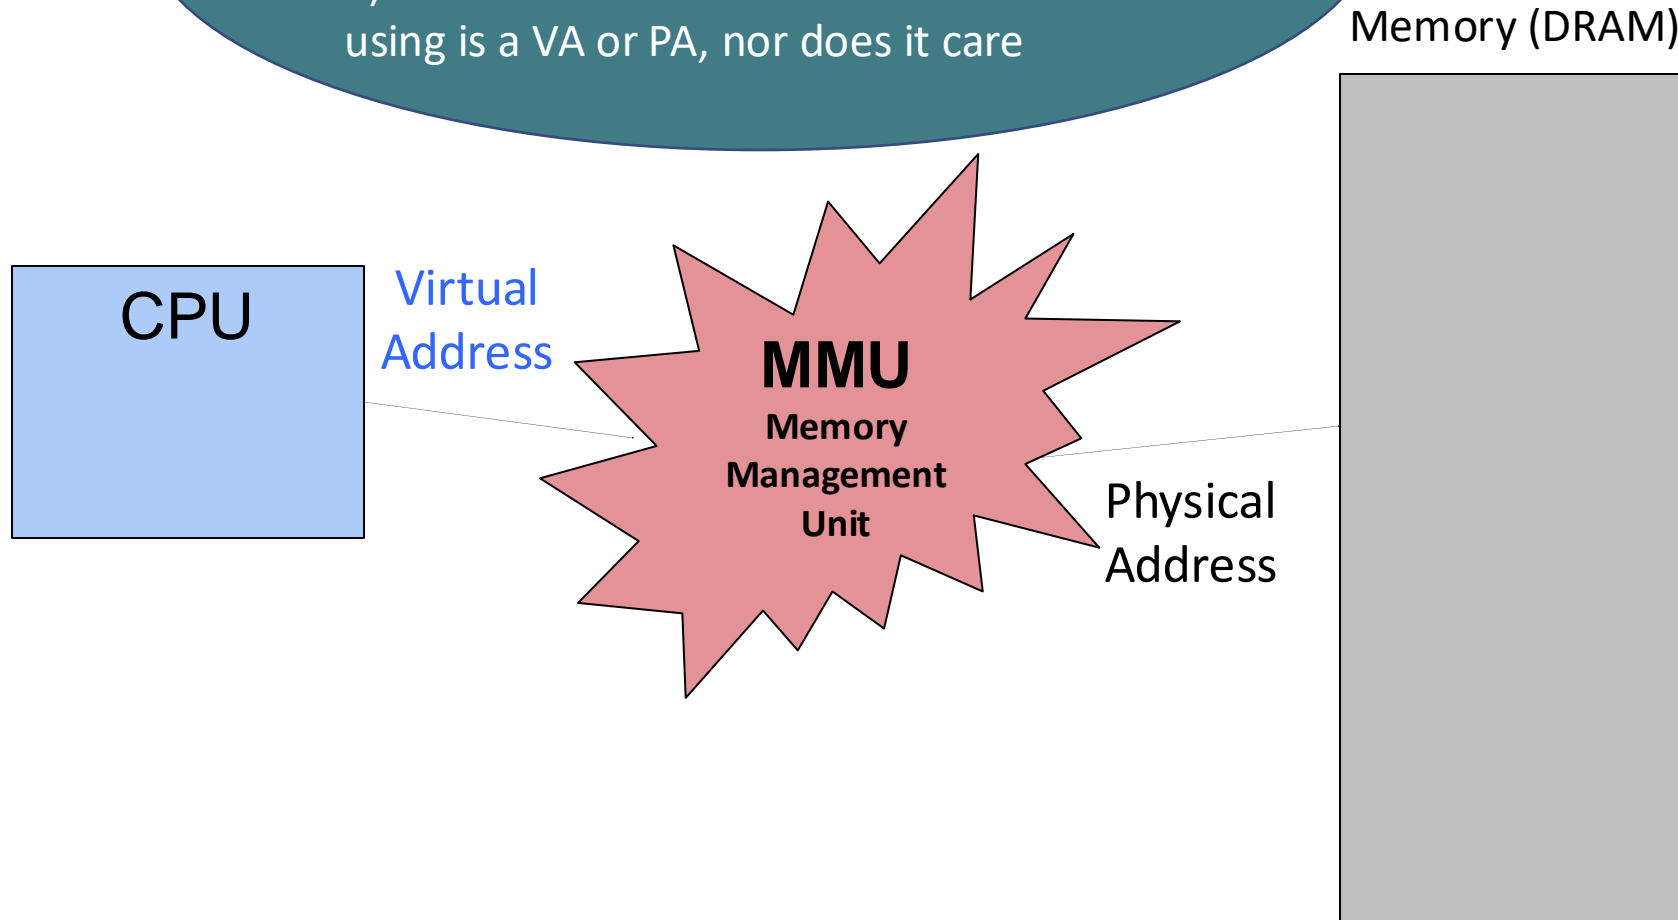
How? – Insert some “magic”



How?

Observe:

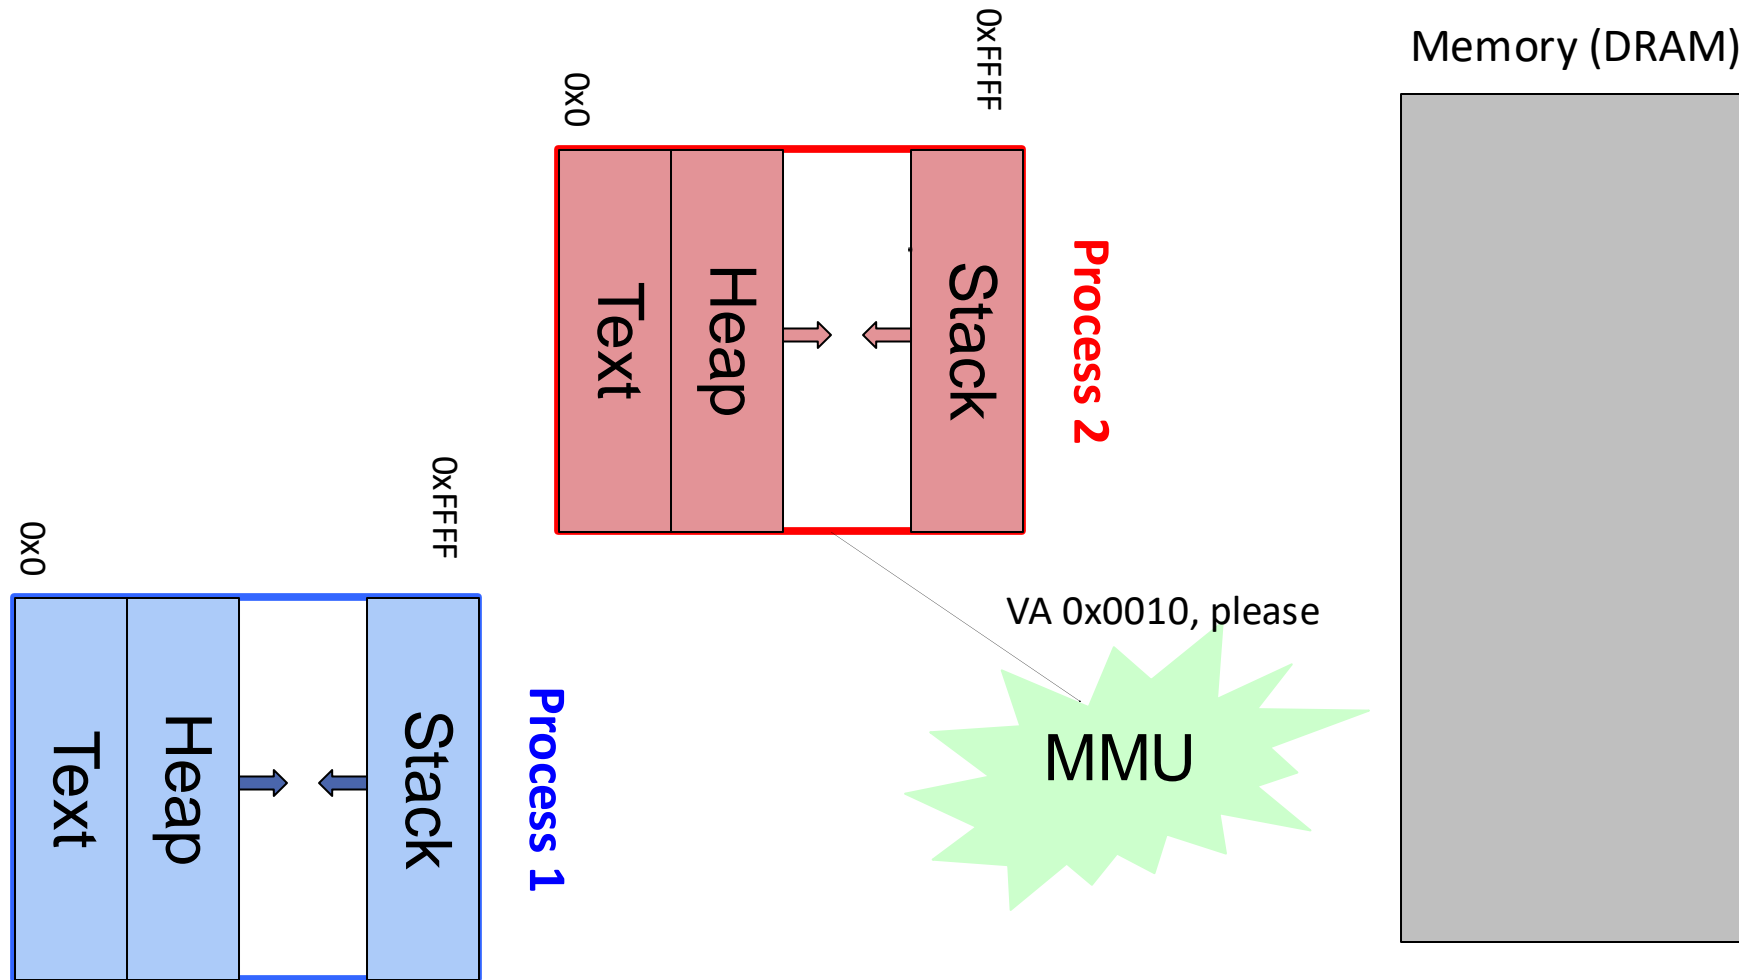
- 1) If MMU takes the VA and simply passes it through then the $VA == PA$
- 2) CPU doesn't know if the address it is using is a VA or PA, nor does it care



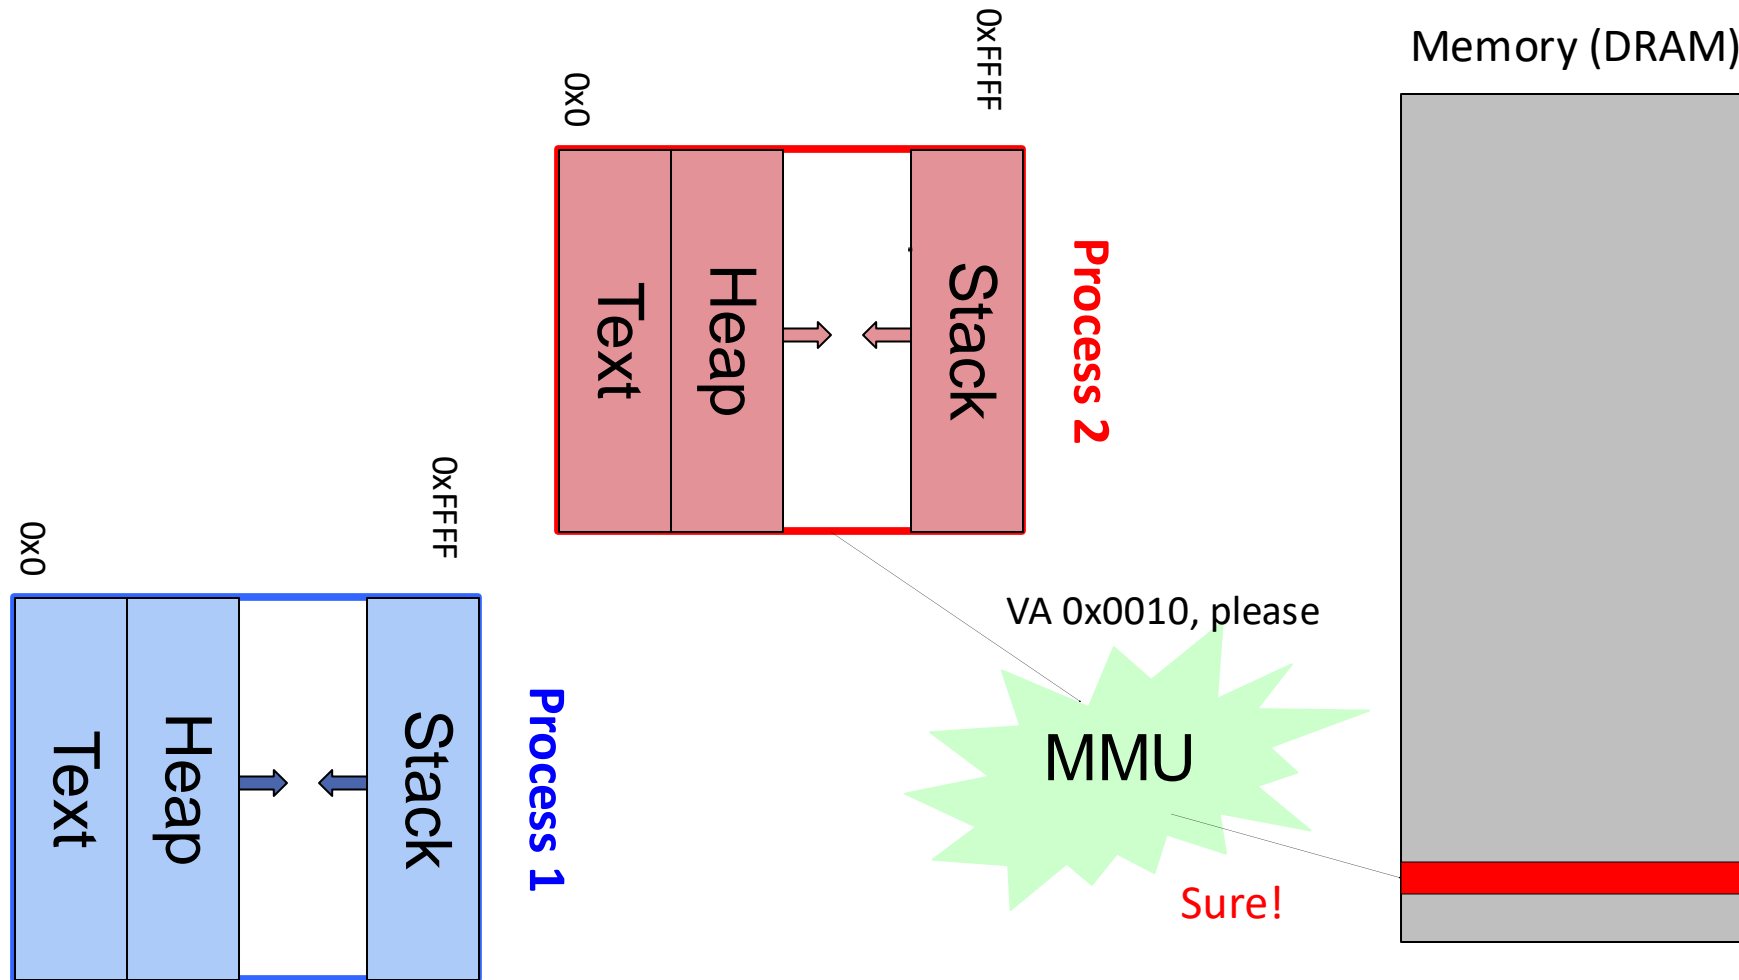
How does this help us?




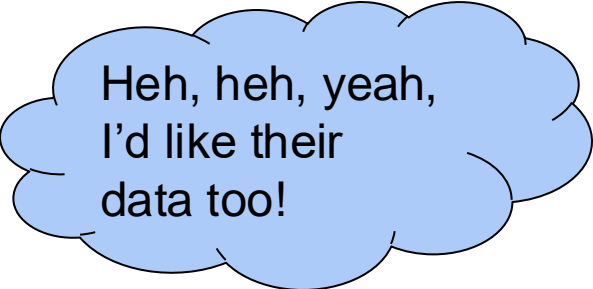
How VM Provides Process Isolation/Protection




How VM Provides Process Isolation/Protection



Protection



Process 1



Process 2


VA 0x0010, please

MMU


Sure!



Protection



Process 1



Process 2


VA 0x0010, please

VA 0x0010, please

MMU

Sure!

Sure!



VM: A Hardware/Software Partnership

- We need hardware support to provide virtual memory. Why?


Speed! Invoking the operating system (OS) on every address access would be **much** too slow.

- The Hardware
 - Provides a fast mechanism to map a virtual address to a physical address.
- The Software (OS)
 - Sets up the mappings that the hardware uses.
 - Manages the allocation of physical memory.
 - Implements policies about how memory is shared.

System calls that manipulate address spaces

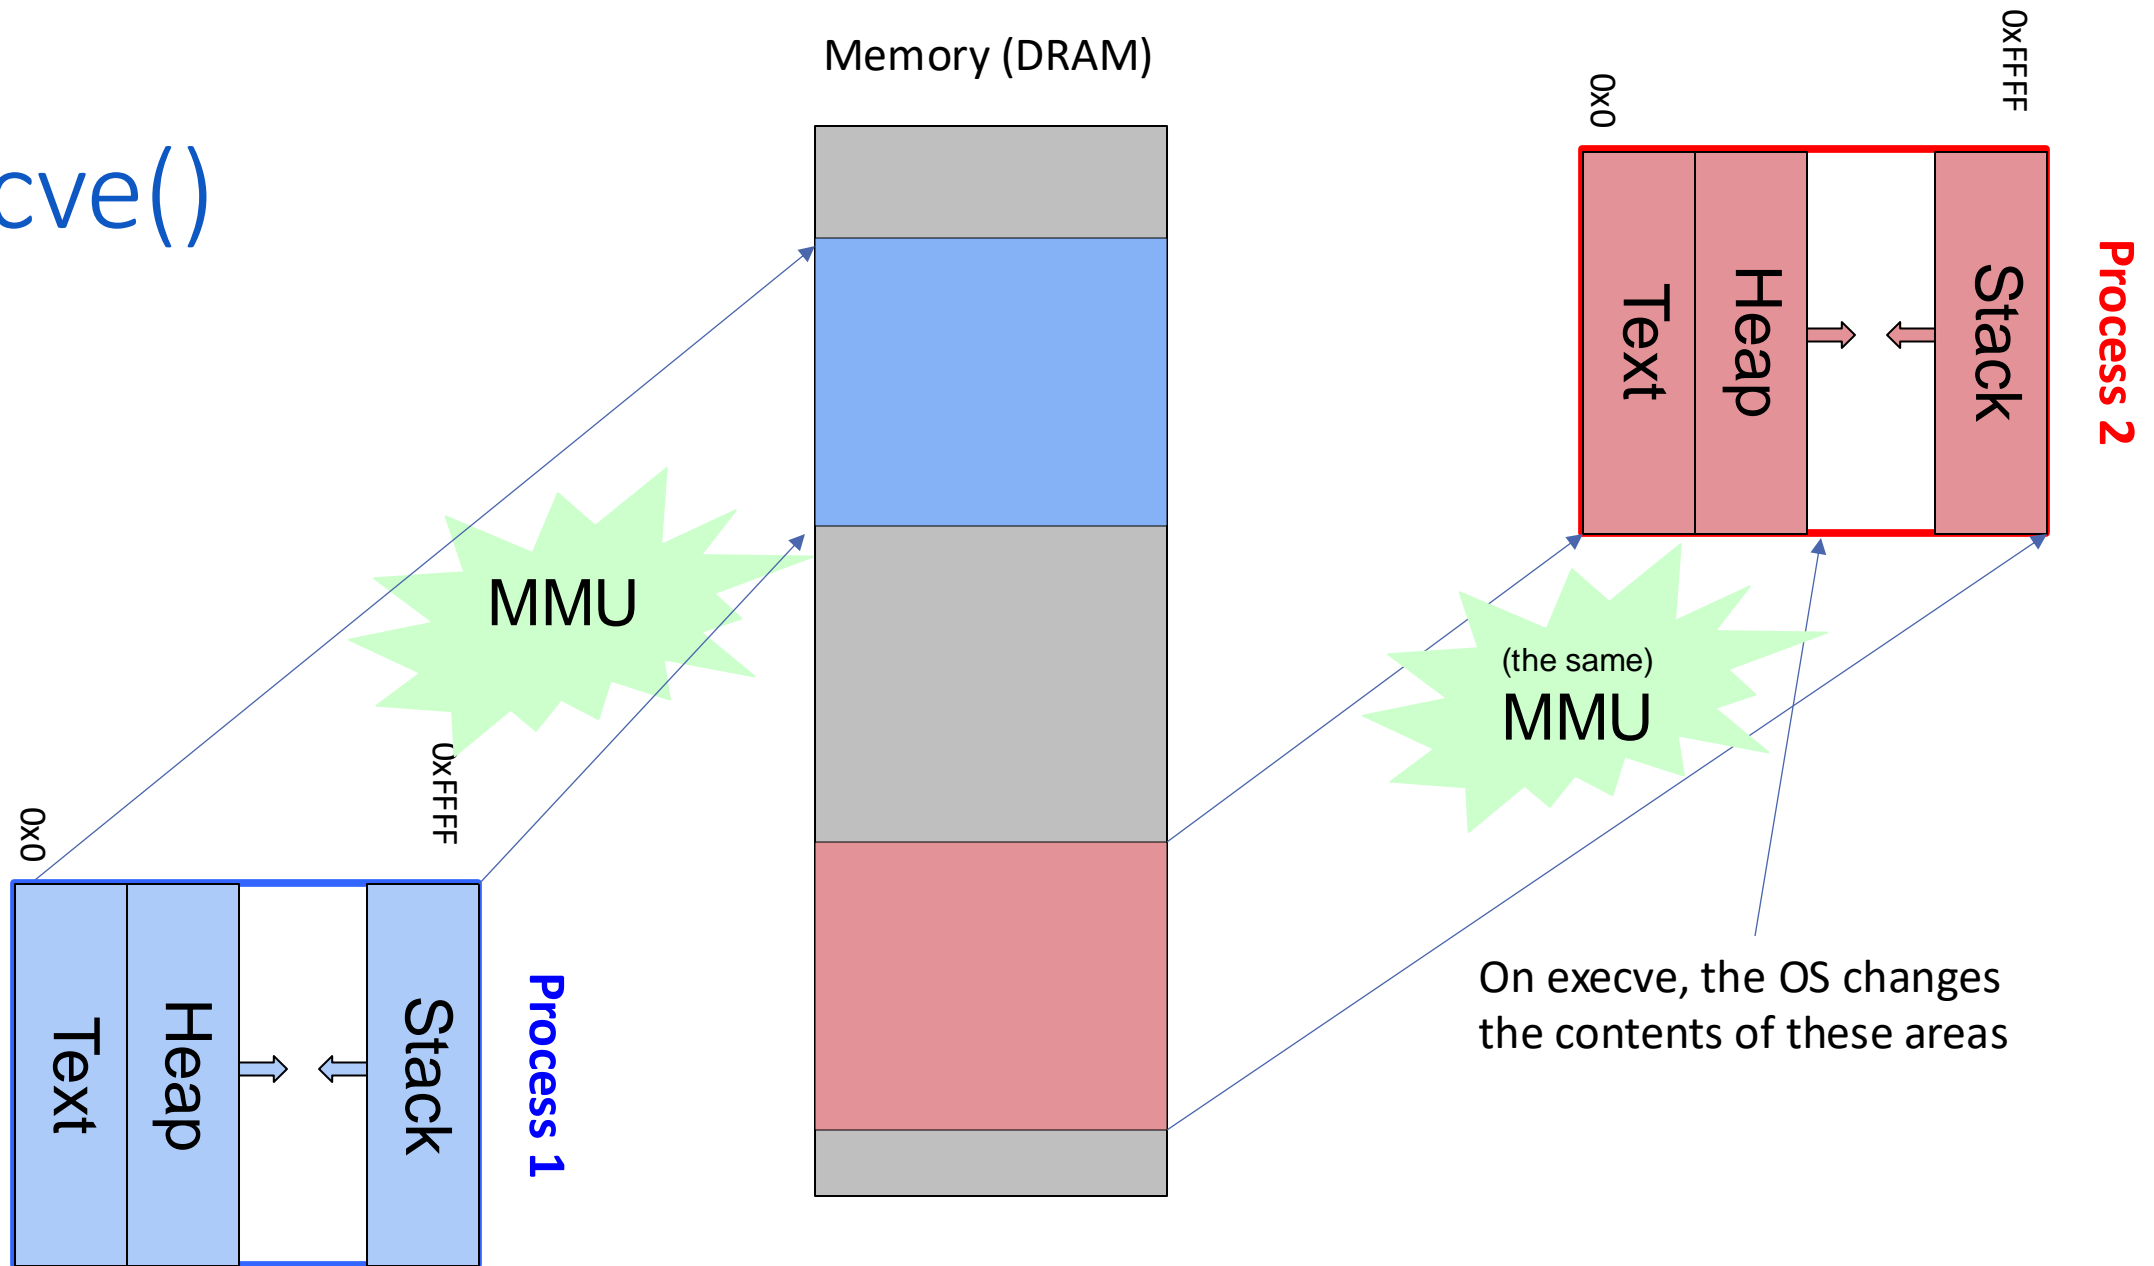
- **`pid_t fork(void)`**
 - Creates a new process (called a child process)
 - The child is an exact (almost) copy of the parent
- **`int execve(const char *path, char *const argv[], char *const envp[])`**
 - Replace the current process image with the program image stored in file.
 - There is a whole family of function calls that all ultimately invoke the `execve` system call.
 - `execl, execlp, execl, exect, exectp, exectp`
 - `execv, execvp, execvpe` (or `execvP` on a MAC, FreeBSD)

pid = fork();



This all happens *instantly*?!
Something here is a lie. Start pondering what.
We'll talk more in a few lectures.

execve()



Exec and friends

- **`int execve(const char *path, char *const argv[], char *const envp[])`**
 - Typically `path` is the pathname of a command you want to execute, e.g.,
`./myprog, /bin/ls`
 - `argv` is an argument vector -- it is what is passed to `main`, e.g.,
 - `int main(int argc, char *argv[])`
 - `envp` is an environment: the environment is a set of name/value pairs that are frequently used to communicate 'environmental' information to processes: where should the process look for commands, what OS are we running, etc.
- **`int execlp(const char *path, char *const argv[])`**
 - If the parameter `path` does not begin with `"/`, or `."`, or `.."`, then `execlp` searches for parameter `path` in each directory listed in the `PATH` environment variable (just like the shell does).



Shell demo

In-class activity – build a baby shell!

