

CPSC313: Computer Hardware and Operating Systems

Unit 5: Virtual Memory

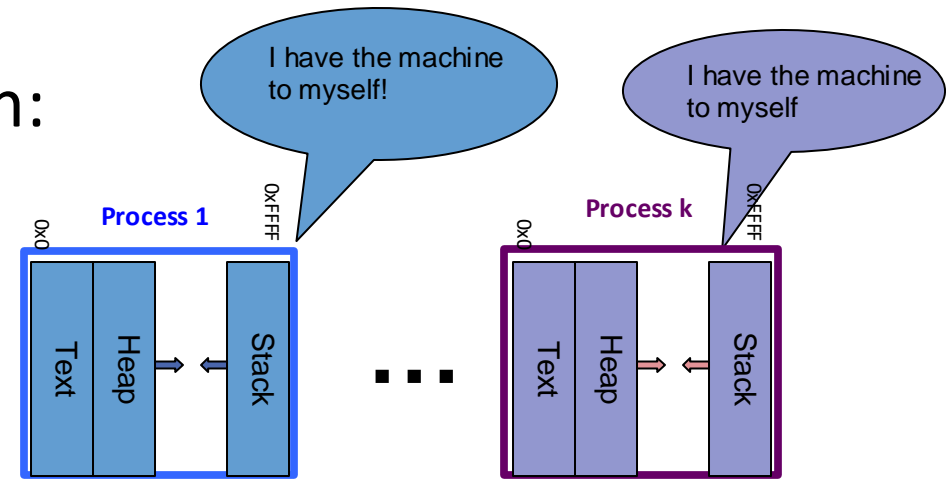
(5.3) Virtual Memory: How does it actually work?

Administrivia

- Quiz 4 retakes finishing this week
- Quiz 5 coming next week
- Lab 10 due Sunday
- Exam period office hours schedule coming soon!
- **As always: check Canvas for all details!**

Context

We want the fiction of process isolation:



How do we get there?

- Virtual memory enforces isolation; translation/faulting supports virtual addressing
- We need to trust a resource “manager”– the OS:
 - Privileged (kernel/supervisor privilege, vs. user privilege)
 - Need to safely escalate privilege

Today


- Roadmap:
 - **Virtual memory (VM)** as a case study of how the OS and hardware work together and how/when control transfers.
 - To do that, we're going to first dive into how VM works.
 - We'll figure out what problem we are trying to solve and approaches to solving it.
 - Then we'll look at existing real world solutions.
- Learning Outcomes
 - Explain why we organize memory in pages
 - Define: Translation Lookaside Buffer (TLB)
 - Figure out how many bits are necessary to represent virtual and physical page numbers

Recall: (You've seen this!)

- We ask the **hardware** to map a triple of:
 - Virtual address (VA)
 - Access mode (read, write, execute)
 - **Privilege** levelto a physical address


Recall: (You've seen this!)

- We ask the **hardware** to map a triple of:
 - Virtual address (VA)
 - Access mode (read, write, execute)
 - **Privilege** levelto a physical address
- The hardware will either:
 - Produce a **physical address**
 - **Fault**, due to:
 - The address is invalid
 - The type of access is not allowed to the memory requested
 - The process requesting access does not have the appropriate privilege level to access the memory
 - The data is not in memory (it is on disk and hasn't been loaded yet)



Recall: (You've seen this!)

- We ask the **hardware** to map a triple of:
 - Virtual address (VA)
 - Access mode (read, write, execute)
 - **Privilege** levelto a physical address
- The hardware will either:
 - Produce a **physical address**
 - **Fault**, due to:
 - The address is invalid
 - The type of access is not allowed to the memory requested
 - The process requesting access does not have the appropriate privilege level to access the memory
 - The data is not in memory (it is on disk and hasn't been loaded yet)
- When the hardware faults: software (the OS) takes over
 - pre-class material explains how



VM Mapping:

- Goal: We need a map from triples to physical address:
 - $(VA, \text{access}, \text{priv}) \Rightarrow PA \mid \text{FAULT}$
 - Physical addresses are not necessarily contiguous or ascending
 - We would like to avoid fragmentation.
- How might you implement this?
Discuss with your neighbor!

VM Mapping: The Dumbest Idea Ever

- Goal: We need a map from triples to physical address:
 - $(VA, \text{access}, \text{priv}) \Rightarrow PA \mid \text{FAULT}$
 - Physical addresses are not necessarily contiguous or ascending
 - We would like to avoid fragmentation.
- How might you implement this?
 - Store a mapping for every byte in the virtual address space
 - Critique:

VM Mapping: The Dumbest Idea Ever

- Goal: We need a map from triples to physical address:
 - (VA, access, priv) => PA | FAULT
 - Physical addresses are not necessarily contiguous or ascending
 - We would like to avoid fragmentation.
- How might you implement this?
 - Store a mapping for every byte in the virtual address space
 - Critique:
 - The mapping table would be huge!
 - And wait -- we know about caching -- could you have a different mapping for every element of a cache line??? What a terrible idea.
 - And you've taught us about file systems where blocks are larger than cache lines. You probably don't even want a mapping for every cache line!

VM Mapping: A Better Idea

- Goal: We need a map from triples to physical address:
 - $(VA, \text{access}, \text{priv}) \Rightarrow PA \mid \text{FAULT}$
 - Physical addresses are not necessarily contiguous or ascending
 - We would like to avoid fragmentation.
- How might you implement this?
 - Just like we divide files into blocks, let's divide both virtual address spaces and memory (DRAM) into blocks. We shall call them pages.
 - Store one mapping for each page.
 - Critique:

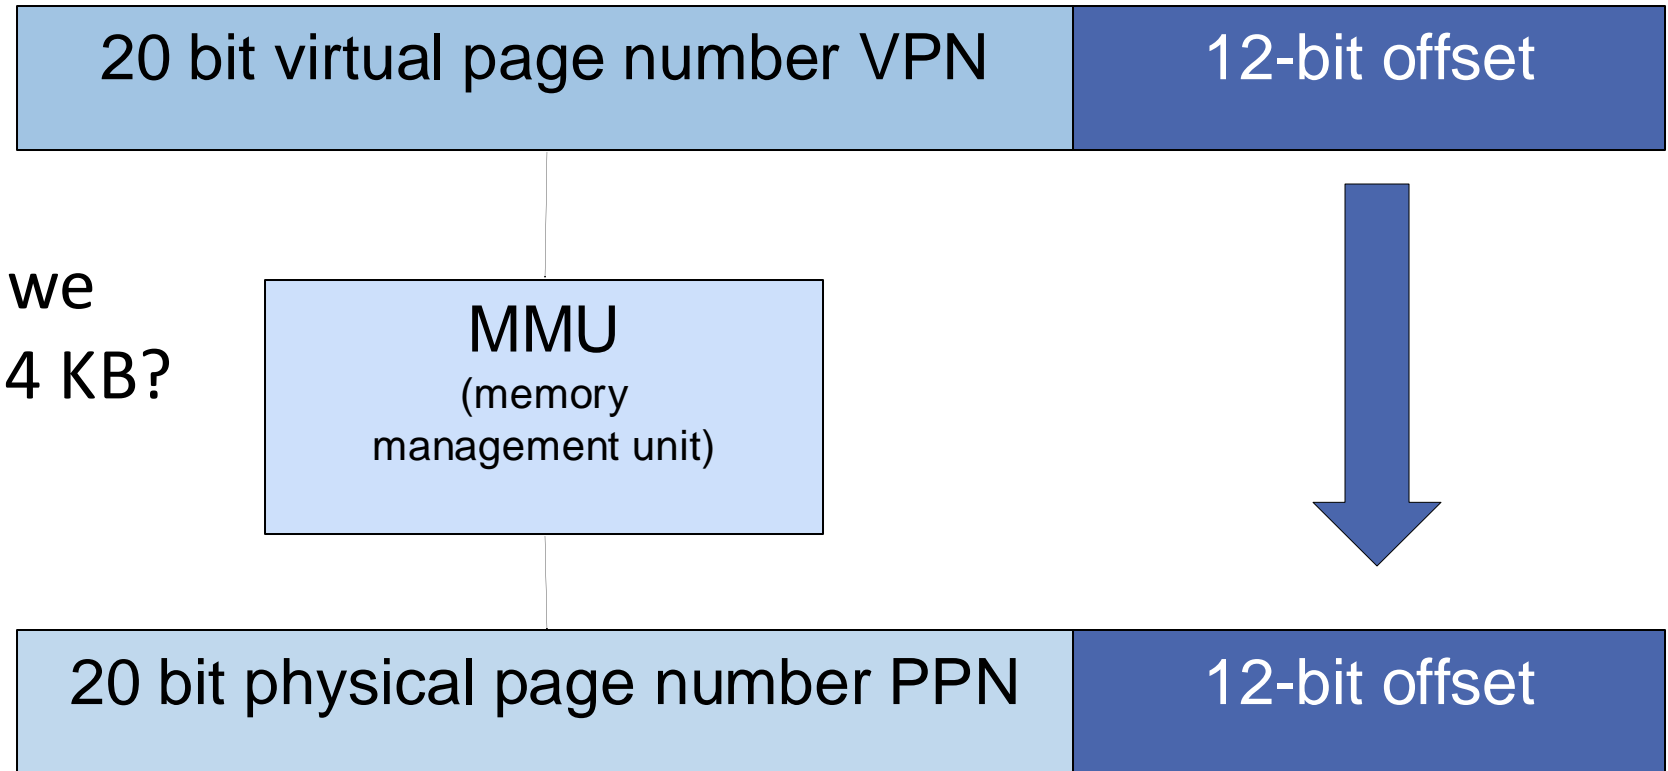
VM Mapping: A Better Idea

- Goal: We need a map from triples to physical address:
 - $(VA, \text{access}, \text{priv}) \Rightarrow PA \mid \text{FAULT}$
 - Physical addresses are not necessarily contiguous or ascending
 - We would like to avoid fragmentation.
- How might you implement this?
 - Just like we divide files into blocks, let's divide both virtual address spaces and memory (DRAM) into blocks. We shall call them pages.
 - Store one mapping for each page.
 - Critique: That seems far more sensible.
 - How large are pages?
 - How do you implement this map?

Dividing the bits in an address: x86-32

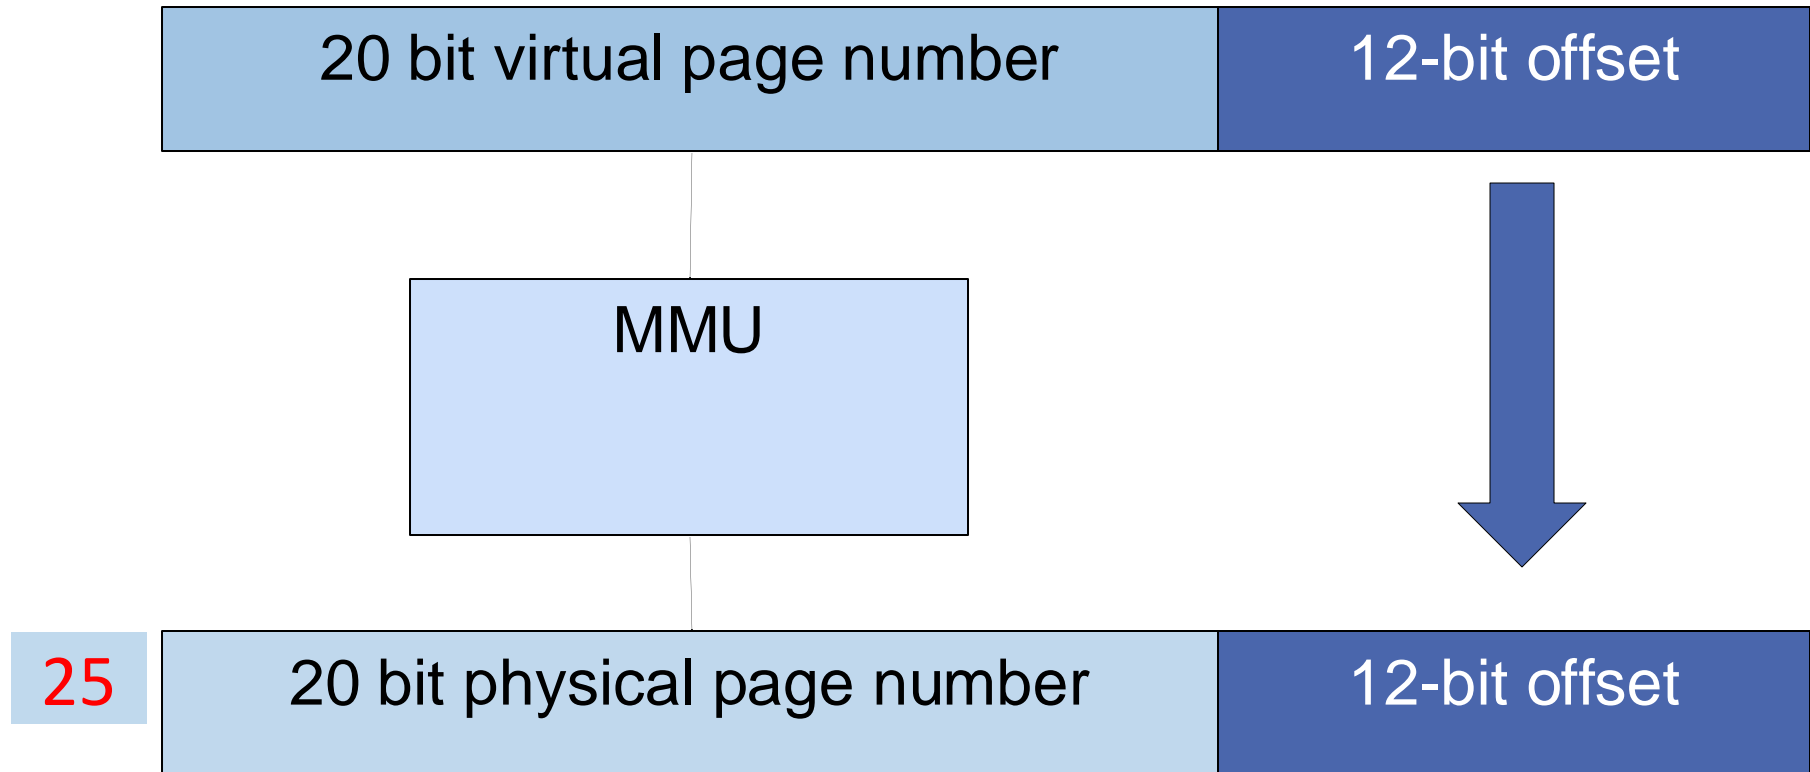
On x86:

- Pages are **4 KB**
 - i.e., 2^{12} B
- How many bits do we need to represent 4 KB?
12



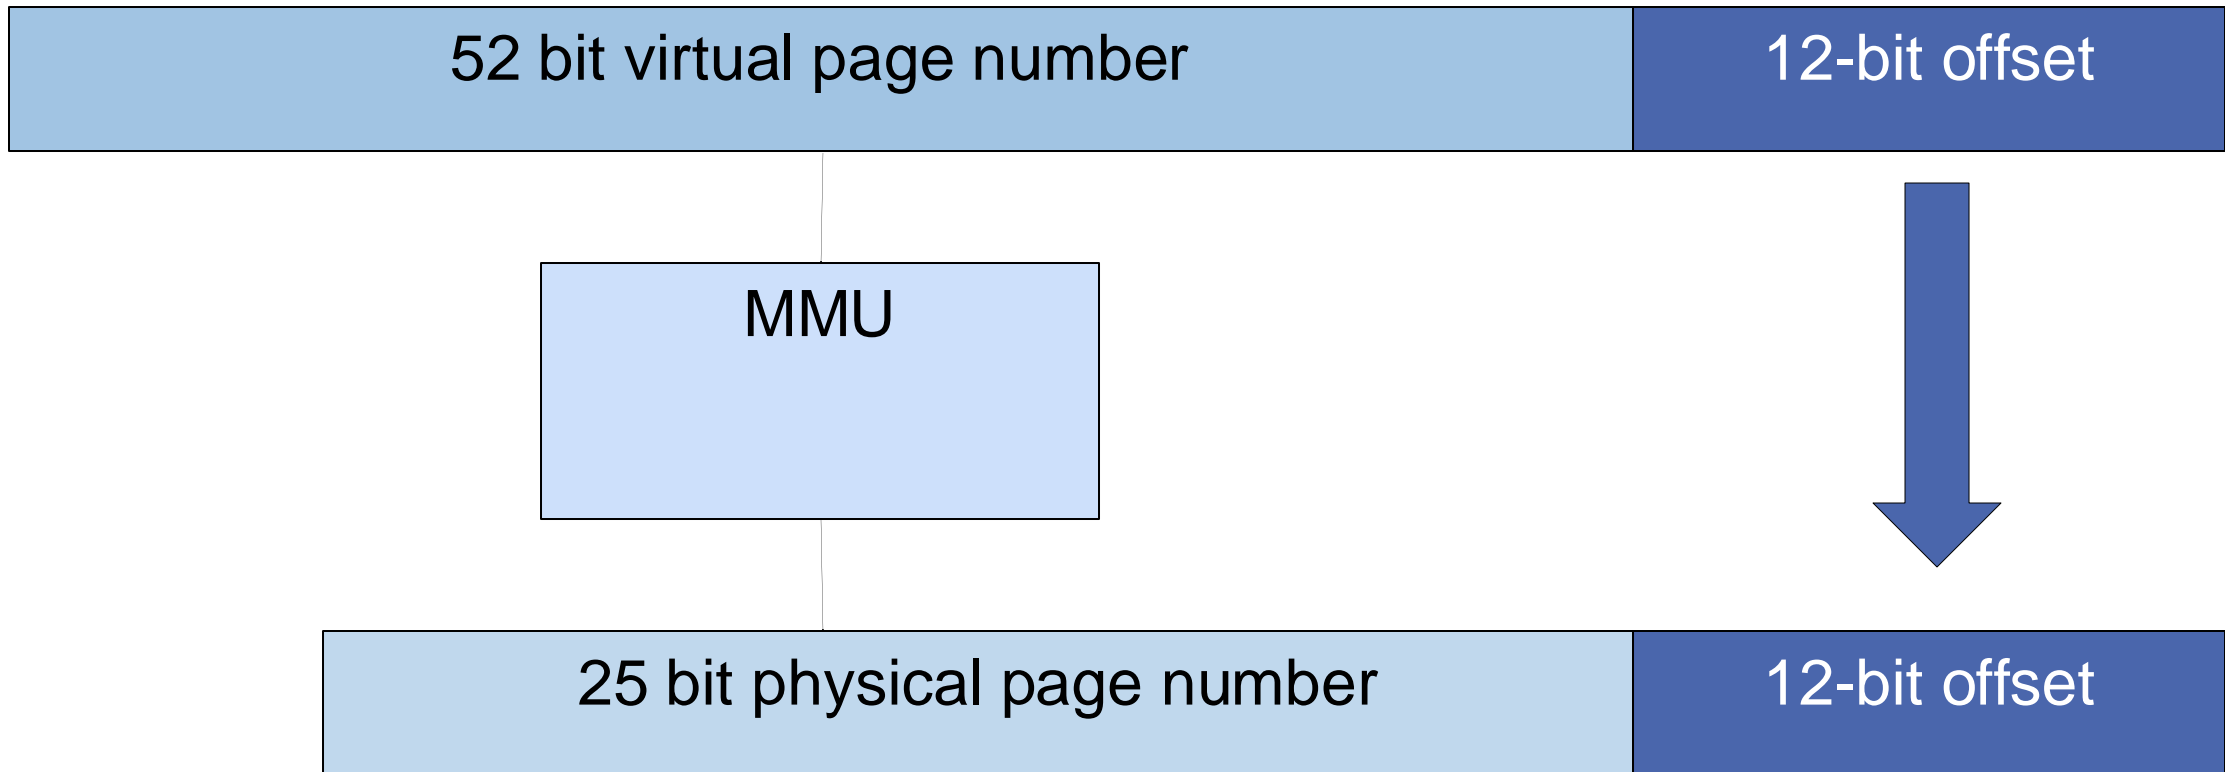
But wait ...

- What if we have a 32-bit VAS, but our machine has more than 4 GB of memory?




But wait – 2 ...

- What if we have a 64-bit VAS, while our machine has much less main/physical memory DRAM?




But wait – 2 ...

- What if we have a 64-bit VAS, while our machine has much less than that much main/physical memory DRAM?
 - That's fine!
 - We run in that regime all the time: the Intel x86-64 runs with a 48-bit address space, and pretty much none of us have 256 TB machines, and it all works.
 - How?
 - Our processes don't use the entire 48-bit address space.
 - Sometimes you **want** "holes" in memory.
 - You can run a process without having all its pages in memory.



Example: x86-64


- An instructor machine:
 - 3.3 GHz Dual-Core Intel Core i5
 - 64-bit processor (Kaby Lake)
 - 4 KB page size
 - Of the 64 bits in the virtual address space:
 - 12 bits are page offset
 - 16 bits are unused
 - **36 bits are the virtual page number**



- 16 GB 2133 MHz LPDDR3 DRAM
 - How many bits of physical page number do we need for this machine?

Example: x86-64

- An instructor machine:
 - 3.3 GHz Dual-Core Intel Core i5
 - 64-bit processor (Kaby Lake)
 - 4 KB page size
 - Of the 64 bits in the virtual address space:
 - 12 bits are page offset
 - 16 bits are unused
 - **36 bits are the virtual page number**

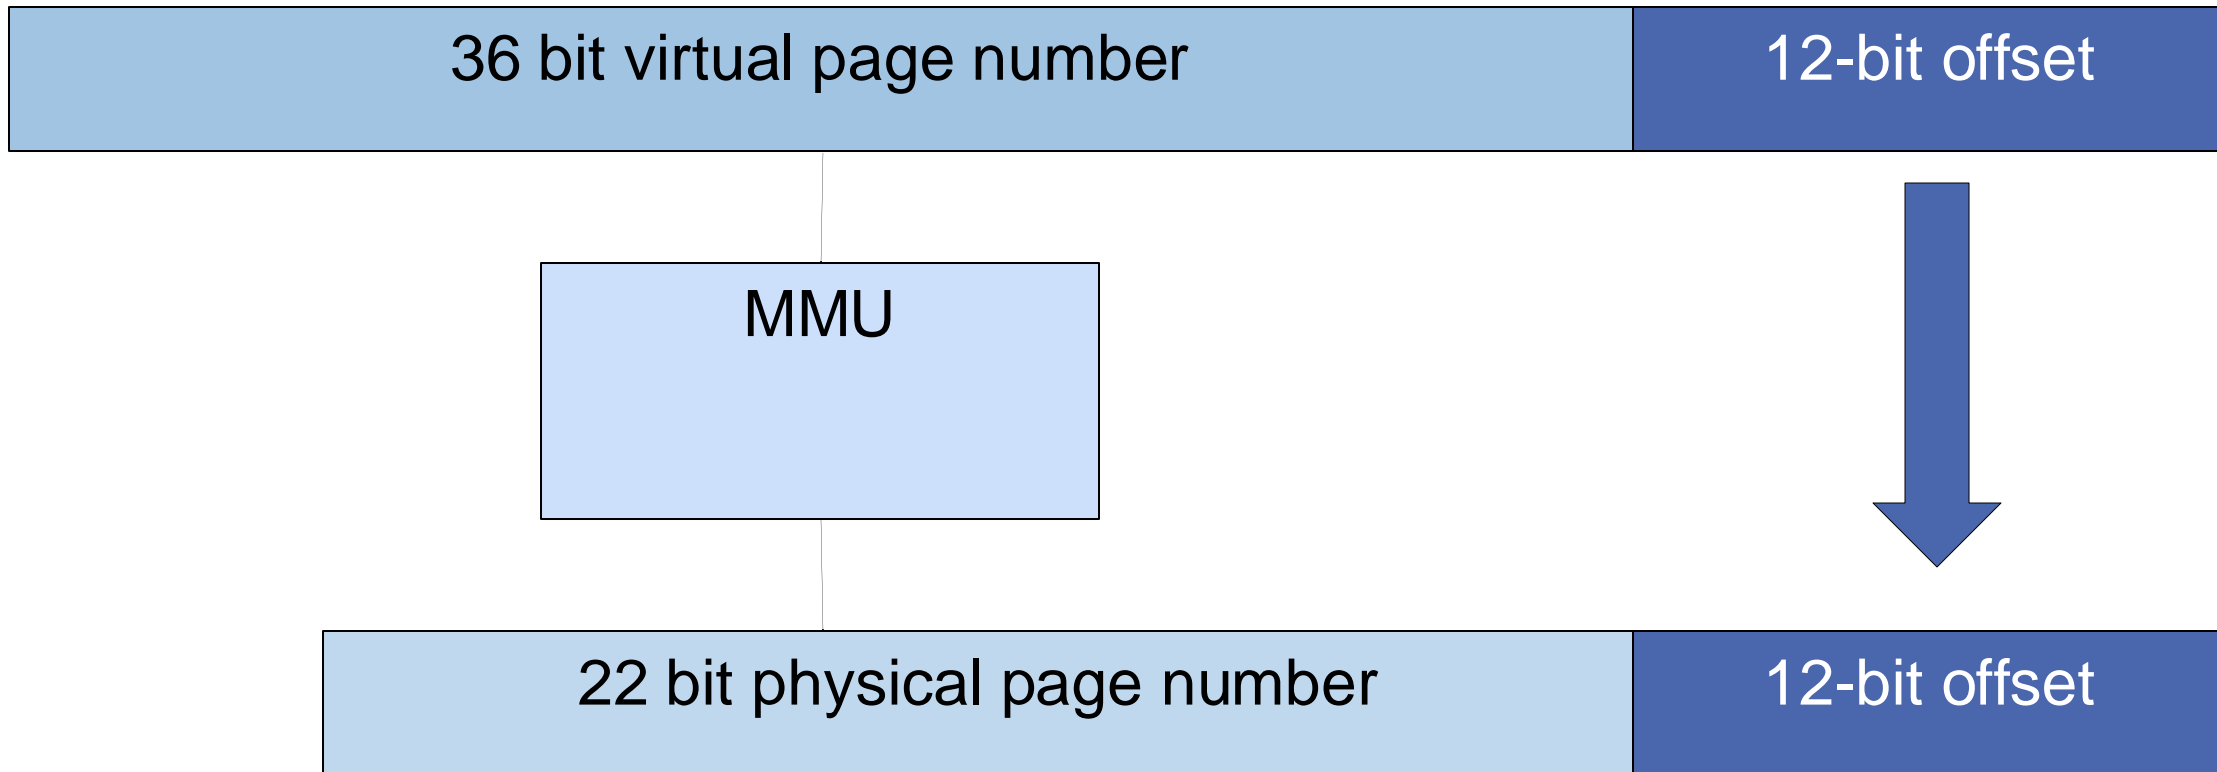


- 16 GB 2133 MHz LPDDR3 DRAM
 - How many bits of physical page number do we need for this machine?
16GB needs 30+4 bits for addressing. 34 total bits – 12 offset bits = 22 phys. page bits

*) Note: we ignore devices for now and assume the physical address space is just DRAM...

A typical 2023 machine

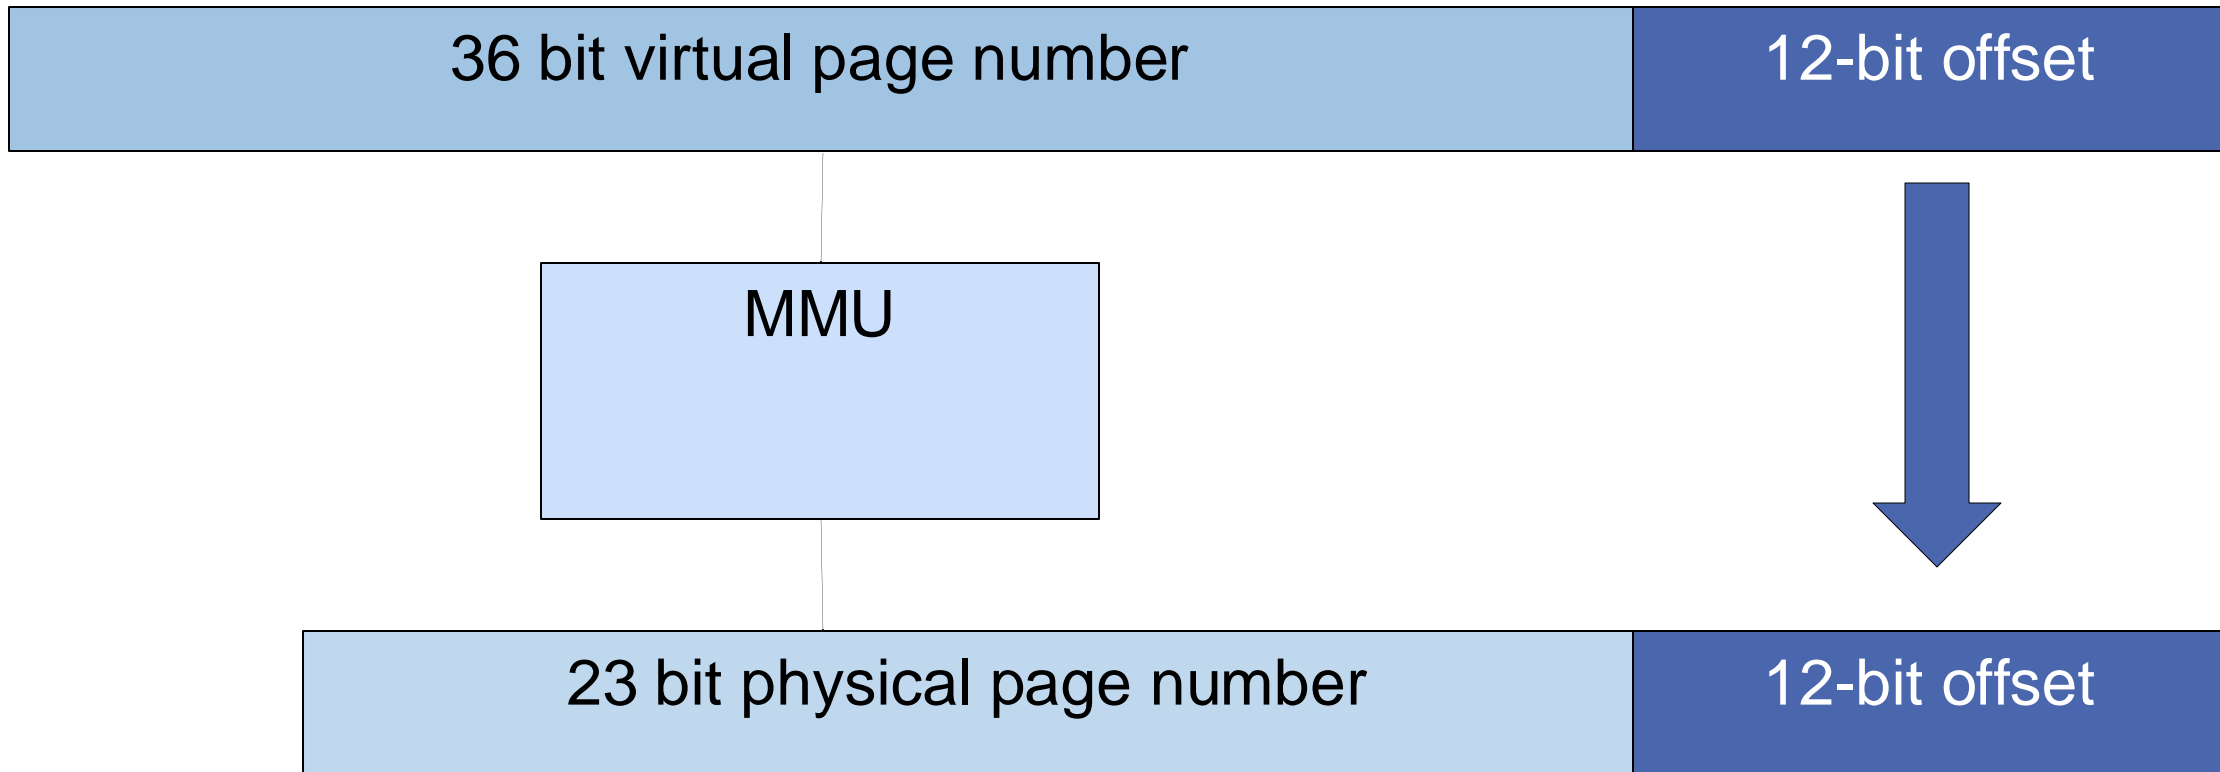
- 48-bit Virtual Address Space (256 TB)
- 34-bit Physical Address Space* (16GB)



*) Note: we ignore devices for now and assume the physical address space is just DRAM...

What if I had more memory?

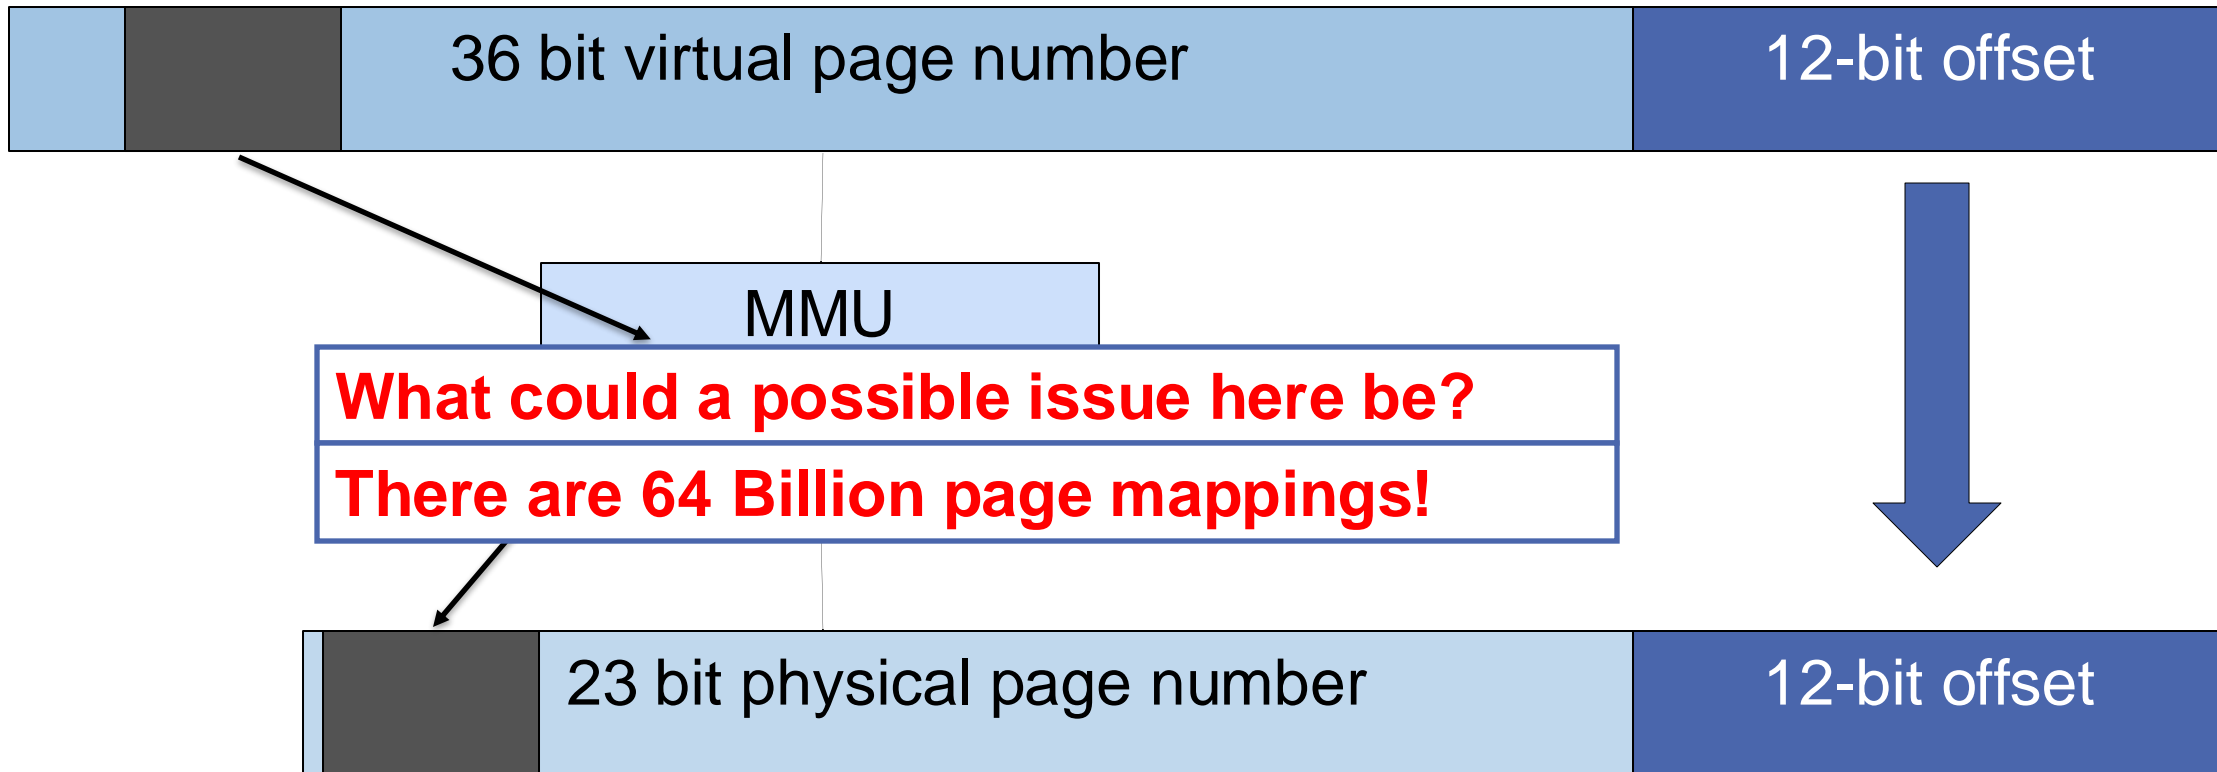
- 48-bit Virtual Address Space (256 TB)
- 35-bit Physical Address Space* (32GB)



*) Note: we ignore devices for now and assume the physical address space is just DRAM...

What if I had more memory?

- 48-bit Virtual Address Space (256 TB)
- 35-bit Physical Address Space* (32GB)



What Does our MMU do?

- Let's start with simple HW to go with our simple map.
- Rather than ask the MMU to store all the mappings, we can have the MMU **cache some mappings**.
- Why might this be OK?

What Does our MMU do?


- Let's start with simple HW to go with our simple map.
- Rather than ask the MMU to store all the mappings, we can have the MMU **cache some mappings**.
- Why might this be OK?
 - Instructions: These are pretty tiny and a lot of them fit in a single page, so a single mapping will likely get a lot of hits.
 - Data: Caching is premised on locality
 - If locality makes caching work, it ought to make caching translations work, too

Translation Lookaside Buffers (TLB)

- The simplest form of MMU is basically just a cache of translations (and permissions and protections).

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor aka Kernel
0x000000010003	0x0CAFE	R/X	User
0x000FFF0000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User

Code (i.e., instructions)




Trying to **write** to this page will cause a FAULT (bad permissions)

Translation Lookaside Buffers (TLB)

- The simplest form of MMU is basically just a cache of translations (and permissions and protections).

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User




Stack

Trying to **execute** data on this page will cause a FAULT (bad permissions)

Translation Lookaside Buffers (TLB)

- The simplest form of MMU is basically just a cache of translations (and permissions and protections).

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User




Heap

Trying to **execute** data on this page will cause a FAULT (bad permissions)

Translation Lookaside Buffers (TLB)

- The simplest form of MMU is basically just a cache of translations (and permissions and protections).

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User



Read only data


Trying to **write** or **execute** data on this page will cause a FAULT (bad permissions)

Translation Lookaside Buffers (TLB)

- The simplest form of MMU is basically just a cache of translations (and permissions and protections).

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User

Operating System (code)



If a **user program** tries to access this page, it will produce a FAULT.

Translating a Virtual Address (1)

```
* thread #1, queue = 'com.apple.main-thread', stop reason = breakpoint 1.1
  frame #0: 0x0000000010003f74 goodbye`main(...)
* This is an execute access from user mode
```

0 0 0 0 0 0 0 1 0 0 0 3 **F 7 4** ← Page offset

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x0000000010003	0x0CAFE	R/X	User
0x000FFF00000000	0x10011	R/W	User
0x0000000010005	0xFA0CE	R/W	User
0x0000000010004	0x54321	R	User

Translating a Virtual Address (2)

```
* thread #1, queue = 'com.apple.main-thread', stop reason = breakpoint 1.1
  frame #0: 0x0000000010003f74 goodbye`main(...)
* This is an execute access from user mode
```

Virtual Page Number

000000010003 F74

Page offset

Check access

Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User

Translating a Virtual Address (3)

```
* thread #1, queue = 'com.apple.main-thread', stop reason = breakpoint 1.1
  frame #0: 0x0000000010003f74 goodbye`main(...)
* This is an execute access from user mode
```

Virtual Page Number



Virtual Page Number	Physical Page Number	Permissions	Mode
0x00000000000000	0x10001	R/X	Supervisor
0x000000010003	0x0CAFE	R/X	User
0x000FFF000000	0x10011	R/W	User
0x000000010005	0xFA0CE	R/W	User
0x000000010004	0x54321	R	User


Physical Page Number



Takeaway


- Maximum virtual page **number** and physical page **number** can be different.
 - i.e., The **virtual address space** and **physical address space** can be different sizes.
- Virtual **pages** and physical **pages** are the same size.
 - The virtual page size is set by the hardware.
- The (maximum) size of the virtual address space is also determined by the hardware.
 - Therefore the number of bits in a virtual page number is also determined by the hardware.
- The (maximum) number of physical pages is also determined by the hardware
 - Therefore the number of bits in a physical page number is also determined by the hardware.
- The actual number of physical pages is specific to a particular machine.
- A **TLB** is a fast, hardware cache that maps virtual addresses to physical addresses (by storing mappings from virtual page numbers to physical page numbers).

What things looked like before:



What things look like now (maybe):

the point in the hierarchy where
we move from virtual to physical may differ



In class activity – becoming the TLB!

Think Ahead to Next Time

Our full page table maps what the process thinks are its page numbers to the **real** page numbers in memory.

There are a **lot** of (potential) pages in our virtual address space.

That means the page table itself will be **big**.

Worse, there are tons of holes/gaps in our virtual address space, and we don't want to waste memory storing those.

I wish we had a data structure that was good for this sort of mapping!!