

CPSC313: Computer Hardware and Operating Systems

Unit 5: Virtual Memory
Virtual Memory: The x86

Administrivia

- Quiz 5 coming soon!
- Lab 10 due Sunday
- Exam period office hours schedule coming soon (or already posted!)

Check Canvas/Piazza for details!

Today

- Roadmap:
 - So far, we've described page tables as huge arrays, cached in the TLB -- let's figure out why representing tables that way could be a problem.
 - We'll solve that problem and examine how the x86 represents page tables.
- Learning Outcomes
 - Describe the similarities and differences between a multi-level index for a file and multi-level page tables.
 - Describe the x86 Virtual Memory architecture
 - Generalize from the x86 to other configurations of virtual memory systems.

Things you know so far:

Virtual **pages** and physical **pages** are the same size.

- The page size is set by the hardware (CPU).
- The (maximum) size of the virtual address space is also determined by the hardware (CPU).
- Therefore, the number of bits or size of a virtual page number is also determined by the CPU.

Things you know so far:

- The **maximum** number of physical pages and therefore the size of physical page numbers is also determined by the hardware (CPU).
- The **actual** number of physical pages is specific to a particular machine (not only determined by its CPU).
- Virtual page **numbers** and physical page **numbers** can be different sizes.

Pre-class video implementation: one flat array

- There is one page-table per process.
- The OS must track page tables of all running processes.
- So, how big would those page tables be?
- Assume Intel x86-64
 - 4 KB pages
 - 48-bit address space
 - How many entries would we need in our page table?

Pre-class video implementation: one flat array

- Assume Intel x86-64
 - 4 KB pages
 - 48-bit address space
 - How many entries would we need in our page table?
 - 48 (bits in address space) – 12 bits (per 4 KB page) = 36 bits
 - 36 bits = 2^{36} = 64 G pages; that's more than 64 billion pages
 - We need 1 PTE per page; if each PTE is 8 bytes, **that's 512 GB per page table! Per process!**

Observations about Address Spaces

- What feature of address spaces might enable a better solution than a huge, VPN-indexed table?

Observations about Address Spaces

- What feature of address spaces might enable a better solution than a huge, VPN-indexed table?
 - Most processes won't use all or even most of their address space.
 - Many only use a few pages! Consider the hello program:
 - It consumes only 1 page each for text, data, stack
 - BUT: it also needs virtual address space for:
 - Shared libraries (code and data)
 - Still: total number of pages is measured in “a few hundred.”
 - **Page Tables are Sparse (i.e., most entries are invalid)**

How do we represent a sparse address space?

- Requirements:
 - **Simple** enough data structure to work with in hardware (on the x86).
 - **Efficient** for a large address space.
 - Supports gaps in the address space: **sparse address spaces**.
- Insight:
 - This is similar to the problem we had representing **files**: it had to be efficient for both small and large files and for sparse files.
 - *We were able to ignore whole chunks of the logical block number space by having NULL (invalid daddr_t) values for indirect block addresses.*
 - Maybe we could do something similar?

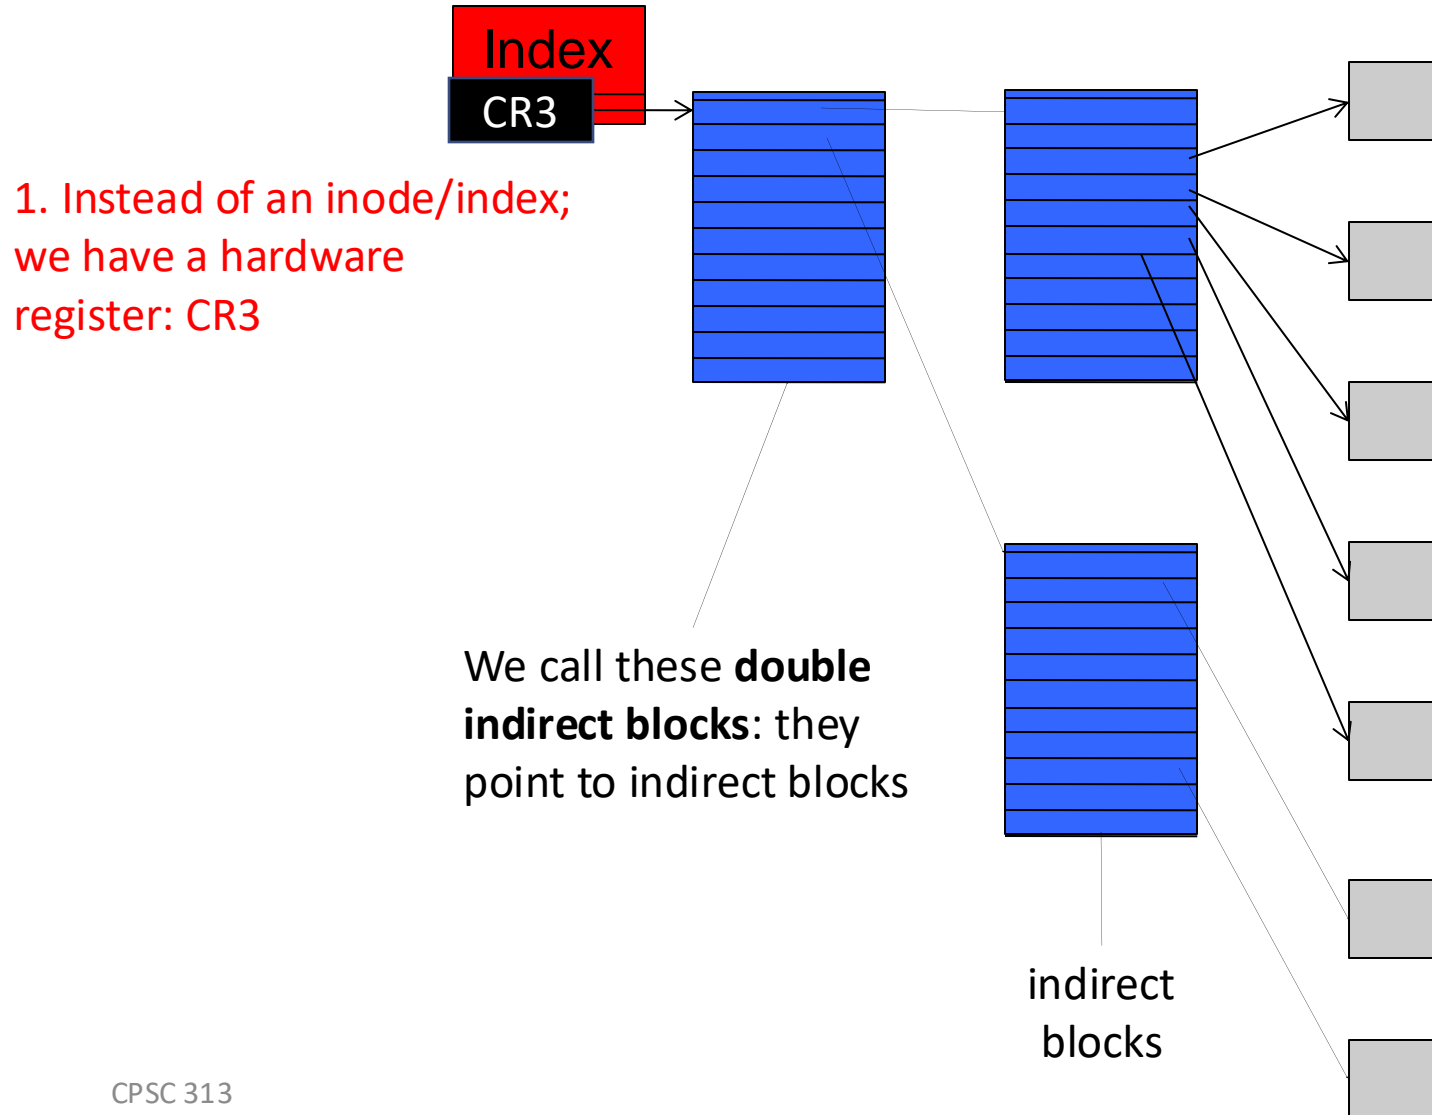
When implementing in HW

- Guiding principles:
 - Using bit values directly is good, cheap, simple
 - Computing things is relatively slower and more expensive

When implementing in HW

- Guiding principles:
 - Using bit values directly is good, cheap, simple
 - Computing things is relatively slower and more expensive
- We are essentially going to use a multi-level index, but we are going to choose the sizes of those indexes to match the hardware.
- In particular:
 - All blocks in the tree to match our virtual memory page size (4 KB)
 - We use specific bits in an address to index into each level of the tree

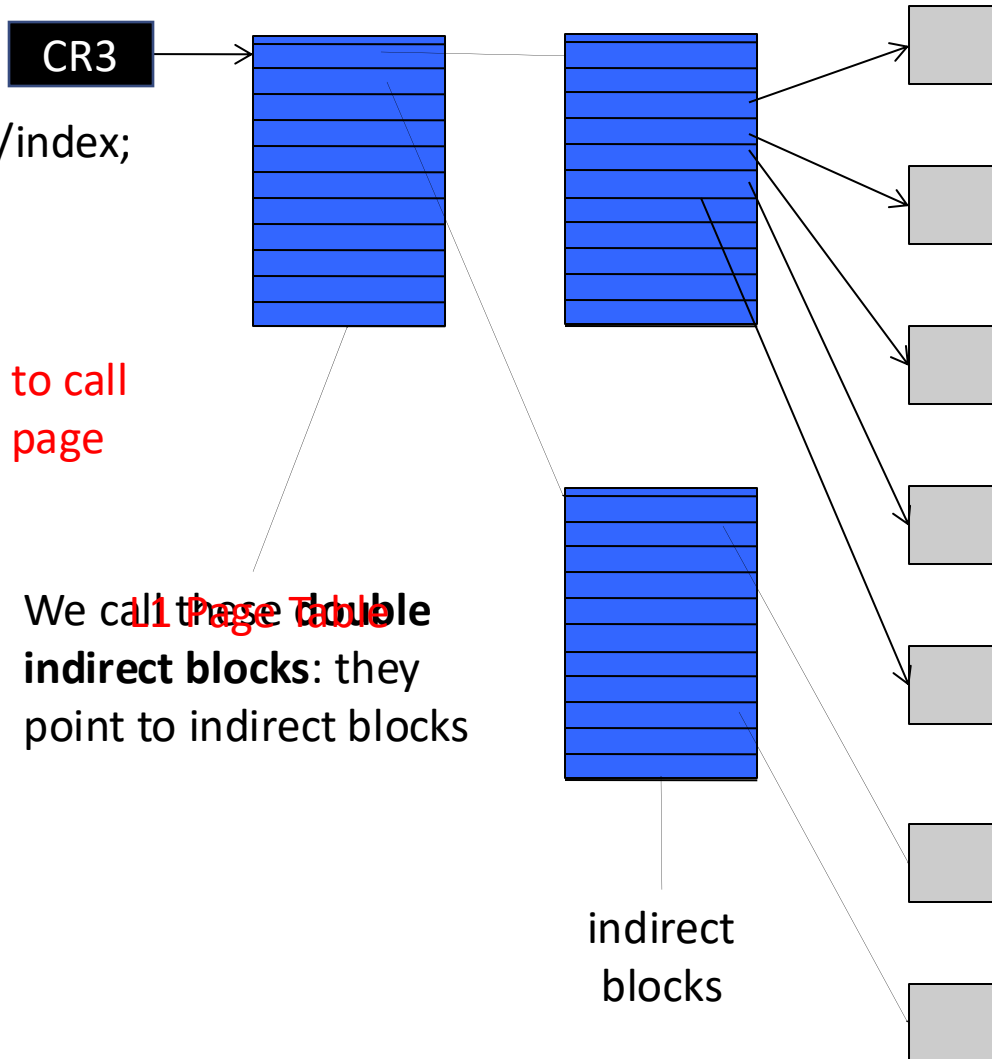
From a file index to a page table!



From a file index to a page table!

1. Instead of an inode/index;
we have a hardware
register: CR3

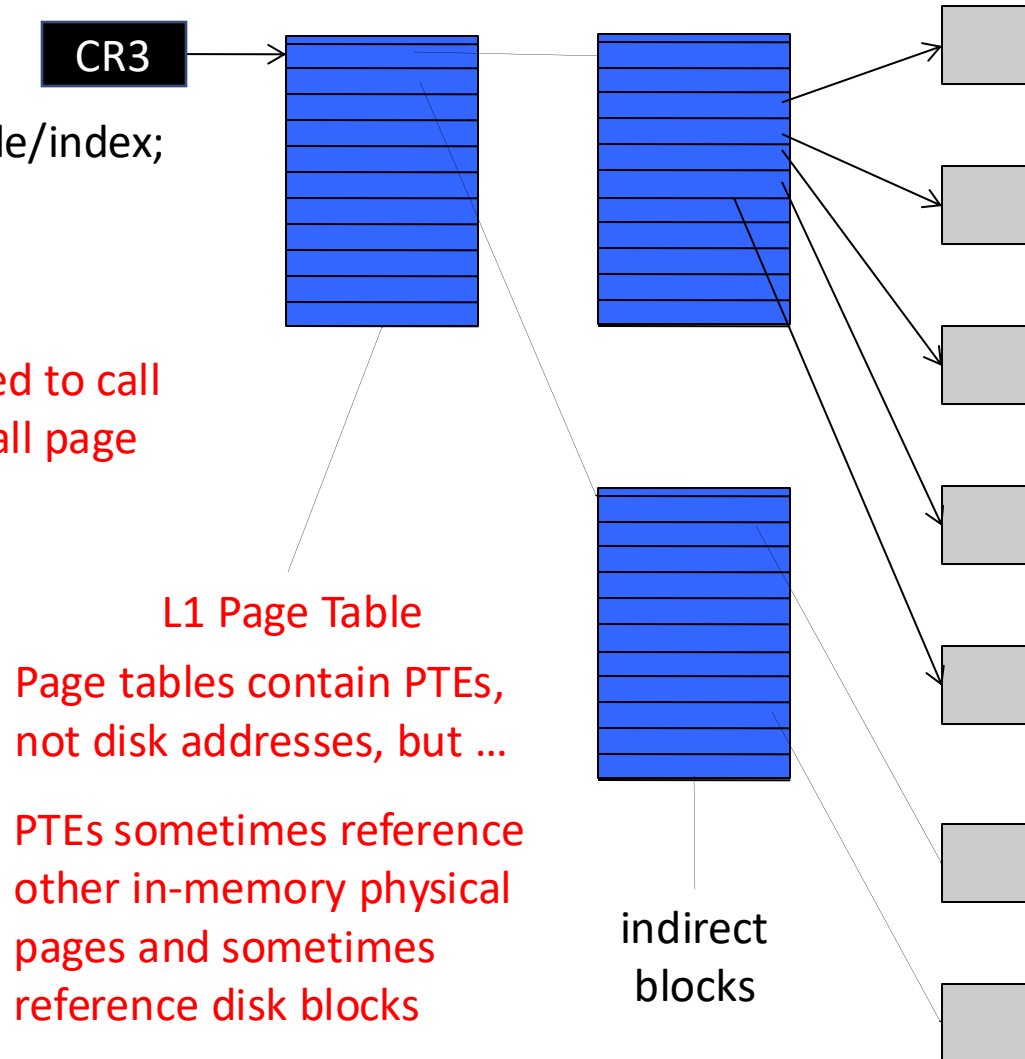
2. Everything we used to call
indirect blocks we call page
tables.



From a file index to a page table!

1. Instead of an inode/index;
we have a hardware
register: CR3

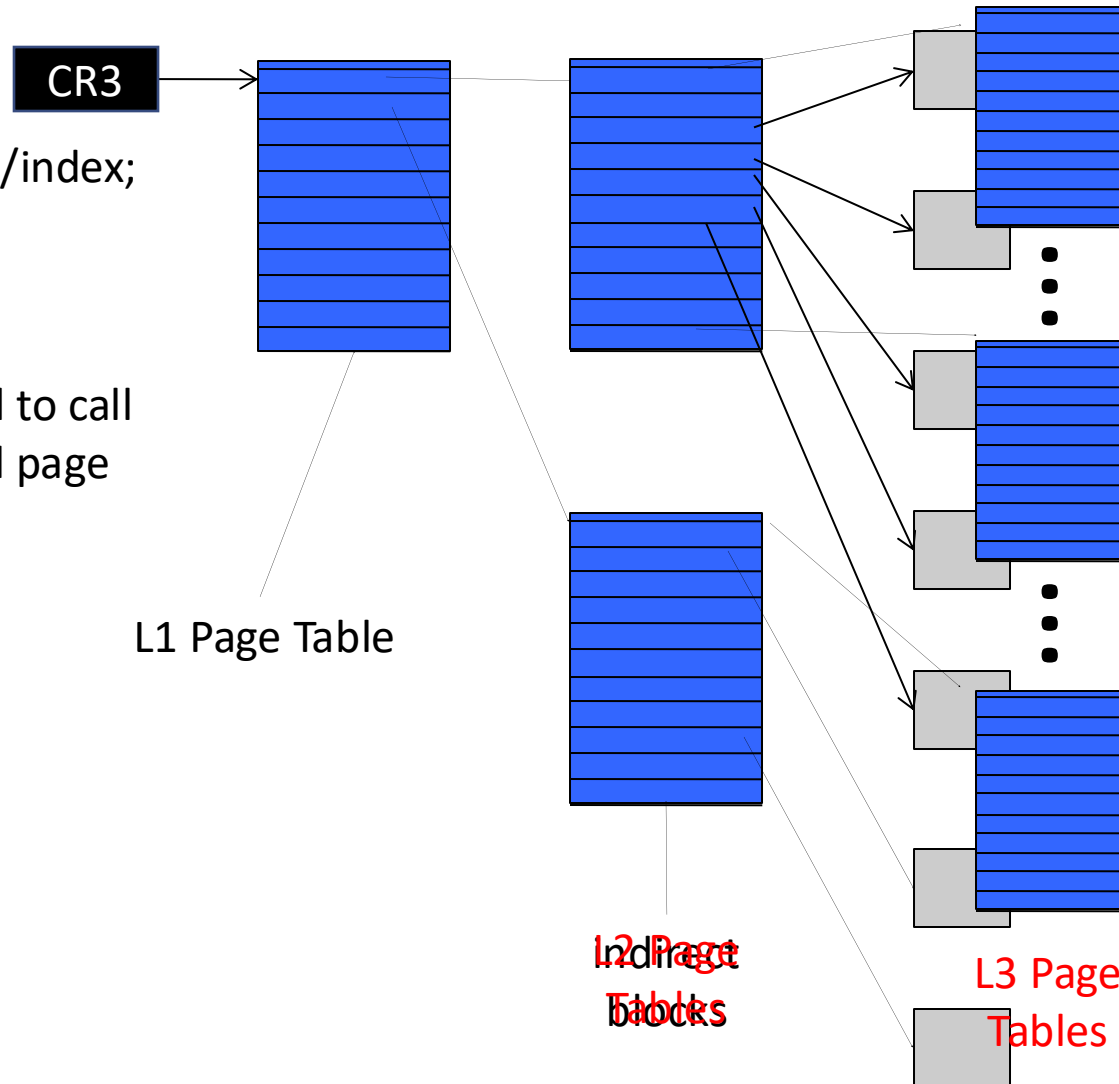
2. Everything we used to call
indirect blocks we call page
tables.



From a file index to a page table!

1. Instead of an inode/index; we have a hardware register: CR3

2. Everything we used to call indirect blocks we call page tables.



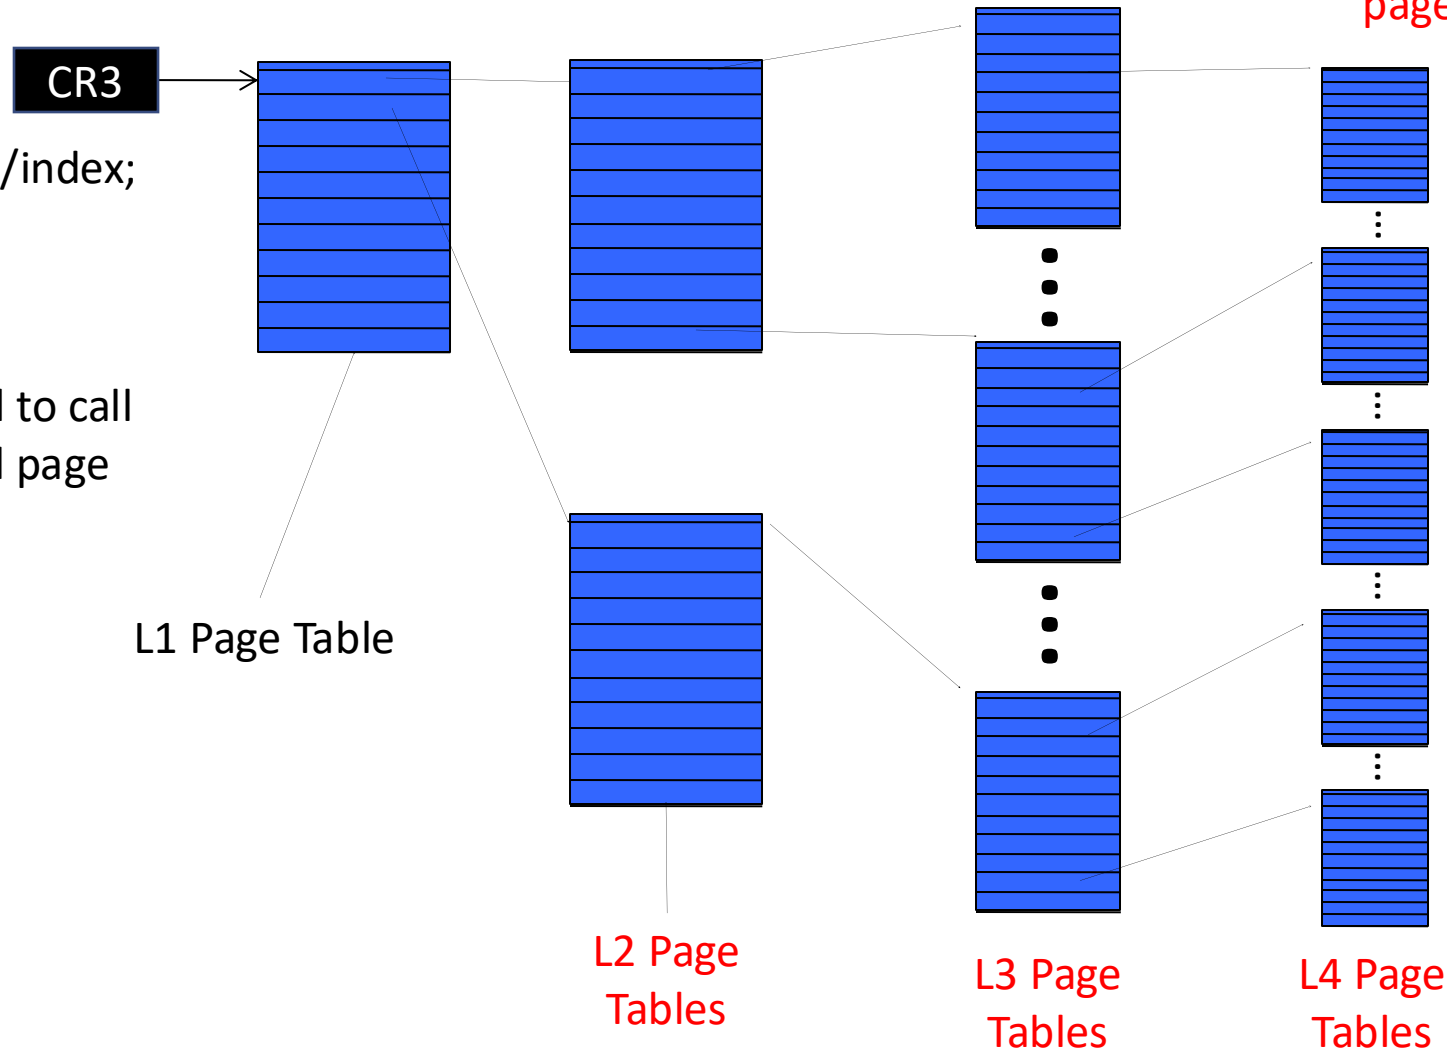
3. The x86-64 has 4-levels of page tables

From a file index to a page table!

1. Instead of an inode/index;
we have a hardware
register: CR3

2. Everything we used to call
indirect blocks we call page
tables.

3. The x86-64 has 4-levels of
page tables



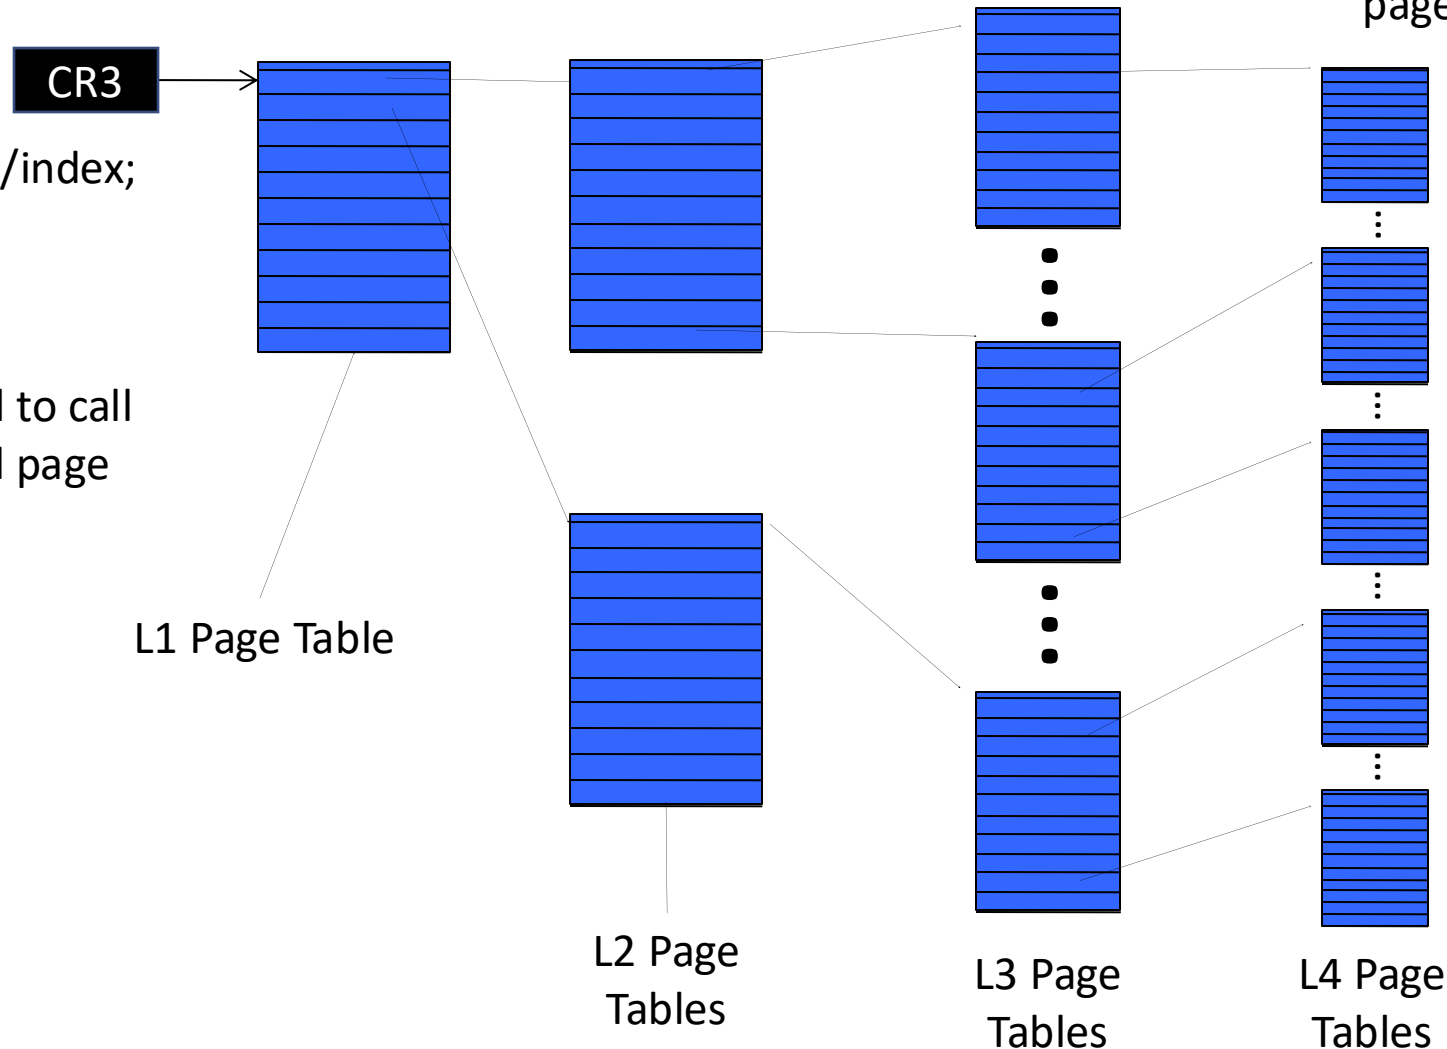
From a file index to a page table!

1. Instead of an inode/index; we have a hardware register: CR3

2. Everything we used to call indirect blocks we call page tables.

3. The x86-64 has 4-levels of page tables

4. All page tables are indexed by bits in the virtual address.



From a file index to a page table!


1. Instead of an inode/index; we have a hardware register: CR3

2. Everything we used to call indirect blocks we call page tables.

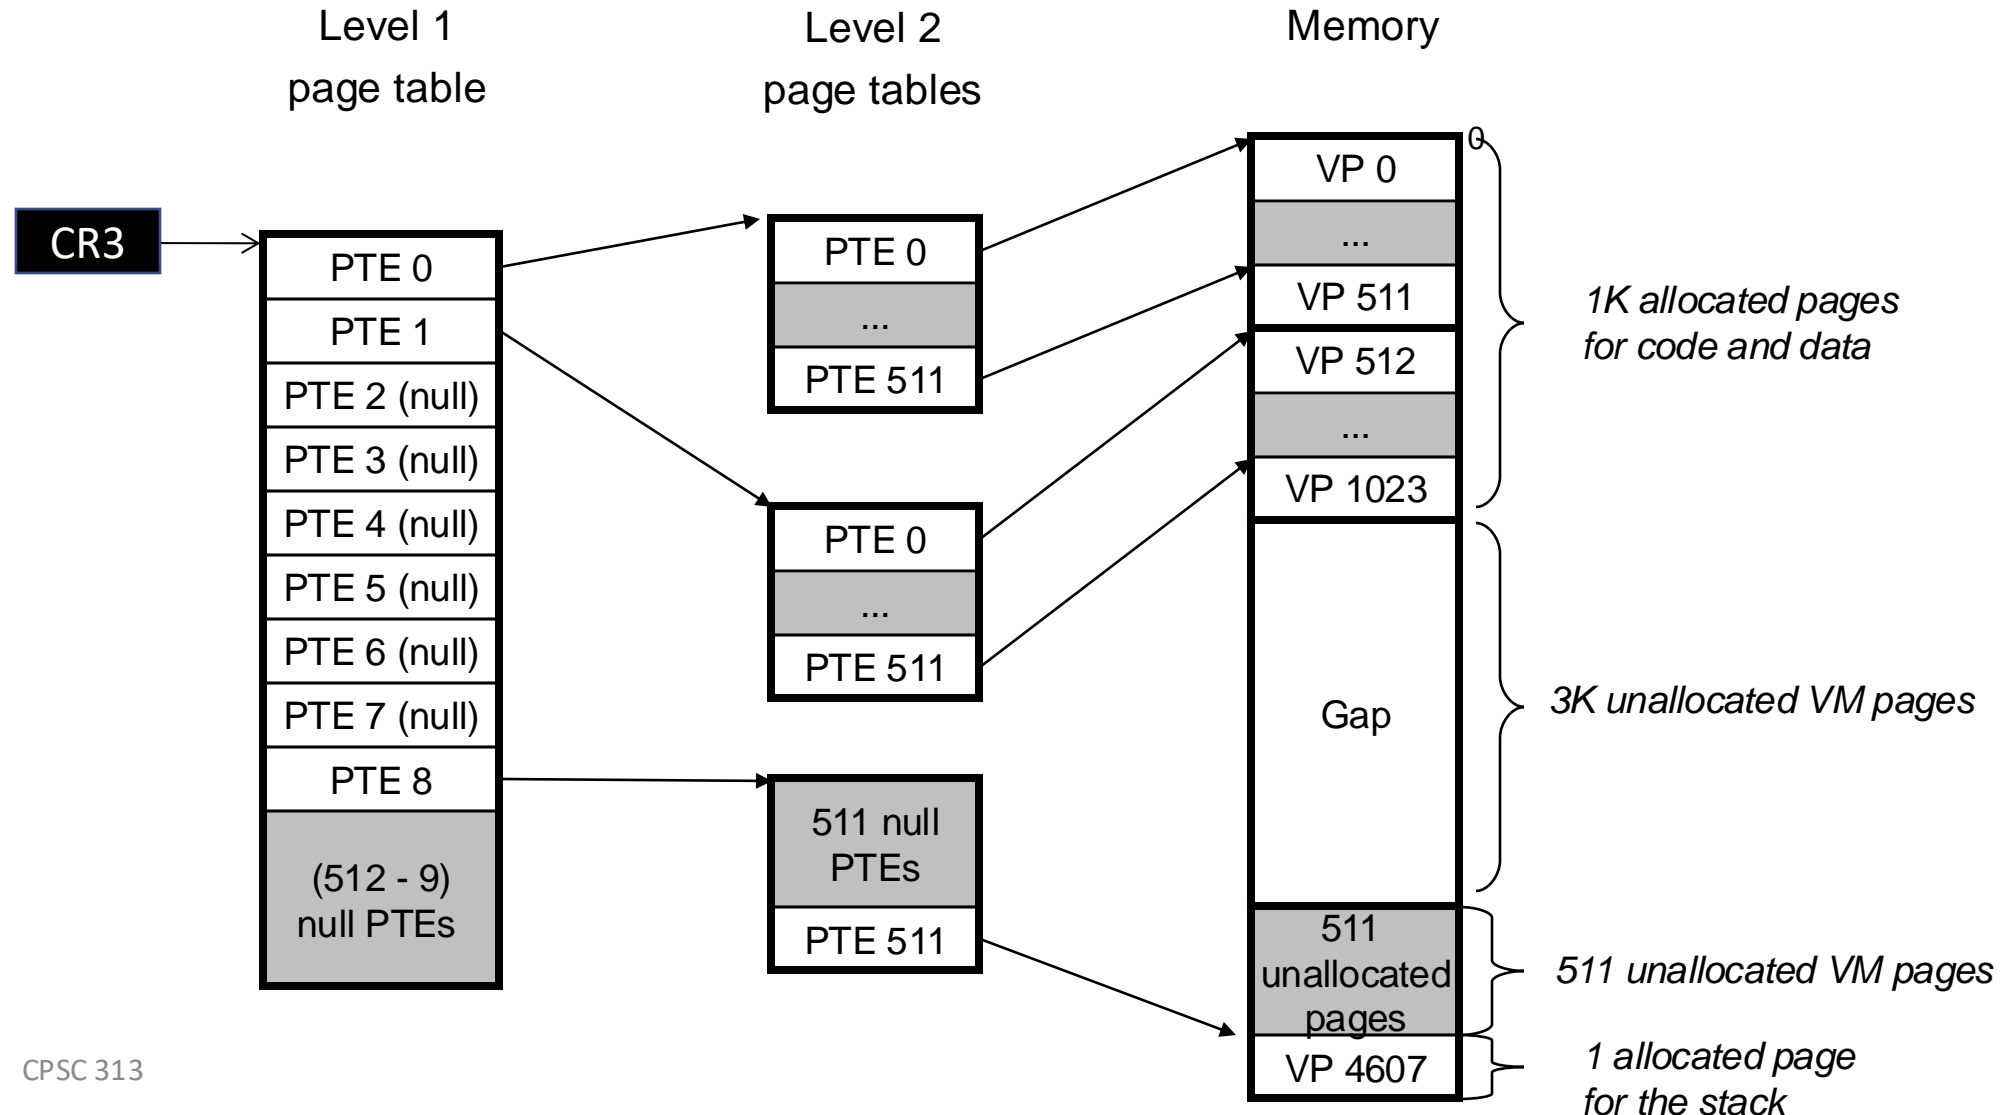
3. The x86-64 has 4-levels of page tables

4. All page tables are indexed by bits in the virtual address.

5. PTEs are 8 bytes



Example for one process (using 2 levels only)



X86 Address Translation

64-bit Virtual Address

X86 Address Translation

64-bit address



X86 Address Translation

64-bit address




X86 Address Translation

64-bit address




X86 Address Translation

64-bit address



X86 Address Translation

64-bit address



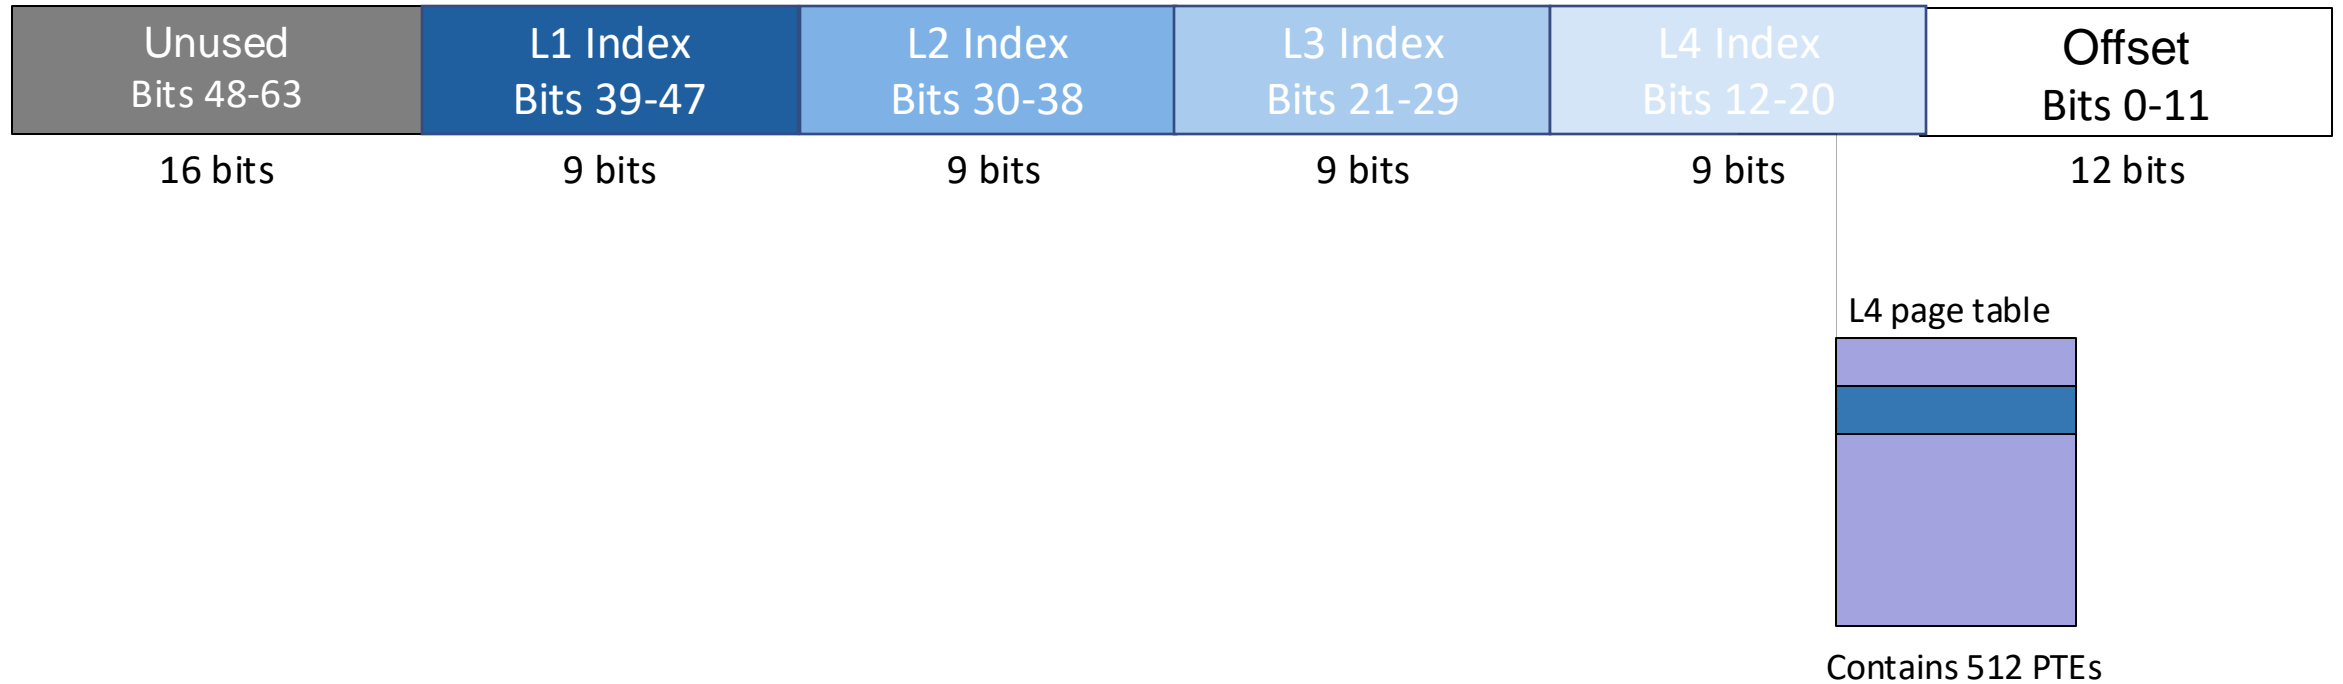
Each of these 9-bit chunks is going to be an index into a part of a page table.

Analogy to indirect blocks:

- An L4 index is for a page table containing PTEs that direct you to blocks in the address space
- An L3 index is for a page table containing PTEs that direct you to L4 page tables
- An L2 index is for a page table containing PTEs that direct you to L3 page tables
- An L1 index is for THE page table containing PTEs that direct you to L2 page tables

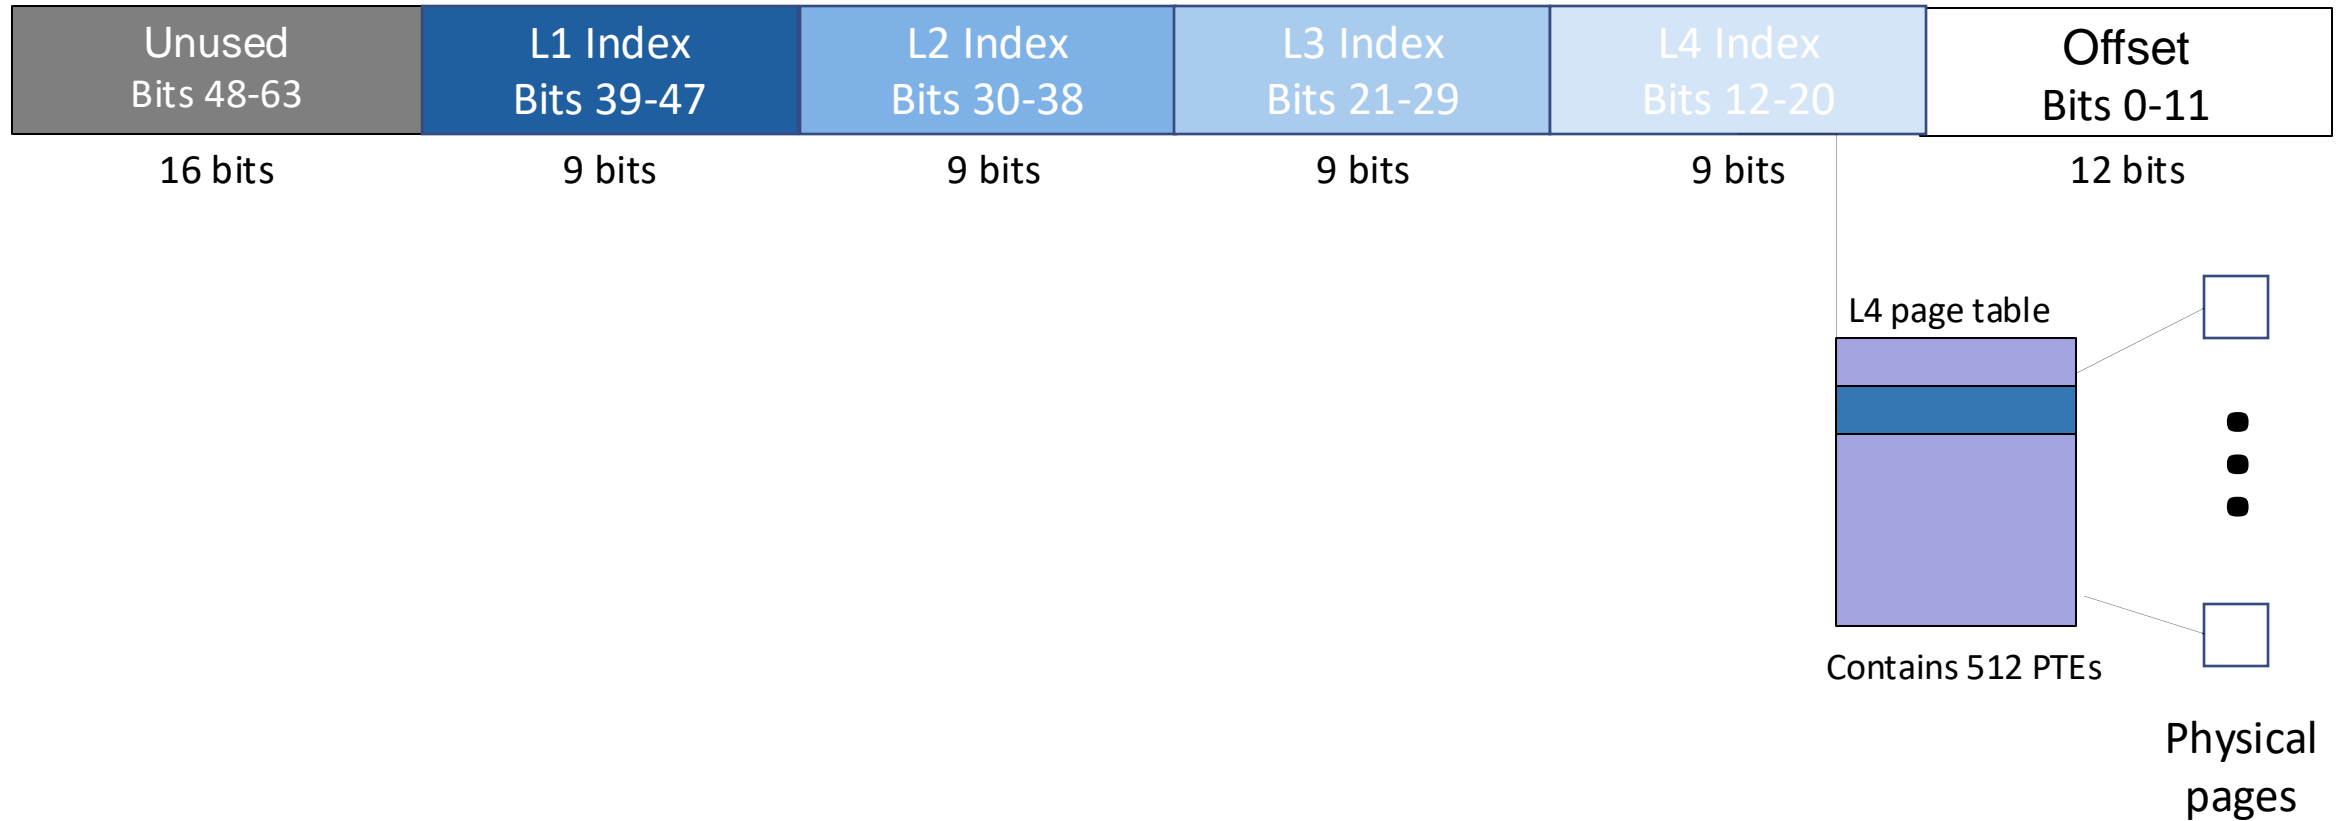
X86 Address Translation

64-bit address




X86 Address Translation

64-bit address




X86 Address Translation

64-bit address




X86 Address Translation

64-bit address




X86 Address Translation

64-bit address



X86 Address Translation

64-bit address



This is **exactly** the division/modulo process we use to find indices in a multi-level index!

A Closer Look at PTEs: Levels 1, 2, and 3

- Levels 1, 2, and 3 each refer to other page tables, if P=1:

XD	Unused	Page table physical base address	Unused		PS		A	CD	WT	U/S	R/ W	P
63	52-62	12-51	9-11	8	7	6	5	4	3	2	1	0

XD: Execute disable (can be used to disable instruction fetches from, e.g., stack)

Bits 12-51 contain the Physical page number of the page table referenced by this PTE

PS: Page size (safe to ignore this)

A: Reference bit (used for page replacement; check out the video); set on every access

CD: Are we allowed to cache the page table we reference

WT: Write through or write-back cache policy for the child page table

U/S: User or Supervisor mode for all pages reachable by the child table

R/W: Read-only or read/write permissions for all reachable pages

P: Is the child table present in main memory?

A Closer Look at PTEs: Levels 1, 2, and 3

- Levels 1, 2, and 3 each refer to other page tables:

XD	Unused	Page table physical base address	Unused		PS		A	CD	WT	U/S	R/ W	P
63	52-62	12-51	9-11	8	7	6	5	4	3	2	1	0

If P = 0, that says that the page table is not in memory (DRAM); the OS is free to use the rest of the entry in whatever way it wants; what kinds of information do you think the OS wants to put there?

The OS can do whatever it wants with these bits!	P
Bits 1-63	0

A Closer Look at PTEs: Level 4

- Levels 4 PTEs refer to pages that belong to the process' VAS, if P=1

XD	Unused	Page physical base address	Unused	G	PS	D	A	CD	WT	U/S	R/ W	P
63	52-62	12-51	9-11	8	7	6	5	4	3	2	1	0

XD: Execute disable (can be used to disable instruction fetches from, e.g., stack)

Bits 12-51 contain the Physical page number of the page referenced by this PTE

G: Global page (is not evicted from TLB on a process switch)

PS: Page size (safe to ignore this)

D: Dirty bit (set on every write)

A: Reference bit (used for page replacement; stay tuned until next week); set on every access

CD: Are we allowed to cache the page table we reference

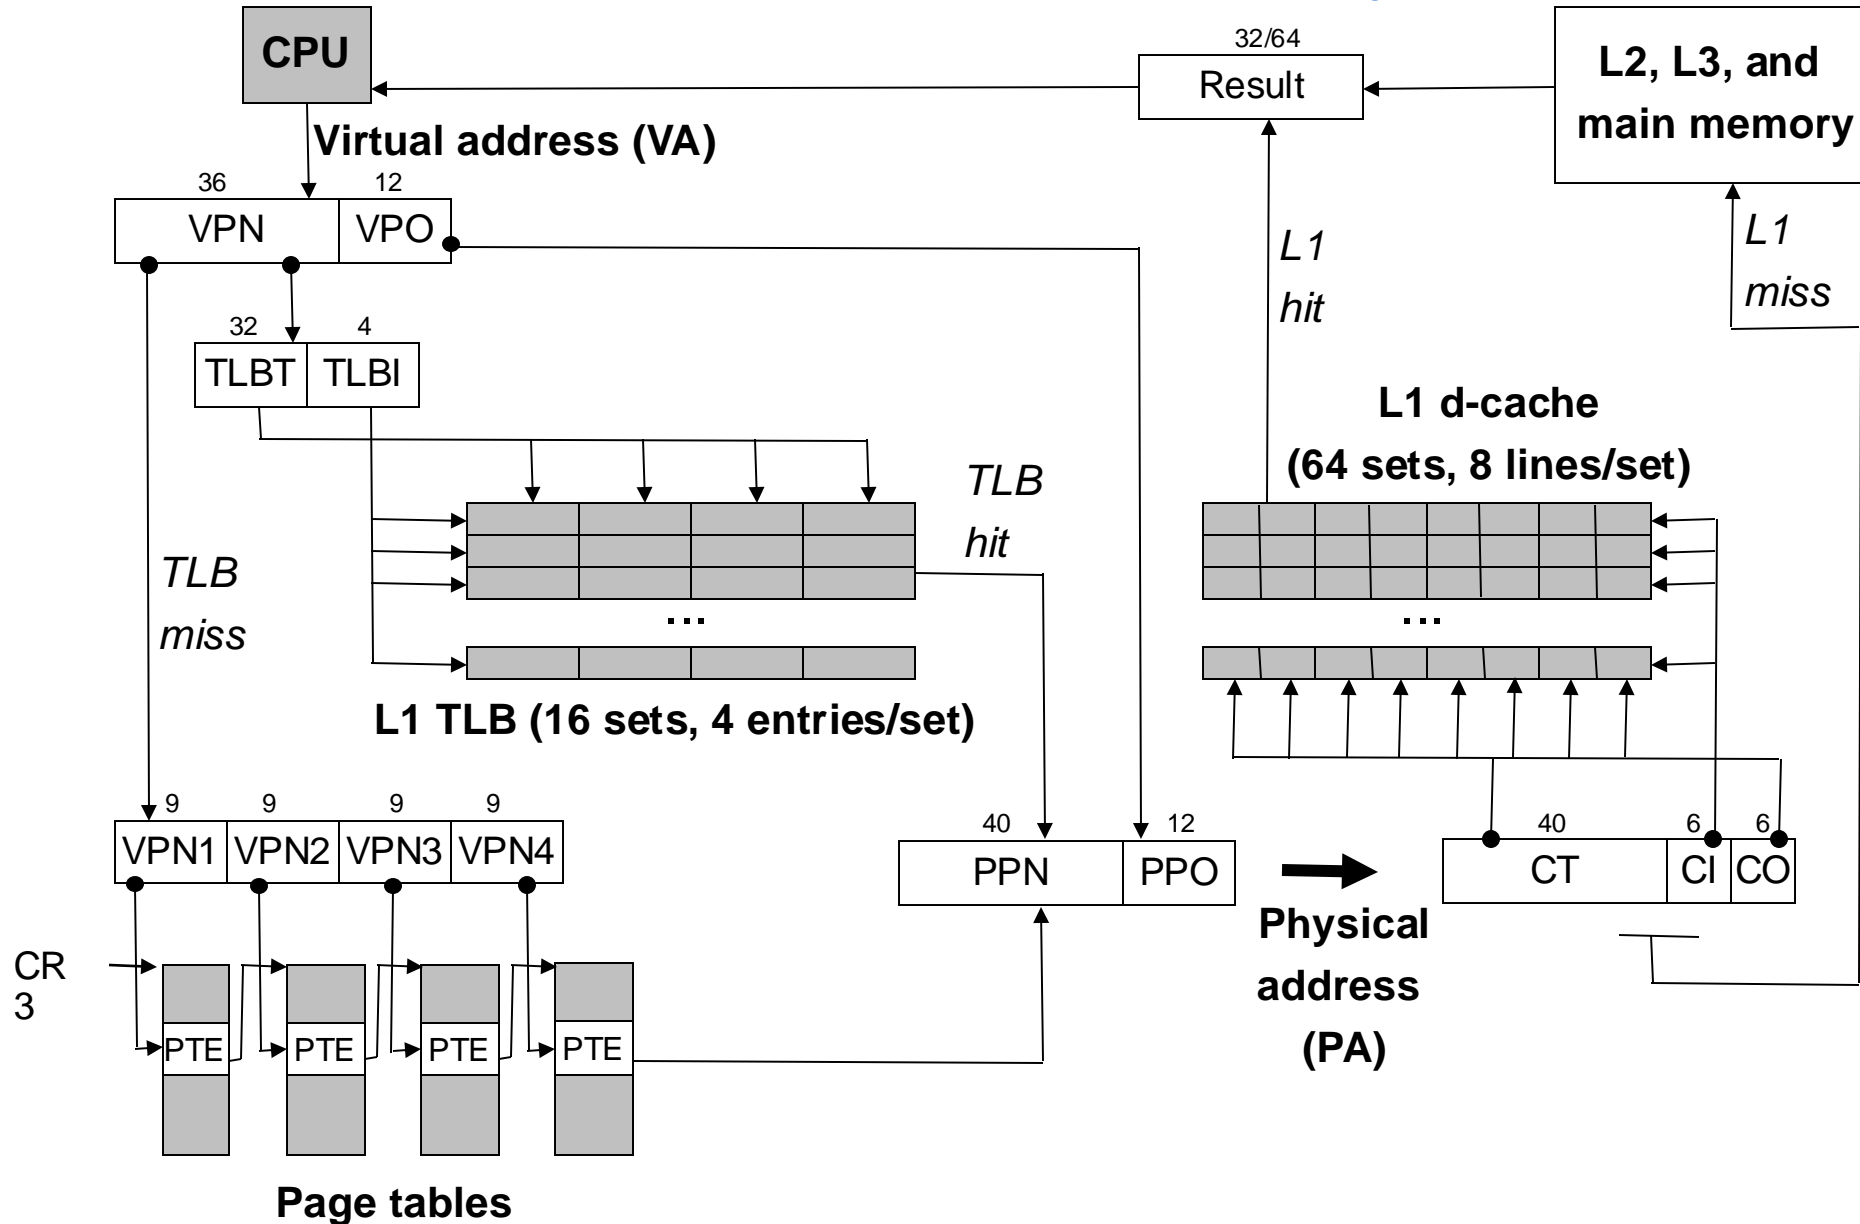
WT: Write through or write-back cache policy for the child page table

U/S: User or Supervisor mode for all pages reachable by the child table

R/W: Read-only or read/write permissions for all reachable pages

P: Is the page present in main memory? (If P=0, the OS again uses the rest for whatever it wants!)

A closer look at the x64-64 VM system



Questions to test your understanding:

- Think about these before the next class!
1. Must cr3 contain a physical or virtual address? Why?
 2. How large is a single page table? Why?
 3. Why do PTEs contain physical page numbers?
 4. Why don't PTEs need an offset? (I.e., why is a page number sufficient?)
 5. For each level page table, how much memory is reachable?
 6. On the previous slide, the CI + CO happened to fit in the VPO. Is that necessary? Is it useful in any way?

Intel x86 VM Historical Notes

- The Intel x86 Virtual Memory Architecture reflects many major revisions that have occurred over various generations of Intel microprocessors.
- While this makes the system a bit more complicated than some others, it is the most widely used platform today, so understanding it will serve you well.
- Note:
 - This presentation did not cover x86 in its full glory.
 - We cover it sufficiently so should you ever need to dig into the details, they will make sense.