


Today


- Roadmap
 - Page Tables: How the OS maps virtual addresses to physical addresses
- Learning Outcomes
 - Define:
 - Page Table
 - Page Table Entry (PTE)
 - Segfault
 - Page Fault
 - Explain what each field in a PTE means
- Reading
 - 9.5

Recall how we represent Address Spaces



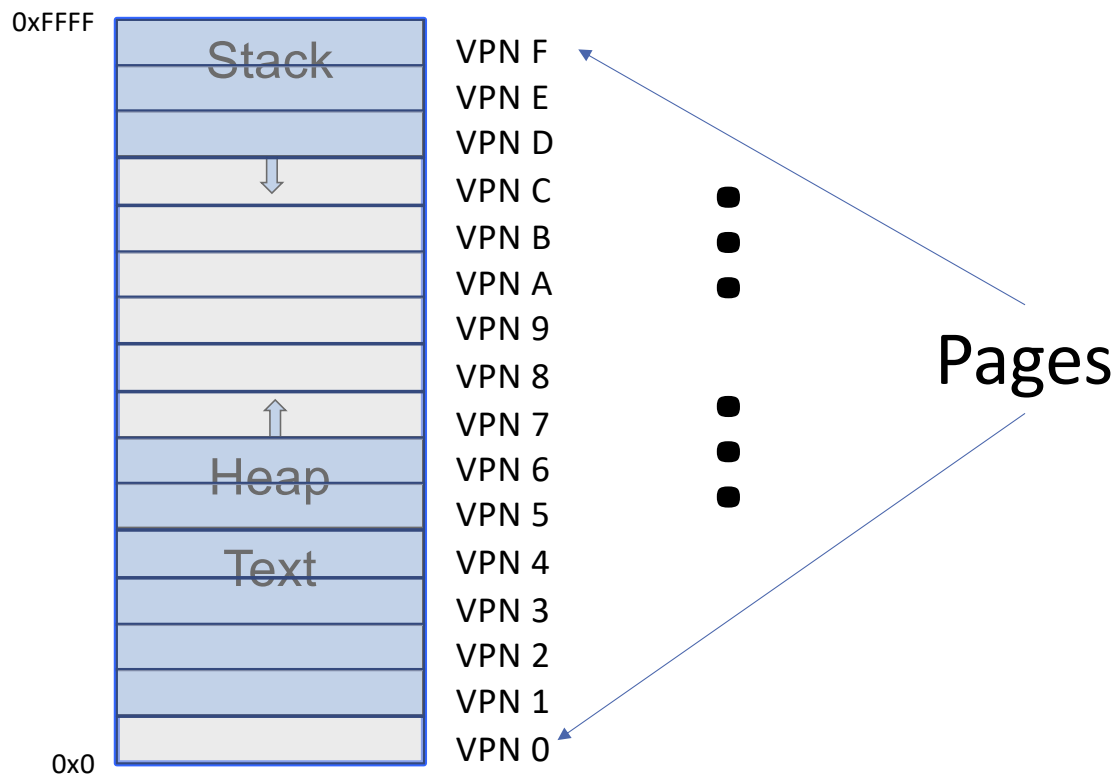
Back to Address Spaces

Program 1




Back to Address Spaces

Program 1



Back to Address Spaces

Program 1




Notice:

- The address space is 16 bits (4 hex digits)
- There are 16 pages (0-F)
- Therefore, this must be a 4 KB page system.

Mappings: Recall

Program 1



(VA, access, privilege) => (PA/fault)

Triple consisting of a:

- Virtual address (VA)
- Access (read, write, execute)
- Privilege (user, supervisor)

Physical address
OR

Fault (turns control
over to the
operating system)

And what would we call that fault?

An exception!

Page Tables: Mapping data structure


- The TLB is simply a cache of mappings.
- The **Page Table** is the data structure that holds all the mappings.

Indexed by virtual page number (VPN)

[illegible]

Page Tables: A pile of PTEs

Program 1



PTE


indices

Page table

| | PPN | Access | Privilege |
|---|---------|---------------|-----------|
| 0 | Invalid | | |
| 1 | 0x0000A | Read, Execute | U |
| 2 | 0x1100B | Read, Execute | U |
| 3 | 0x98765 | Read, Execute | U |
| 4 | 0xCAFE0 | Read, Execute | U |
| 5 | 0xFACE1 | Read, Write | U |
| 6 | 0xC0FFF | Read, Write | U |
| 7 | Invalid | | |
| 8 | Invalid | | |
| 9 | Invalid | | |
| A | Invalid | | |
| B | Invalid | | |
| C | Invalid | | |
| D | 0x50505 | Read, Write | U |
| E | 0x12345 | Read, Write | U |
| F | 0x24680 | Read, Write | U |

Address Translation

Program 1



Assume we are running in user mode

`xor %rax, %rax`

`mrmovq 0xFEED(%rax), %rbx`

What virtual address are we accessing?

`0xFEED`

What is the VPN to translate?

`0xF`

What is the corresponding PTE?

`0x24680/RW/U`

Do we have the proper permissions?

`Yes`

What is the PA for the data?

`0x24680EED`

| | PPN | Access | Privilege |
|---|---------|---------------|-----------|
| 0 | Invalid | | |
| 1 | 0x0000A | Read, Execute | U |
| 2 | 0x1100B | Read, Execute | U |
| 3 | 0x98765 | Read, Execute | U |
| 4 | 0xCAFE0 | Read, Execute | U |
| 5 | 0xFACE1 | Read, Write | U |
| 6 | 0xC0FFF | Read, Write | U |
| 7 | Invalid | | |
| 8 | Invalid | | |
| 9 | Invalid | | |
| A | Invalid | | |
| B | Invalid | | |
| C | Invalid | | |
| D | 0x50505 | Read, Write | U |
| E | 0x12345 | Read, Write | U |
| F | 0x24680 | Read, Write | U |

Page Table Entries: PTEs

- A PTE can be in one of three states
 1. Invalid: There is no mapping
 - Exception: Segfault (usually kill the process)
 2. Valid and Memory-Resident: Permissions dictate if the access is allowed
 - Execution proceeds as normal
 3. Valid but not Memory-Resident: Permissions still dictate if the access is allowed. While the VPN is valid, the page we want to access is not yet in memory; the OS must make it memory resident.
 - Could require reading the page from disk
 - Could require creating a page full of 0's
 - Exception: Page fault (read the page from disk and then proceed as normal)

Page Tables: Summary

- The page table is the collection of mappings that describe an address space.
- The page table contains a set of page table entries (PTEs).
- A PTE can be in one of three states:
 - invalid -- physical page number and permissions should be ignored
 - valid and present -- contains physical page number, access and mode info
 - valid and not present -- location of physical page on disk, access and mode info